

openSUSE-KIWI Image System

Cookbook

Marcus Schäfer

openSUSE-KIWI Image System: Cookbook

by Marcus Schäfer

Thomas Schraitle <toms@suse.de>

Robert Schweikert <rjschwei@suse.com>

KIWI Version 5.05

License

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.2 or (at your option) version 1.3; with the Invariant Section being this copyright notice and license. A copy of the license version 1.2 is included in the appendix entitled “GNU Free Documentation License”.

SUSE®, openSUSE®, the openSUSE® logo, Novell®, the Novell® logo, the N® logo, are registered trademarks of Novell, Inc. in the United States and other countries. Linux® is a registered trademark of Linus Torvalds. All other third party trademarks are the property of their respective owners.

All information found in this book has been compiled with utmost attention to detail. However, this does not guarantee complete accuracy. Neither Novell, Inc., SUSE Linux Products GmbH, the authors, nor the translators shall be held liable for possible errors or the consequences thereof.

Table of Contents

1. Introduction	1
1.1. What is KIWI?	1
1.2. What does KIWI do?	1
1.3. How do I use KIWI?	1
2. Installation	3
2.1. Installing using packages	3
2.2. Installing from source	4
3. Basic Workflow	5
3.1. Introduction	5
3.2. Build Process	7
3.3. Boot image Hook-Scripts	10
3.4. Boot image customization	13
3.5. Using pre-built boot images	14
3.6. Boot Parameters	15
3.7. Common and Distribution Specific Code	16
4. Image Caches	17
4.1. Introduction	17
4.2. Example	19
5. KIWI Image Description	21
5.1. The config.xml File	23
6. Creating Appliances with KIWI	41
6.1. Overview	41
6.2. The KIWI Model	42
6.3. Cross Platform Appliance Build	43
7. Maintenance of Operating System Images	45
8. System to Image Migration	49
8.1. Create a Clean Repository Set First	49
8.2. Watch the Overlay and Unpackaged Files	50
8.3. Checklist	50
8.4. Turn my System Into an Image...	50
9. Installation Source	53
9.1. Adapt the Example's config.xml	53
9.2. Create a Local Installation Source	53
10. ISO Image—Live Systems	55
10.1. Building the suse-live-iso Example	55
10.2. Using the Image	55
10.3. Flavours	55
10.4. USB stick images	56
11. VMX Image—Virtual Disks	59
11.1. Building the suse-vm-guest Example	59
11.2. Using the Image	59
11.3. Flavours	59

12. Linux Container	63
12.1. Building the suse-lxc-guest Example	63
12.2. Using the Image	63
12.3. Image Configuration Details	64
13. PXE Image—Thin Clients	65
13.1. Setting Up the Required Services	65
13.2. Building the suse-pxe-client Example	66
13.3. Using the Image	66
13.4. Flavours	67
13.5. Hardware Grouping	76
14. OEM Image—Preload Systems	83
14.1. Building the suse-oem-preload Example	83
14.2. Using the Image	83
14.3. Flavours	84
15. Xen Image—Paravirtual Systems	87
15.1. Building the suse-xen-guest Example	87
15.2. Using the Image	87
15.3. Flavours	88
16. EC2 Image—Amazon Elastic Compute Cloud	89
16.1. Building the suse-ec2-guest Example	90
16.2. Using EC2 and the created image	91
A. KIWI Man Pages	99
kiwi	100
kiwi::config.sh	107
kiwi::images.sh	111
kiwi::kiwirc	114
Index	115

1 Introduction

Table of Contents

1.1. What is KIWI?	1
1.2. What does KIWI do?	1
1.3. How do I use KIWI?	1

1.1. What is KIWI?

KIWI is an image build system for Linux.

A Linux image may present itself in many different formats, for example the *.iso file you download to burn a distribution installation file to optical media is an image. A file used by virtualization systems such as KVM, Xen, or VMware is an image. The installation of a Linux system on your hard drive can be turned into an image using the **dd** command.

Basically, you can think of an image as a Linux system in a file. Depending on the type of the image you are dealing with you have different options for using the image. For example you can burn an ISO image to optical media and then boot your computer from the CD/DVD, or you can run a Virtual Machine from the *.iso file (image) stored on your hard drive.

1.2. What does KIWI do?

KIWI builds images in a variety of formats.

As an image build tool, KIWI builds images in a relatively large number of supported image formats. The details of the image creation process are explained in the Chapter 3, *Basic Workflow* chapter. The image format of the image produced by KIWI is defined within a configuration file named `config.xml` as described in Chapter 5, *KIWI Image Description*.

Note that not all elements and attributes that may be used in the KIWI `config.xml` configuration file are listed or described in this document. The complete schema documentation can be accessed on the web at <http://doc.opensuse.org/projects/kiwi/schema-doc/>, latest version, or on your local system using the `file:///usr/share/doc/packages/kiwi/schema/kiwi.html` path as the URL in the browser.

1.3. How do I use KIWI?

KIWI is a command line tool that is invoked with the **kiwi** command in your shell. KIWI needs to be executed as the root user, as administrative privileges are required for many operations

that need to take place to create an image. Therefore, when using KIWI you need to be aware of what you are doing and a certain amount of caution is in order. Running KIWI on your system is not inherently dangerous to your system, just keep in mind that you are running as the root user.

An image is created in a two step process as described in the Chapter 3, *Basic Workflow* chapter. Use **kiwi --prepare** for the first step and **kiwi --create** for the second step. For user convenience KIWI also has the **--build** that combines the *prepare* and *create* steps.

Additional introductory information can also be found on the web at http://en.opensuse.org/SDB:KIWI_Cookbook_Start_Cooking.

2 Installation

Table of Contents

2.1. Installing using packages	3
2.2. Installing from source	4

2.1. Installing using packages

Once you have added the appropriate repositories (more on this below) to your system you can search for the kiwi packages through the YaST interface or using **zypper** as shown below.

```
zypper se kiwi
```

The list of packages returned by zypper contains the main package, simply named `kiwi-`, the `-doc` package containing the documentation files, and the `-desc-` packages containing the boot descriptions for the various image types. Installing this set of packages is sufficient to build your images.

Adding repositories to your system can be accomplished using the YaST interface or the **zypper ar** command.

2.1.1. Distribution provided packages

The simplest and most straight forward way to install KIWI is to use the packages that are part of the SUSE distribution you are running. In openSUSE the kiwi packages are part of the "standard" distribution and in SUSE Linux Enterprise kiwi packages are available in the SDK channel.

2.1.2. Packages used by SUSE Studio

If you use SUSE Studio to set up your configuration and then export it to build locally on your machine you want to make sure to use the same version of KIWI that SUSE Studio uses to build images. This version of KIWI which most often differs from the version released with a given distribution is available from the openSUSE Build Service. The repository you want to add to your system is <http://download.opensuse.org/repositories/home:/ctso/DISTRO>.

Once you have the repository added to your system you can search for the kiwi packages and install them as described above.

2.1.3. Packages for development releases

KIWI is under active development and changes almost on a daily basis. The development code is generally released once a week on Friday. Sometimes the development releases contain new bugs that break existing builds. Therefore, this is not necessarily the best release stream to track if you are looking for critical on time builds of already configured appliances. However, tracking this stream provides a great opportunity for you to help in detecting such bugs and by reporting them on the mailing list you can help the developers. Any regression fixes are generally released as soon as they are completed. Thus, there is no need to wait until the next scheduled release on a Friday. If you add the `http://download.opensuse.org/repositories/Virtualization:/Appliances/DISTRO` repository to your system you can track the development release.

Once you have the repository added to your system you can search for the kiwi packages and install them as described above.

2.2. Installing from source

KIWI is developed and maintained in a git repository on GitHub. You can clone the source code using the following command.

```
git clone https://github.com/openSUSE/kiwi.git
```

Before installing from source you want to verify that all the dependencies are satisfied. The best way to accomplish this is to install all packages listed as *BuildRequires* in the `.spec` file found in the `rpm` directory. Once all dependent packages are installed change your working directory to the `kiwi` directory and build and install from source.

```
make
```

```
make install
```

The KIWI self tests are executed using:

```
make test
```

If you want to refresh your source with the latest checked in code you can simply pull the latest sources from the GitHub repository using the command shown below.

```
git pull
```


3 Basic Workflow

Table of Contents

3.1. Introduction	5
3.2. Build Process	7
3.3. Boot image Hook-Scripts	10
3.4. Boot image customization	13
3.5. Using pre-built boot images	14
3.6. Boot Parameters	15
3.7. Common and Distribution Specific Code	16

3.1. Introduction

KIWI creates images in a two step process, as mentioned previously. The first step, the *prepare* operation, generates a so called *unpacked image* tree (directory) using the information provided in the `config.xml` configuration file. The `config.xml` file is part of the *configuration directory (tree)* that describes the image to be created by KIWI. The second step, the *create* operation, creates the *packed image* or *image* in the specified format based on the unpacked image, information provided in the `config.xml` and the *boot image* description specified in the `config.xml` file. Generally the KIWI provided boot image description is sufficient to meet the needs of the image to be created. KIWI also supports the use of custom boot images.

Figure 3.1. Image Creation Architecture

- ❶ Encapsulated system reachable via chroot
- ❷ Encapsulated system reachable via kernel filesystem/extension drivers

Prior to building an image with KIWI it is important to understand the composition of an image, the general concepts of Linux, including the boot process, and distribution concepts such as package management.

Installation of a Linux system generally occurs by booting a target system from an installation source such as an install CD/DVD, a live CD/DVD, or entering the PXE boot environment. The installation process is often driven by an installer that interacts with the user to collect information about the installation. This information generally includes the *software to be installed*, the *timezone*, system *user* data, and other information. Once all the information is collected the installer installs the necessary and specified software onto the target system using packages from the available software sources (repositories). After the installation is complete the system generally reboots and enters a configuration procedure upon startup. The configuration may be fully automatic or it may include user interaction.

A system image, or image, is a *complete installation* of a Linux system in a file. The image represents an operational system and may or may not contain the "final" configuration. The behavior of the image upon deployment varies depending on image type and image configuration. With KIWI it is possible to completely customize the initial start up behavior of the image. This may include behavior that allows the image to simply be deployed inside an existing virtual environment with no required configuration at start up. It is also possible to create images that automatically configure themselves in a known target environment. Further, the startup of an interactive configuration procedure can be integrated into the image to allow the user to configure the image when it is booted for the first time. The image configuration possibilities are practically unlimited. The image creation process with KIWI is automated and does not require any user interaction. The required information for the image creation process is provided in the primary configuration file named `config.xml`. The image can optionally be customized using the `config.sh` and `images.sh` scripts. Additional customization can be accomplished with the use of an optional *overlay tree (directory)* called *root*. The configuration information is stored in the so called *image description or configuration directory (tree)*.

3.2. Build Process

The creation of an image with KIWI is a two step process, the first step is called the *prepare* step and it must complete successfully before the second step, the *create* step can be executed. During the prepare step KIWI creates a new root tree or so called *unpacked image*. The new root tree is created in a directory specified on the command line with the `--root` argument or the value of the `defaultroot` element in the `config.xml` file. This directory will be the target for any software packages to be installed during the image creation process. For package installation KIWI relies on the package manager specified with the `packagemanager` element in the `config.xml` file. KIWI supports the *smart* and *zypper* package managers. The prepare step executes the following major stages:

- **Create Target Root Directory.**

KIWI will exit with an error if the target root tree already exists to prevent accidental deletion of an existing unpacked image. Using the `--force-new-root` command line argument will force kiwi to delete the existing target directory and create a new unpacked image in a new directory with the same name.

- **Install Packages.**

Initially KIWI configures the package manager (zypper by default) to be used for the image creation to use the repositories specified in the configuration file and/or specified on the command line. Following the repository setup the packages specified in the `bootstrap` section are installed in a temporary workspace external to the target root tree. This establishes the initial environment, to support the completion of the process in chroot setting. The essential packages to specify as part of the bootstrap environment are the *filesystem* and *glibc-locale* packages. The dependency chain of these two packages is sufficient to populate the bootstrap environment with all required software to support the installation of packages into the new root tree. The installation of software packages through the selected package manager may install packages that you do not want in your image. Removing undesired packages can be accomplished by specifying the packages you would like to remove from the image as children of a `packages` element where the value of the `type` attribute of the `packages` element is set to `delete`.

- **Apply The Overlay Tree.**

After the package installation with the package manager is complete, KIWI will apply all files and directories present in the overlay directory named *root* inside the configuration

directory to the target root tree. This allows you to over write any file that was installed by one of the packages installed during the installation phase. Files and directories will appear in the unpacked image tree in the same location as they are found in the directory named *root*.

- **Apply Archives.**

Any archives specified with the `archive` element in the `config.xml` file are applied in the specified order (top to bottom) after the overlay tree copy operation is complete. Archives are unpacked at the top level of the new root tree and files will be located according to their path in the archive. As with the overlay tree, it is possible to over write any file in the target root tree.

- **Execute User Defined `config.sh` Script.**

At the end of the preparation stage the optional script named `config.sh` is executed at the root level of the target root tree. The primary intended use of this script is to complete system configuration such as service activation. For detailed description pre-defined configuration functions consult the `kiwi::config.sh(1)` man page.

- **Manage The New Root Tree.**

The unpacked image directory is just a directory, as far as the build system is concerned and you can manipulate the content of this directory to your liking. Further, as this directory represents a system installation you can `chroot` into this directory and run in the `chroot` environment to make changes. However, it is strongly discouraged to apply changes directly to the unpacked root, as any changes you apply will be lost when the *prepare* step for the image is repeated. In addition you may introduce errors into the unpacked root tree that may lead to very difficult to track kiwi build issues during the *create* step of the image build process. The best practice is to apply any necessary changes to the configuration directory followed by a new *prepare* operation. If you inspect the created unpacked root tree you will find a directory named `image` at the top level that you would not find on a system installed with the distribution installer. This directory contains information KIWI requires during the *create* step, including a copy of the `config.xml` file. You can make modifications to data in this directory to influence the *create* step, however, as mentioned previously this is discouraged and changes will be lost once the *prepare* step is repeated.

Successful completion of the *prepare* step is a the pre-requisite for the *create* step of the image build process. With the successful completion of the image preparation the unpacked root tree is considered complete and consistent. Creating the packed, or final image requires the execution of the *create* step. Multiple images can be created using the same unpacked root tree, for example it is possible to create a self installing OEM image and a virtual machine image from one unpacked root tree, under the condition that both image types are specified in the `config.xml` when the *prepare* step is executed. During the *create* step the following major operations are performed by kiwi:

- **Execute User Defined `images.sh` Script.**

At the beginning of the image creation process the optional `images.sh` script is executed at the top level of the unpacked image directory. Unlike the `config.sh` script, the `images.sh` script does not have a target use case. The script is most often used to remove files that are no needed in the final image. For example if an appliance is being built that is targeted for specific hardware one can remove all unnecessary kernel drivers from the image using this script. Consult the `kiwi::images.sh(1)` man page for a detailed description of pre-defined functions available in the `images.sh` script.

- **Create Requested Image Type.**

The image types that can be created from a prepared image tree depend on the types specified in the image description `config.xml` file. The configuration file must contain at least one type element. The figure below shows the currently image types:

Figure 3.2. Image Types




- ❶ Live Image on CD, DVD or USB stick
- ❷ Virtual system that can be used in VMware, Xen, Amazon Cloud, KVM, etc. virtual environments. Depending on the format a guest configuration file is created.
- ❸ Preload system for install media CD/DVD or USB stick
- ❹ Network boot image. KIWI also provides the bootp environment via the package `kiwi-pxeboot`

Detailed information, including step by step instructions about using the **kiwi** command and building specific images, as well as the configuration of the supported image types is provided later.

Most Linux systems use a special boot image to control the system boot process after the system firmware, BIOS or UEFI, hands control of the hardware to the operating system. This boot image is called the *initrd*. The Linux kernel loads the *initrd*, a compressed cpio initial ramdisk, into RAM and executes *init* or, if present, *linuxrc*. KIWI creates the boot image as part of the *create* step in the image build process. Each image type has a specialized image description that describes the boot image. Common functionality is shared between the boot images through a set of functions. The boot image descriptions follow the same principles as the system image descriptions and are provided by KIWI. The boot image descriptions provided by KIWI cover almost all use cases and it should not be necessary for the majority of KIWI users to implement their own boot descriptions.

Figure 3.3. Image Descriptions



- ❶ Boot image descriptions are provided by KIWI, use is recommended but not required
- ❷ The system image description is created by the KIWI user, or a KIWI provided template may be used

The boot image descriptions are stored in the `/usr/share/kiwi/image/*boot` directories. KIWI selects the boot image to build based on the value of the `boot` attribute of the type element. The attribute value is expected in the general form of `boottype/distribution`. For example to select the OEM boot image for openSUSE version 12.1 the value of the `boot` attribute should be `oemboot/suse-12.1`. The boot image description only represent the initrd and as such serves a limited purpose. The system image description created by the person building the image is ultimately the image that runs on the target system. Boot image descriptions are complete and consistent descriptions that allow you to build the boot image outside of the system image build process. The resulting boot image can be stored and re-used as described in the Section 3.5, “Using pre-built boot images” section.

3.3. Boot image Hook-Scripts

All KIWI created boot images contain kiwi boot code that gets executed when the image is booted for the first time. This boot code is different for the various image types and provides hooks to execute user defined custom shell scripts. The shell scripts provided by the user may extend the first boot process and are expected to exist inside the boot image in a specific

location with specific names. The naming and timing of the execution of the hook scripts is image type dependent and described later. The instructions below explain the concepts of hook scripts, which is common to all image types, and how to include the scripts in the initrd.

- All hook scripts must be located in the `kiwi-hooks` directory at the top level of the initrd. The best approach to including the hook scripts in the initrd is to create an archive of a `kiwi-hooks` directory that contains the custom boot scripts.

```
mkdir kiwi-hooks
--> place all scripts inside kiwi-hooks
tar -cf kiwi-hooks.tgz kiwi-hooks/
```

The tarball must be located at the top level of the image description directory, this is the same level that contains the `config.xml` file.

- Hook scripts are executed using a predetermined name that is hard coded into the kiwi boot code. This name is extended using the `.sh` extension and differs by boot image type. Therefore, the boot script naming in the archive must be exact. Boot scripts are sourced in the kiwi boot code. This provides the hook script access to all variables set in the boot environment. This also implies that no separate shell process is started and the boot scripts do not have to have the executable bit set. Encoding the interpreter location with the `#!` comment is superfluous.
- Hook scripts are only executed from within kiwi's boot code and must therefore be part of the KIWI created boot image. Including the content of a tarball in the initrd is accomplished by setting the value of the `bootinclude` attribute of the archive element to `true` in the `config.xml` file as shown below:

```
<packages type="image">
  <archive name="kiwi-hooks.tgz" bootinclude="true"/>
</packages>
```

The concept of including an archive in the boot image follows the same concepts described for the system image previously. The setting in the system image description will have no effect if a pre-built boot image is being used. In order to use an archive in a pre-built boot image the archive must be part of the boot image description in which case it is not necessary to set the `bootinclude` attribute.

The following list provides information about the hook names, timing of the execution, and the applicable boot image.

- **init.** This hook is called before `udev` is started. The hook exists only for the `pxe` image type.
- **preconfig|postconfig.** The hooks are called before and after the client configuration files (`CONF` contents) are setup, respectively. The hooks exist only for the `pxe` image type.
- **predownload|postdownload.** The hooks are called before and after the client image receives the root filesystem, respectively. The hooks exist only for the `pxe` image type.
- **preImageDump|postImageDump.** The hooks are called before and after the install image is dumped on the target disk, respectively. The hooks exist only for the `oem` image type.
- **preLoadConfiguration|postLoadConfiguration.** The hooks are called before and after the client configuration file `config.MAC` is loaded, respectively. The hooks exist only for the `pxe` image type.

- **premount|postmount.** The hooks are called before and after the client root filesystem is mounted, respectively. The hooks exist only for the *pxe* image type.
- **prenetwork|postnetwork.** The hooks are called before and after the client network is setup, respectively. The hooks exist only for the *pxe* image type.
- **prepartition|postpartition.** The hooks are called before and after the client creates the partition table on the target disk, respectively. The hooks exist only for the *pxe* image type.
- **preprobe|postprobe.** The hooks are called before and after the loading of modules not handled by udev, respectively. The hooks exist only for the *pxe* image type.
- **preswap|postswap.** The hooks are called before and after the creation of the swap space, respectively. The hooks exist only for the *pxe* image type.
- **preactivate.** This hook is called before the root filesystem is moved to / The hook exists only for the *pxe* image type.
- **preCallInit.** This hook is called in before the initialization process, init or systemd, is started. At call time the root filesystem has already been moved to /. The hook exists only for the *oem* and *vmx* image types.
- **preException.** This hook is called before a system error is handled, the actual error message is passed as parameter. This hook can be used for all image types.
- **preHWdetect.** This hook is called before the install image boot code searches for the possible target storage device(s). The hook exists only for the *oem* image type.
- **preNetworkRelease.** This hook is called before the network connection is released. The hook exists only for the *pxe* image type.

The execution of hooks can be globally deactivated by passing the following variable to the kernel commandline:

```
KIWI_FORBID_HOOKS=1
```

In addition to the hook script itself it's also possible to run a post command after the hook script was called. This allows to run commands tied to a hook script without changing the initrd and thus provides a certain flexibility when writing the hook. The post command execution is based on variables one can pass to the kernel commandline to extend an existing hook script. There are the following rules for the processing of these information

- **The hook must activate the command post processing.** Post hook commands are only processed if the corresponding hook script activates this. The variable the hook script has to set follows the naming schema: `KIWI_ALLOW_HOOK_CMD_|HOOKNAME| = 1` For example:

```
KIWI_ALLOW_HOOK_CMD_preHWdetect=1
```

If this is set as part of the `preHWdetect.sh` hook script code the post command execution is activated

- **KIWI_HOOK_CMD_|HOOKNAME|.** The variable containing the command to become executed must match the following naming schema. For example:

```
KIWI_HOOK_CMD_preHWdetect="ls -l"
```


This would cause the preHWdetect hook to call `ls -l` at the end of the hook script code

- **KIWI_FORBID_HOOK_CMDS.** If this variable is set to something non empty the post hook command execution is deactivated however the basic hook script invocation is still active unless `KIWI_FORBID_HOOKS` is also set

3.4. Boot image customization

The KIWI provided boot image descriptions should satisfy the requirements for a majority of image builds and the environments in which these images are deployed. For the circumstances that require customized boot images KIWI provides mechanisms in the system image `config.xml` file to influence the boot image content. Using these mechanisms allows the user to still base the boot image on the KIWI provided descriptions rather than defining a completely new and custom boot image description. Creating a custom boot image that is not based on the KIWI provided descriptions is also possible. The following question and answer section provides solutions to the most common customization needs from the initrd created by kiwi.

- **Why is the boot image so big and can I reduce it's size ?** KIWI includes all required tools and libraries to boot the image in all circumstances in the target environment for the image type. If target environment is well defined it is possible to remove data from that is known not to be needed. This will decrease the size of the initrd to and decrease boot time. Removing files in the boot image is accomplished by adding a `strip` section to the system image `config.xml` file, with the `type` attribute set to `delete`, as shown below.

```
<strip type="delete"/>
  <file name="..." />
</strip>
```

Removing files that are needed my result in an image that cannot be booted.

- **Can drivers be added to the boot image?** KIWI uses a subset of the kernel. Should you encounter problems due to a missing driver that is part of the "standard" kernel but has not been included by the kiwi build process you can add the driver by adding a `drivers` section to the system image `config.xml` file, as shown below.

```
<drivers>
  <file name="drivers/..." />
</drivers>
```

If the driver is provided by a package, the package itself needs to be specified as part of the `image` package section and it must be marked for boot image inclusion by setting the value of the `bootinclude` attribute of the package element to `true`, as shown below.

```
<packages type="image"/>
  <package name="..." bootinclude="true" />
</packages>
```

- **How to add missing tools/libraries?** Additional software can be added to the boot image with the use of the `bootinclude` attribute of the package element or the archive element. At the end of the boot image creation process kiwi attempts to reduce the size of the boot image by removing files that are not part of a known list of required files, any detectable dependencies of the files listed are preserved as well. The list of known required files is hard coded in the `/usr/share/kiwi/modules/KIWIConfig.txt` file. If you added

files to the boot image that are needed in your specific use case you need to instruct kiwi to not strip the files you have added to the boot image. This is accomplished by adding a `strip` section to the system image `config.xml` file, with the `type` attribute set to `tools`, as shown below.

```
<strip type="tools"/>
  <file name="..." />
</strip>
```

the removal/preservation of files is name base and the path is immaterial. Therefore, you only have to specify the file name that is to be preserved.

- **Is it possible to add boot code?** Yes, as described in the Section 3.3, “Boot image Hook-Scripts” section above, KIWI supports the execution of boot code at various times for various image types using *hook* scripts.
- **Is it possible to include completely custom boot code?** No. In cases where the provided hooks are insufficient and the KIWI provided boot code needs to be replaced completed it is necessary to create a custom boot image description. In this case, all parts of the boot image description must be created by the user. It is best to use one of the KIWI provided boot descriptions as a template.

3.5. Using pre-built boot images

During the create step of the KIWI image building process kiwi, creates the so called boot image, as described previously, based on the specified boot image description in the `config.xml` configuration file. This creation process takes time and can be short circuited by using pre-built boot images.

As described earlier, the KIWI provided boot images can be found in the `/usr/share/kiwi/image/*boot` directories. Located within the `*boot` directories are boot image description trees named for the applicable distribution. For example the `oemboot/suse-SLES11` directory is the boot image description for an OEM image for SUSE Linux Enterprise Server 11. The boot image configuration trees are complete image descriptions, very similar in nature to the system image descriptions created most commonly for image building, that kiwi uses to create the boot image during the system image creation process. Therefore, it is possible to build these boot images outside of the system image build process. The result of a build of one of the boot image descriptions is a pre-built boot image that can be used in many image builds for the same distribution and type. The following commands show the creation of a pre-built boot image for openSUSE 12.1 for the OEM image type.

```
kiwi --prepare /usr/share/kiwi/image/oemboot/suse-12.1 --root /tmp/oem121_initunpacked
```

```
kiwi --create /tmp/oem121_initunpacked -d /mystore/kiwiprebuiltboot
```

The commands above result in the creation of the OEM boot image for openSUSE 12.1 in the directory `/mystore/kiwiprebuiltboot`. This boot image can readily be used by any kiwi build for an openSUSE 12.1 OEM image. Using the pre-built image requires that the value of the `checkprebuilt` attribute of the `type` element be set to `true` and that the location of the boot image is provided with the `--prebuiltbootimage` command line argument, or the `defaultprebuilt` element in the `config.xml` file.

Using pre-built boot images has the advantage that the boot image does not have to be re-created every time a specific image type for a given distribution is rebuilt. Additionally, this

process provides a convenient way to maintain customized boot images. One disadvantage to the use of pre-built images is that it is not possible to integrate the latest updates of tools that are part of the initrd in the image as the pre-built boot image will contain only the latest versions available in the specified repositories on the build date. However, in most cases this does not represent a concern/issue as the initrd in the image generally gets replaced once the image is deployed.

3.6. Boot Parameters

A KIWI created initrd based on one of the KIWI provided boot image descriptions recognizes kernel parameters that are useful for debugging purposes, should the image not boot. These parameters may not work if the image contains a custom boot image where the kiwi boot code has been replaced, and the parameters are not recognized after the initial KIWI created initrd has been replaced by the "regular" distribution created initrd after the initial boot of the image.

- ***kiwidebug=1***. If the boot process encounters a fatal error, the default behavior is to re-boot the system 120 seconds. The “exception” behavior is changed by setting the *kiwidebug* parameter. With the value of the parameter set to 1 the system will enter a limited shell environment should a fatal error occur during boot. The shell contains the standard basic commands. The `/var/log/kiwi.boot` boot log file may be consulted to develop a better understanding of the boot failure. In addition to the spawned shell process kiwi also starts the dropbear ssh server if the environment is suitable. Support for ssh into the boot image is possible in the netboot and oemboot (in PXE boot mode) boot images. For isoboot and vmxboot boot images there is no remote login support because they don't setup a network. In order to have dropbear installed as part of the boot image the following needs to be added to the system image configuration:

```
<packages type="image"/>
  <package name="dropbear" bootinclude="true"/>
</packages>
```

It's required that the repo setup provides dropbear. Once dropbear is there the kiwi boot code will start the service. In order to access the boot image via ssh it's required to provide a public key on the pxe server in the directory: `server-root/KIWI/debug_ssh.pub`. kiwi only searches for that filename so it's required to name it “debug_ssh.pub”. Adding more than one public key to this file is possible exactly like the common SSH file “authorized_keys”. The path “server-root” depends on what server type was configured to download the image. By default this is done via tftp. In that case the complete path to put the public key to is `/srv/tftpboot/KIWI/debug_ssh.pub`. on the pxe server. If ftp or http is used the server-root path is different. If a public key was found you can login as follows:

```
ssh root@<ip>
```

It might be useful to have a copy tool like scp or rsync as part of the boot image as well. Adding rsync as bootincluded package does not increase the size of the initrd much and would allow to extract e.g the kiwi boot log as follows:

```
RSYNC_RSH='ssh -l root'
rsync -avz <ip>:/var/log/boot.kiwi .
```

- ***kiwistderr=/dev/...*** During boot, the kiwi boot code writes messages to tty1 and tty3. The tty1 messages are high-level summary messages, whereas the shell debug output messages, which may also contain error information, are written to tty3. With the *kiwistderr*

parameter one can combine both message streams and specify the device the messages should be written to. It is common to set `/dev/console` as an alternative target and change the default logging behavior.

3.7. Common and Distribution Specific Code

KIWI is designed to be in principal distribution independent and the majority of the kiwi implementation follows this design principal. However, Linux distributions differ from each other, primarily in the package management area as well as the creation and composition of the boot image.

Within the KIWI code base major areas of Linux distribution differences are isolated into specific regions of the code. The remainder of the code is common and distribution independent.

KIWI provided functions that are distribution specific contain the distribution name as a prefix, such as `suseStripKernel`. Scripts that are part of the boot code and are distribution specific are identified by a prefix of the distribution name followed by a “-”, **suse-linuxrc** for example. When kiwi creates a boot image for a SUSE distribution the **suse-linuxrc** file from the boot discription is used as the **linuxrc** file that the Linux kernel calls.

With this design and implementation it is possible to maintain distribution specific code in the same project while also providing explicit hints to the user when distribution specific code is being used. The implemented SUSE specific code can be used as a guideline to support other distributions.

4 Image Caches

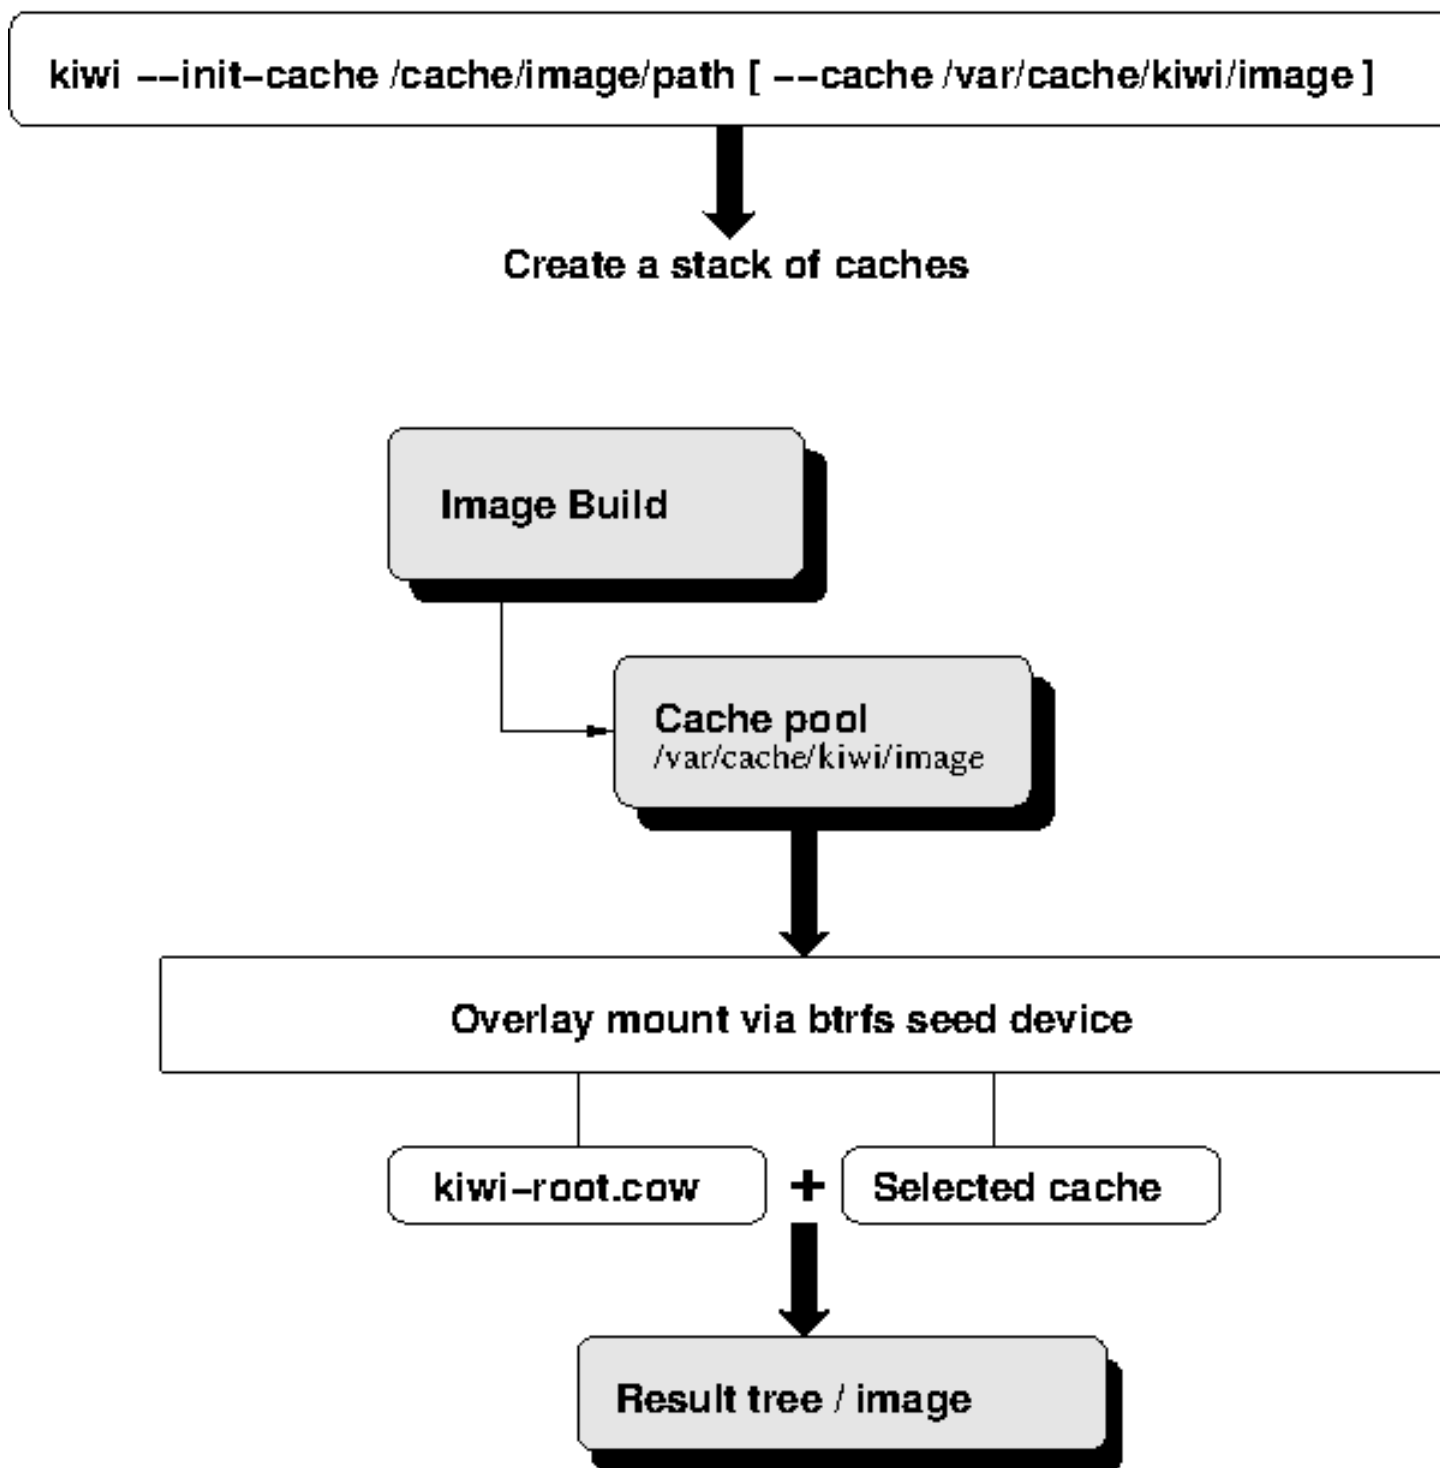
Table of Contents

4.1. Introduction	17
4.2. Example	19

4.1. Introduction

The process of creating an appliance could take quite some time and often the same software is installed over and over again. In order to speed up that process kiwi is able to create and re-use so called image caches. An image cache in kiwi is an ext3 based filesystem image representing a pre-installed software bundle.

Figure 4.1. Image Caching Architecture



Before one can use a cache it needs to be created. A cache can be created from any standard kiwi image description. That means you can simply use one of the template descriptions and create a cache from it. But it might be more clever to create image descriptions for the purpose of caching. Such descriptions could represent a set of patterns for example. The less special a cache is the more often it can be re-used

Once there are caches in the system kiwi selects the best match and mounts the cache in a way that all write actions (copy-on-write cache) are redirected to a cow file. That way the

cache itself is never changed and can be re-used simultaneously for other build processes. As result the build process doesn't start with an empty tree but with a tree almost complete. Only the missing parts are now installed and according to how much the cache already covered this process is done fastly

4.2. Example

Let's say we know that we want to build some images of type 'vmx' and based on the openSUSE 12.2 JeOS image description. In order to create image caches for the system and the boot image the following steps needs to be done:

- 1) copy the base image descriptions used in the build:

```
cp -a /usr/share/kiwi/image/vmxboot/suse-12.2 /tmp/boot-cache
cp -a /usr/share/kiwi/image/suse-12.2-JeOS /tmp/image-cache
```
- 2) modify boot-cache and image-cache config.xml files to add a btrfs type and size

```
<type image="btrfs">
  <size unit="G">5</size>
</type>
```
- 3) modify boot-cache and image-cache to contain the packagemanager. This is required for the later use of the caches

```
<package name="zypper"/>
```
- 4) now build the caches:

```
kiwi --init-cache /tmp/image-cache
kiwi --init-cache /tmp/boot-cache
```

By default those caches will be created in '/var/cache/kiwi/image'. To run a build which makes use of the cache the following command is used:

```
kiwi --build suse-12.2-JeOS -d /tmp/myimage --type vmx \
  --cache /var/cache/kiwi/image
```

This call is about 50% faster compared to the creation without a cache. It's important to understand that a cache based build will create a root tree which contains only the cow file containing the differences compared to the used cache. So at any time you want to create an image out of it you have to make sure that the cache also exists

5 KIWI Image Description

Table of Contents

5.1. The config.xml File	23
--------------------------------	----

In order to be able to create an image with KIWI, a so called image description must be created. The image description is represented by a directory which has to contain at least one file named `config.xml` or `*.kiwi`. A good start for such a description can be found in the examples provided in `/usr/share/doc/packages/kiwi/examples`.

Figure 5.1. Image Description Directory



The following additional information is optional for the process of building an image, but most often mandatory for the functionality of the created operating system:

`images.sh`

Optional configuration script while creating the packed image. This script is called at the beginning of the image creation process. It is designed to clean-up the image system. Affected are all the programs and files only needed while the unpacked image exists.

`config.sh`

Optional configuration script while creating the unpacked image. This script is called at the end of the installation, but *before* the package scripts have run. It is designed to configure

the image system, such as the activation or deactivation of certain services (insserv). The call is not made until after the switch to the image has been made with chroot.

root

Subdirectory that contains special files, directories, and scripts for adapting the image environment *after* the installation of all the image packages. The entire directory is copied into the root of the image tree using **cp -a**.

config-yast-firstboot.xml

Configuration file for the control of the YaST firstboot service. Similar to the AutoYaST approach, YaST also provides a boot time service called firstboot. Unfortunately there is no GUI available to setup the firstboot, but good documentation in `/usr/share/doc/packages/yast2-firstboot`. Once you have created such a firstboot file in your image description directory, KIWI will process the file and setup your image as follows:

1. KIWI enables the firstboot service.
2. While booting the image, YaST is started in firstboot mode.
3. The firstboot service handles the instructions listed in the file `config-yast-firstboot.xml`.
4. If the process finished successfully, the environment is cleaned and firstboot will not be called at next reboot.

config-yast-autoyast.xml

Configuration file which has been created by AutoYaST. To be able to create such an AutoYaST profile, run:

```
yast2 autoyast
```

Once you have saved the information from the AutoYaST UI as `config-yast-autoyast.xml` file in your image description directory KIWI will process on the file and setup your image as follows:

1. While booting the image YaST is started in AutoYaST mode automatically
2. The AutoYaST description is parsed and the instructions are handled by YaST. In other words the *system configuration* is performed
3. If the process finished successfully the environment is cleaned and AutoYaST won't be called at next reboot.

config-cdroot.tgz

Archive which is used for ISO images only. The data in the archive is uncompressed and stored in the CD/DVD root directory. This archive can be used, for example, to integrate a license file or information directly readable from the CD or DVD.

config-cdroot.sh

Along with the `config-cdroot.tgz` one can provide a script which allows to manipulate the extracted data.

config/

Optional subdirectory that contains Bash scripts that are called after the installation of all the image packages, primarily in order to remove the parts of a package that are not

needed for the operating system. The name of the Bash script must resemble the package name listed in the config.xml.

5.1. The config.xml File

The mandatory image definition file is divided into different sections which describes information like the image name and type as well as the packages and patterns the image should consist of.

The following information explains the basic structure of the XML document. When KIWI is executed, the XML structure is validated by the KIWI RELAX NG based schema. For details on attributes and values please refer to the schema documentation file at `/usr/share/doc/packages/kiwi/kiwi.rng.html`.

5.1.1. image Element

```
<image schemaversion="5.2" name="iname"
  displayname="text"
  inherit="path"
  kiwirevision="number"
  id="10 digit number">
  <!-- ... -->
</image>
```

The image definition starts with an image tag and requires the schema format at version 2.0. The attribute `name` specifies the name of the image which is also used for the filenames created by KIWI. Because we don't want spaces in filenames the `name` attribute must not have any spaces in its name.

The following optional attributes can be inserted in the image tag:

`displayname`

allows setup of the boot menu title for isolinux and grub. So you can have *suse-SLED-foo* as the image name but something like *my cool Image* as the boot display name.

`inherit`

inherits the packages information from another image description

`kiwirevision`

specifies a KIWI git revision number which is known to build a working image from this description. If the KIWI git revision doesn't match the specified value, the process will exit. The currently used git revision can be queried by calling `kiwi --version`.

`id`

sets an identification number which appears as file `/etc/ImageID` within the image.

Inside the image section the following mandatory and optional subelements exists. The simplest image description must define the elements description, preferences, repository and packages (at least one of `type="bootstrap"`).

5.1.2. description Element

```
<description type="system">
  <author>an author</author>
```

```
<contact>mail</contact>
<specification>short info</specification>
</description>
```

The mandatory `description` section contains information about the creator of this image description. The attribute `type` could be either of the value system which indicates this is a system image description or at value boot for boot image descriptions.

5.1.3. profiles Element

```
<profiles>
  <profile name="name" description="text"/>
  <!-- ... -->
</profiles>
```

The optional `profiles` section lets you maintain one image description while allowing for variation of the sections `packages` and `drivers` that are included. A separate profile element must be specified for each variation. The profile child element, which has `name` and `description` attributes, specifies an alias name used to mark sections as belonging to a profile, and a short description explaining what this profile does.

To mark a set of packages/drivers as belonging to a profile, simply annotate them with the `profiles` attribute. It is also possible to mark sections as belonging to multiple profiles by separating the names in the `profiles` attribute with a comma. If a `packages` or `drivers` tag does not have a `profiles` attribute, it is assumed to be present for all profiles.

5.1.4. preferences Element

```
<preferences profiles="name">
  <version>1.1.2</version>
  <packagemanager>zypper</packagemanager>
  <type image="name" ...>
    <ec2config|systemdisk|oemconfig|pxedeploy|size|split|machine>
  </type>
</preferences>
```

The mandatory `preferences` section contains information about the supported image type(s), the used package manager, the version of this image, and optional attributes. The image version must be a three-part version number of the format: **Major.Minor.Release**. In case of changes to the image description the following rules should apply:

- For smaller image modifications that do not add or remove any new packages, only the release number is incremented. The `config.xml` file remains unchanged.
- For image changes that involve the addition or removal of packages the minor number is incremented and the release number is reset.
- For image changes that change the size of the image file the major number is incremented.

By default, KIWI uses the **zypper** package manager but it is also possible to use the non SUSE native package manager called **smart**.

In general the specification of one `preferences` section is sufficient. However, it's possible to specify multiple `preferences` sections and distinguish between the sections via the `profiles` attribute. Data may also be shared between different profiles. Using profiles it is possible to, for example, configure specific preferences for OEM image generation. Activation of a given

preferences during image generation is triggered by the use of the `--add-profile` command line argument.

For each preferences block at least one type element must be defined. It is possible to specify multiple type elements in any preferences block. To set a given type description as the default image use the boolean attribute `primary` and set its value to `true`. The image type to be created is determined by the value of the `image` attribute. The following list describes the supported types and possible values of the image attribute:

`image="cpio"`

Use the cpio image type to specify the generation of a boot image (initrd). When generating a boot image, it is possible to specify a specific boot profile and boot kernel using the optional `bootprofile="default"` and `bootkernel="std"` attributes.

A boot image should group the various supported kernels into profiles. If the user chooses not to use the profiles supplied by KIWI, it is required that one profile named `std` be created. This profile will be used if no other bootkernel is specified. Further it is required to create a profile named `default`. This profile is used when no bootprofile is specified.

It is recommended that special configurations that omit drivers, use special drivers and/or special packages be specified as profiles.

The bootprofile and bootkernel attribute are respected within the definition of a system image. Use the attribute and value `type="system"` of the description element to specify the creation of a system image. The values of the bootprofile and bootkernel attributes are used by KIWI when generating the boot image.

`image="iso"`

Specify the key-value pair `image="iso"` to generate a live system suitable for deployment on optical media (CD or DVD). Use the `boot="isoboot/suse-*` attribute when generating this image type to select the appropriate boot image for optical media. In addition the optional `flags` attribute may be set to the following values with the effects described below:

`clirc`

Creates a fuse based compressed read-only filesystem which allows write operations into a cow file. If the iso is bigger than 4G you can use the `clirc_udf` flag instead

`seed`

Creates a btrfs based compressed read-only filesystem which allows write operations into a btrfs seed device.

`compressed`

Creates a split ext3 plus squashfs filesystem and combines them via a symlink system to a complete system it is recommended to specify a `split` section as a child of this type element.

If the flags attribute is not used the filesystem will be squashfs compressed for `/bin /boot /lib /lib64 /opt /sbin` and `/usr`. The rest of the filesystem is packed into a tmpfs and linked via symbolic links

`image="oem"`

Use this type to create a virtual disk system suitable in a preload setting. In addition specify the attributes `filesystem`, and `boot="oemboot/suse-*` to control the filesystem used for the virtual and to specify the proper boot image. Using the optional `format`

attribute and setting, the value to iso or usb will create self installing images suitable for optical media or a USB stick, respectively. Booting from the media will deploy the OEM preload image onto the selected storage device of the system. It is also possible to configure the system to use logical volumes. Use the optional `lvm` attribute and specify the logical volume configuration with the `systemdisk` child element. The default volume group name is `kiwiVG`. Further configuration of the image is performed using the appropriate `*config` child block.

`image="pxe"`

Creating a network boot image is supported by KIWI with the `image="pxe"` type. When specifying the creation of a network boot image use the `filesystem` and `boot="netboot/suse-"` attributes to specify the filesystem of the image and the proper boot image. To compress the image file set the `compressed` boolean attribute to `true`. This setting will compress the image file and has no influence on the filesystem used within the image. The compression is often use to support better transfer times when the pxe image is pushed to the boot server over a network connection. The pxe image layout is controlled by using the `pxedeploy` child element.

`image="split"`

The split image support allows the creation of an image as split files. Using this technique one can assign different filesystems and different read-write properties to the different sections of the image. The `oem`, `pxe`, `usb`, and `vmx` types can be created as a split system image. Use the `boot="oem|netboot|usb|vmx/suse-"` attribute to select the underlying type of the split image. The attributes `fsreadwrite`, `fsreadonly` are used to control the read-write properties of the filesystem specified as the attributes value. Use the appropriate `*config` child block to specify the properties of the underlying image. For example when building a OEM based split image use the `oemconfig` child section.

`image="vmx"`

Creation of a virtual disk system is enabled with the `vmx` value of the image attribute. Set the filesystem of the virtual disk with the `filesystem` attribute and select the appropriate boot image by setting `boot="vmxboot/suse-"` The optional format attribute is used to specify one of the virtualization formats supported by QEMU, such as `vmdk` (also the VMware format) or `qcow2`. For the virtual disk image the optional `vga` attribute may be used to configure the kernel framebuffer device. Acceptable values can be found in the Linux kernel documentation for the framebuffer device (see `Documentation/fb/vesafb.txt`). KIWI also supports the selection of the bootloader for the virtual disk according to the rules indicated for the USB system. Last but not least the virtual disk system may also be created with a LVM based layout by using the `lvm` attribute. The previously indicated rules apply. Use the `machine` child element to specify appropriate configuration of the virtual disk system.

All of the mentioned types can specify the `boot` attribute which tells KIWI to call itself to build the requested boot image (`initrd`). It is possible to tell KIWI to check for an already built boot image which is a so called *prebuilt boot image*. To activate searching for an appropriate prebuilt boot image the type section also provides the attribute `checkprebuilt="true|false"`. If specified KIWI will search for a prebuilt boot image in a directory named `/usr/share/kiwi/image/*boot/*-prebuilt`. Example: If the `boot` attribute was set to `isoboot/suse-10.3` and `checkprebuilt` is set to `true` KIWI will search the prebuilt boot image in `/usr/share/kiwi/image/isoboot/suse-10.3-prebuilt`. The directory KIWI searches for the prebuilt boot images can also be specified at the commandline with the `--prebuiltbootimage` parameter.

Within the preferences section, there are the following optional elements:

showlicense

Specifies the base name of a license file which is displayed in oem images before the installation happens. It's possible to add more showlicense sections to display more licenses one after the other. If no such element is specified the default 'license' and 'EULA' files are searched. The search algorithm will append the .txt or .locale.txt suffix to the license name to form the license file name. You should make sure that your license files contain this suffix.

rpm-check-signatures

Specifies whether RPM should check the package signature or not

rpm-excludedocs

Specifies whether RPM should skip installing package documentation

rpm-force

Specifies whether RPM should be called with --force

keytable

Specifies the name of the console keymap to use. The value corresponds to a map file in /usr/share/kbd/keymaps. The KEYTABLE variable in /etc/sysconfig/keyboard file is set according to the keyboard mapping.

timezone

Specifies the time zone. Available time zones are located in the /usr/share/zoneinfo directory. Specify the attribute value relative to /usr/share/zoneinfo. For example, specify Europe/Berlin for /usr/share/zoneinfo/Europe/Berlin. KIWI uses this value to configure the timezone in /etc/localtime for the image.

locale

Specifies the name of the UTF-8 locale to use, which defines the contents of the LC_LANG system environment variable in /etc/sysconfig/language. Please note only UTF-8 locales are supported here which also means that the encoding must *not* be part of the locale information. The KIWI schema validates the locale string according to the following pattern: `[a-z]{2}_[A-Z]{2}([a-z]{2}_[A-Z]{2})*`. This means you have to specify the locale like the following example: `en_US` or `en_US,de_DE`

bootsplash-theme

Specifies the name of the bootsplash theme to use

bootloader-theme

Specifies the name of the gfxboot theme to use

defaultdestination

Used if the --destdir option is not specified when calling KIWI

defaultroot

Used if the option --root is not specified when calling KIWI

kernelcmdline

Specifies additional kernel parameters. The following example disables kernel messages: `kernelcmdline="quiet"`

The type element may contain child elements to provide specific configuration information for the given type. The following lists the supported child elements:

ec2config

The optional ec2config block is used to specify information relevant only to AWS EC2 images. The following information can be provided:

```
<ec2config>
  <ec2accountnr> Your AWS account number </ec2accountnr>
  <ec2certfile> Path to the AWS cert-*.pem file </ec2certfile>
  <ec2privatekeyfile> Path to the AWS pk-*.pem file </ec2privatekeyfile>
</ec2config>
```

systemdisk

Using the optional systemdisk section it is possible to create a LVM (Logical Volume Management) based storage layout. By default, the volume group is named *kiwiVG*. It is possible to change the name of the group by setting the [name](#) attribute to the desired name. Individual volumes within the volume group are specified using the volume element.

The following example shows the creation of a volume named *usr* and a volume named *var* inside the volume group systemVG.

```
<systemdisk name="systemVG">
  <volume name="usr" freespace="100M"/>
  <volume name="var" size="200M"/>
</systemdisk>
```

The optional attribute [freespace](#) controls the amount of unused space available after software has been installed in the given volume. By default the available space of a created volume is between 10% and 20%. Using the optional [size](#) attribute the absolute size of the given volume is specified. The [size](#) attribute takes precedence over the [freespace](#) attribute. If the specified size is insufficient, based on the estimated software install size for the given volume, the specified value will be ignored and a volume with default settings will be created. This implies that the volume will be 80% to 90% full.

oemconfig

By default, the oemboot process will create or modify a swap, and / partition. It is possible to influence the behavior by the oem-* elements explained below. KIWI uses this information to create the file /config.oempartition as part of the automatically created oemboot boot image. The format of the file is a simple key=value format and created by the **KIWIConfig.sh** function named baseSetupOEMPartition.

```
<oemconfig>
  <oem-systemsize>2000</oem-systemsize>
  <oem-... >
</oemconfig>
```

[<oem-align-partition>>true|false</oem-align-partition>](#)

Kiwi attempts to align the start sector of the disk partition on a 4K boundary.

[<oem-boot-title>text</oem-boot-title>](#)

By default, the string OEM will be used as the boot manager menu entry when KIWI creates the GRUB configuration during deployment. The oem-boot-title element allows you to set a custom name for the grub menu entry. This value is represented by the OEM_BOOT_TITLE variable in config.oempartition.

[<oem-bootwait>true|false</oem-bootwait>](#)

Specify if the system should wait for user interaction priot to continuing the boot process after the oem image has been dumped to the designated storage device (default value is false). This value is represented by the OEM_BOOTWAIT variable in config.oempartition.

`<oem-inplace-recovery>true|false</oem-inplace-recovery>`

Specify if the recovery archive is stored as part of the image or if the archive is to be created at the time the image is deployed to the target storage device. OEM_RECOVERY_INPLACE variable in config.oempartition.

`<oem-kiwi-initrd>true|false</oem-kiwi-initrd>`

If this element is set to true (default value is false) the oemboot boot image (initrd) will *not* be replaced by the system (mkinitrd) created initrd. This option is useful when the system is installed on removable storage such as a USB stick or a portable external drive. For movable devices it is potentially necessary to detect the storage location during every boot. This detection process is part of the oemboot boot image. This value is represented by the OEM_KIWI_INITRD variable in config.oempartition.

`<oem-partition-install>true|false</oem-partition-install>`

Specify if the image is to be installed into a free partition on the target storage device. By default the value is false and Kiwi installs images to a target device which causes data loss on the device. With oem-partition-install set to true any other settings that have influence on the partition table, such as oem-swap are ignored. This value is represented by the OEM_PARTITION_INSTALL variable in config.oempartition.

`<oem-reboot>true|false</oem-reboot>`

Specify if the system is to be rebooted after the oem image has been deployed to the designated storage device (default value is false). This value is represented by the OEM_REBOOT variable in config.oempartition.

`<oem-reboot-interactive>true|false</oem-reboot-interactive>`

Specify if the system is to be rebooted after the oem image has been deployed to the designated storage device (default value is false). Prior to reboot a message is posted and must be acknowledged by the user in order for the system to reboot. This value is represented by the OEM_REBOOT_INTERACTIVE variable in config.oempartition.

`<oem-recovery>true|false</oem-recovery>`

If this element is set to true (default value is false), KIWI will create a recovery archive from the prepared root tree. The archive will appear as /recovery.tar.bz2 in the image file. During first boot of the image a single recovery partition will be created and the recovery archive will be moved to the recovery partition. An additional boot menu entry is created that when selected restores the original root tree on the system. The user information on the /home partition or in the /home directory is not affected by the recovery process. This value is represented by the OEM_RECOVERY variable in config.oempartition.

`<oem-recoveryID>partition-id</oem-recoveryID>`

Specify the partition type for the recovery partition. The default is to create a Linux partition (id = 83). This value is represented by the OEM_RECOVERY_ID variable in config.oempartition.

`<oem-silent-boot>true|false</oem-silent-boot>`

Specify if the system should boot in silent mode after the oem image has been deployed to the designated storage device (default value is false). This value is represented by the OEM_SILENTBOOT variable in config.oempartition.

`<oem-shutdown>true|false</oem-shutdown>`

Specify if the system is to be powered down after the oem image has been deployed to the designated storage device (default value is false). This value is represented by the OEM_SHUTDOWN variable in config.oempartition.

`<oem-shutdown-interactive>true|false</oem-shutdown-interactive>`

Specify if the system is to be powered down after the oem image has been deployed to the designated storage device (default value is false). Prior to shutdown a message is posted and must be acknowledged by the user in order for the system to power off. This value is represented by the OEM_SHUTDOWN_INTERACTIVE variable in config.oempartition.

`<oem-swap>true|false</oem-swap>`

Specify if a swap partition should be created. The creation of a swap partition is the default behavior. This value is represented by the OEM_WITHOUTSWAP variable in config.oempartition.

`<oem-swapsize>number in MB</oem-swapsize>`

Set the size of the swap partition. If a swap partition is to be created and the size of the swap partition is not specified with this optional element, KIWI will calculate the size of the swap partition and create a swap partition equal to two times the RAM installed on the system at initial boot time. This value is represented by the OEM_SWAPSIZE variable in config.oempartition.

`<oem-systemsize>number in MB</oem-systemsize>`

Set the size of the root partition. This value is represented by the variable OEM_SYSTEMSIZE in config.oempartition.

`<oem-unattended>true|false</oem-unattended>`

The installation of the image to the target system occurs automatically without requiring user interaction. If multiple possible target devices are discovered the image is deployed to the first device. OEM_UNATTENDED in config.oempartition.

pxedeploy

Information contained in the optional pxedeploy section is only considered if the `image` attribute of the type element is set to `pxe`. In order to use a PXE image it is necessary to create a network boot infrastructure. Creation of the network boot infrastructure is simplified by the KIWI provided package kiwi-pxeboot. This package configures the basic PXE boot environment as expected by KIWI pxe images. The kiwi-pxeboot package creates a directory structure in /srv/tftpboot. Files created by the KIWI create step need to be copied to the /srv/tftpboot directory structure. For additional details about the PXE image please refer to the PXE Image chapter later in this document.

In addition to the image files it is necessary that information be provided about the client setup. This information, such as the image to be used or the partitioning, is contained in a file with the name config.MAC in the directory /srv/tftpboot/KIWI. The content of this file is created automatically by KIWI if the pxedeploy section is provided in the image description. A pxedeploy section is outlined below:

```
<pxedeploy server="IP" blocksize="4096">
  <timeout>seconds</timeout>
  <kernel>kernel-file</kernel>
  <initrd>initrd-file</initrd>
  <partitions device="/dev/sda">
    <partition type="swap" number="1" size="MB"/>
    <partition type="L" number="2" size="MB"
      mountpoint="/" target="true"/>
    <partition type="fd" number="3"/>
  </partitions>
  <union ro="dev" rw="dev" type="clicfs"/>
  <configuration source="/KIWI/./file" dest="/../file" arch="..."/>
```

```
<configuration .../>
</pxedeploy>
```

- The `server` attribute is used to specify the IP address of the PXE server. The `blocksize` attributes specifies the blocksize for the image download. Other protocols are supported by KIWI but require the `kiwiserver` and `kiwiservertype` kernel parameters to be set when the client boots.
- The value of the optional `timeout` element specifies the grub timeout in seconds to be used when the KIWI initrd configures and installs the grub boot loader on the client machine after the first deployment to allow standalone boot.
- Passing kernel parameters is possible with the use of the optional `kernelcmdline` attribute in the type section. The value of this attribute is a string specifying the settings to be passed to the kernel by the GRUB bootloader. The KIWI initrd includes these kernel options when installing grub for standalone boot
- The optional `kernel` and `initrd` elements are used to specify the file names for the kernel and initrd on the boot server respectively. When using a special boot method not supported by the distribution's standard mkinitrd, it is imperative that the KIWI initrd remains on the PXE server and also be used for local boot. If the configured image uses the `split` type or the `pxedeploy` section includes any union information the kernel and initrd elements must be used.
- The `partitions` section is required if the system image is to be installed on a disk or other permanent storage device. Each partition is specified with one partition child element. The mandatory `type` attribute specifies the partition type id.

The required `number` attribute provides the number of the partition to be created. The size of the partition may be specified with the optional `size` attribute. The optional `mountpoint` attribute provides the value for the mount point of the partition. The optional boolean `target` attribute identifies the partition as the system image target partition. KIWI always generates the swap partition as the first partition of the netboot boot image. By default, the second partition is used for the system image. Use the boolean `target` attribute to change this behavior. Providing the value `image` for the `size` attribute triggers KIWI into calculating the required size for this partition. The calculated size is sufficient for the created image.

- If the system image is based on a read-only filesystem such as squashfs and should be mounted in read-write mode use the optional `union` element. The `type` attribute is used to specify one of the supported overlay filesystem `clircfs` Use the `ro` attribute to point to the read only device and the `rw` attribute to point to the read-write device.
- The optional `configuration` element is used to integrate a network client's configuration files that are stored on the server. The `source` attribute specifies the path on the server for the file to be downloaded. The `dest` attribute specifies destination of the downloaded file on the network client starting at the root (`/`) of the filesystem. Multiple configuration elements may be specified such that multiple files can be transferred to the network client. In addition configuration files can be bound to a specific client architecture by setting the optional `arch` attribute. To specify multiple architectures use a comma separated string.

size

Use the `size` element to specify the image size in Megabytes or Gigabytes. The `unit` attribute specifies whether the given value will be interpreted as Megabytes (`unit="M"`) or

Gigabytes (`unit="G"`). The optional boolean attribute `additive` specifies whether or not the given size should be added to the size of the generated image or not.

In the event of a size specification that is too small for the generated image, KIWI will expand the size automatically unless the image size exceeds the specified size by 100 MB or more. In this case KIWI will generate an error and exit.

Should the given size exceed the necessary size for the image KIWI will not alter the image size as the free space might be required for proper execution of components within the image.

If the size element is not used, KIWI will create an image containing approximately 30 % free space.

```
<size unit="M">1000</size>
```

split

For images of type `split` or `iso` the information provided in the optional `split` section is considered if the `compressed` attribute is set to `true`. With the configuration in this block it is possible to determine which files are writable and whether these files should be persistently writable or temporarily. Note that for ISO images only temporary write access is possible.

When processing the provided configuration KIWI distinguishes between directories and files. For example, providing `/etc` as the value of the `name` attribute indicates that the `/etc` directory should be writable. However, this does not include any of the files or sub-directories within `/etc`. The content of `/etc` is populated as symbolic links to the read-only files. The advantage of setting only a directory to read-write access is that any newly created files will be stored on the disk instead of in `tmpfs`. Creating read-write access to a directory and its files requires two specifications as shown below.

```
<split>
  <temporary>
    <!-- read/write access to -->
    <file name="/var"/>
    <file name="/var/*"/>
    <!-- but not on this file: -->
    <except name="/etc/shadow"/>
  </temporary>
  <persistent>
    <!-- persistent read/write access to: -->
    <file name="/etc"/>
    <file name="/etc/*"/>
    <!-- but not on this file: -->
    <except name="/etc/passwd"/>
  </persistent>
</split>
```

Use the `except` element to specify exceptions to previously configured rules.

machine

The optional `machine` section serves to specify information about a VM guest machine. Using the data provided in this section, KIWI will create a guest configuration file required to run the image on the target machine.

If the target is a VMware virtual machine indicated by the `format` attribute set to `vmdk`, KIWI creates a VMware configuration file. If the target is a Xen virtual machine indicated by the `domain` attribute in the `machine` section KIWI will create a Xen guest config file.

The sample block below shows the general outline of the information that can be specified to generate the configuration file

```
<machine arch="arch" memory="MB"
  HWversion="number" guestOS="suse|sles"
  domain="dom0|domU"/>
  <vmconfig-entry>Entry_for_VM_config_file</vmconfig-entry>
  <vmconfig-entry .../>
  <vmnic driver="name" interface="number" mode="mode"/>
  <vmnic ...>
  <vmdisk controller="ide|scsi" id="number"/>
  <vmdvd controller="ide|scsi" id="number"/>
</machine>
```

arch

The virtualized architecture. Supported values are **ix86** or **x86_64**. The default value is **ix86**.

memory

The mandatory **memory** attribute specifies how much memory in MB should be allocated for the virtual machine

HWversion

The VMware hardware version number, the default value is **3**.

guestOS

The guest OS identifier. For the ix86 architecture the default value is **suse** and for the x86_64 architecture **suse-64** is the default. At this point only the SUSE and SLES guestOS types are supported.

domain

The Xen domain setup. This could be either a **dom0** which is the host machine hosting the guests and therefore doesn't require a configuration file, or it could be set to **domU** which indicates this is a guest and also requires a guest configuration which is created by KIWI.

Use the **vmconfig-entry** element to create entries in the virtual machine's configuration file; **.vmx** for VMware images and **.xenconfig** for Xen images. You may specify as many configuration options as desired. The value of the **vmconfig-entry** element is expected to be specified in the syntax required by the VM configuration file to be written. The value is free format text and is not validated by Kiwi in any way. The entry is written to the VM configuration file verbatim.

Use the **vmdisk** element to setup the virtual main storage device.

controller

Supported values for the mandatory **controller** attribute are **ide** and **scsi**.

id

The mandatory **id** attribute specifies the disk id. If only one disk is set the id value should be set to 0.

device

The device attribute specifies the disk that should appear in the para virtual instance. Therefore only relevant for Xen

Use the **vmdvd** element to setup a virtual optical drive (CD/DVD) connection

controller

Supported values for the mandatory **controller** attribute are **ide** and **scsi**.

id

The mandatory **id** attribute specifies the disk id. If only one disk is set the id value should be set to 0.

Use the **vmnic** element to setup the virtual network interface. Multiple **vmnic** child elements may be specified to setup multiple virtual network interfaces.

driver

The mandatory **driver** attribute specifies the driver to be used for the virtual network card. The supported values are **e100**, **vlance**, and **vmxnet**. If the **vmxnet** driver is specified the vmware tools must be installed in the image.

interface

The mandatory **interface** attribute specifies the interface number. If only one interface is set the value should be set to 0.

mode

The network mode used to communicate outside the VM. In many cases the bridged mode is used.

5.1.5. users Element

```
<users group="group_name" id="number">
  <user home="dir" id="number" name="user" pwd="..."
    pwdformat="encrypted|plain" realname="string" shell="path/>
  <!-- ... -->
</users>
```

The optional **users** element lists the users belonging to the group specified with the **group** attribute. At least one user child element must be specified as part of the **users** element. Multiple **users** elements may be specified.

The attributes **home**, **id**, **name**, **pwd**, **realname**, and **shell** specify the created users home directory, the user name, the user's password, the user's real name, and the user's login shell, respectively. By default, the value of the password attribute is expected to be an encrypted string. An encrypted password can be created using **kiwi --createpassword**. It is also possible to specify the password as a non encrypted string by using the **pwdformat** attribute and setting it's value to "plain". KIWI will then encrypt the password prior to the user being added to the system.

All specified users and groups will be created if they do not already exist. By default, the defined users will be part of the group specified with the **group** attribute of the **users** element and the default group called "users". If it is desired to have the specified users to only be part of the given group it is necessary to specify the **id** attribute. It is recommended to use a group id greater than 100.

5.1.6. drivers Element

```
<drivers profiles="name">
  <file name="filename"/>
  <!-- ... -->
```

```
</drivers>
```

The optional `drivers` element is only useful for boot images (`initrd`). As a boot image doesn't need to contain the complete kernel one can save a lot of space if only the required drivers are part of the image. Therefore the `drivers` section exists. If present only the drivers which matches the file names or glob patterns will be included into the boot image. Each file is specified relative to the `/lib/modules/Version/kernel` directory.

According to the driver element the specified files are searched in the corresponding directory. The information about the driver names is provided as environment variable named like the value of the `type` attribute and is processed by the function `suseStripKernel`. According to this along with a boot image description a script called **images.sh** must exist which calls this function in order to allow the driver information to have any effect.

5.1.7. repository Element

```
<repository type="type" alias="name" imageinclude="true|false"
            password="password" priority="number" status="replaceable"
            username="user-name"> <source path="URL"/>
</repository>
```

The mandatory `repository` element specifies the location and type of a repository to be used by the package manager as a package installation source. The mandatory `type` attribute specifies the repository type. A specified repository can only be accessed by the chosen package manager if the given type is supported by the specified package manager. KIWI supports `smart` or `zypper` as package managers, specified with the `packagemanager` element. The default package manager is `zypper`. The following table shows the possible supported repository types for each package manager:

Table 5.1. Supported Types for zypper and smart

Type	smart Support	zypper Support
apt-deb	yes	no
apt-rpm	yes	no
deb-dir	yes	no
mirrors	yes	no
red-carpet	yes	yes
rpm-dir	yes	yes
rpm-md	yes	yes
slack-site	yes	no
up2date-mirrors	yes	no
urpmi	yes	no
yast2	yes	yes

The `repository` element has the following optional attributes:

`alias="name"`

Specifies an alternative name for the configured repository. If the attribute is not specified KIWI will generate an alias name by replacing any “/” in the given repository location with an “_”. It is helpful to set an alias name if the repository path is insufficient in expressing the purpose of the contained packages.

`imageinclude="true|false"`

Specifies whether the given repository should be configured as a repository in the image or not. The default behavior is that repositories used to build an image are not configured as a repository inside the image. This feature allows you to change the behavior by setting the value to `true`. The repository is configured in the image according to the source path as specified with the `path` attribute of the source element. Therefore, if the path is not a fully qualified URL, you may need to adjust the repository file in the image to accommodate the expected location. It is recommended that you use the `alias` attribute in combination with the `imageinclude` attribute to avoid having unpredictable random names assigned to the repository you wish to include in the image. This also facilitates modification of the "baseurl" entry in the .repo file from the config.sh script if you need to make adjustments to the path.

`password="string"`

Specifies a password for the given repository. The `password` attribute must be used in combination with the `username` attribute. Dependent on the repository location this information may not be used.

`prefer-license="true|false"`

The repository providing this attribute will be used primarily to install the license tarball if found on that repository. If no repository with a preferred license attribute exists, the search happens over all repositories. It's not guaranteed in that case that the search order follows the repository order like they are written into the XML description.

`priority="number"`

Specifies the repository priority for this given repository. Priority values are treated differently by different package managers. Repository priorities allow the package management system to disambiguate packages that may be contained in more than one of the configured repositories. The smart package manager treats packages from repositories with the *highest* priority number as preferable to packages from a repository with a lower priority number. The value 0 means "no priority is set". The zypper package manager prefers packages from a repository with a *lower* priority over packages from a repository with higher priority values. The value 99 means "no priority is set".

`status="replaceable"`

This attribute should only be applied in the context of a boot image description. Setting the `status` to `replaceable` indicates that the specified repository may be replaced by the repositories specified in the image description. This is important as the KIWI generated boot image, if required, should be created based on packages from the same repositories used to build the system image.

`username="name"`

Specifies a user name for the given repository. The `username` attribute must be used in combination with the `password` attribute. Dependent on the repository location this information may not be used.

When specifying an https location for a repository it is generally necessary to include the "openssl-certs" and "cracklib-dict-full" packages in the `bootstrap` section of the image configuration.

The location of a repository is specified by the `path` attribute of the mandatory source child element. The location specification may include the %arch macro which will expand to the architecture of the image building host. The value for the `path` attribute may begin with any of the following location indicators:

`dir:///local/path`

An absolute path to a directory accessible through the local file system. The “`dir:///`” prefix may be omitted.

`ftp://URL`

A ftp protocol based network location.

`http://URL`

A http protocol based network location.

`https://URL`

A https protocol based network location. See the comment above about the handling of certificates and additional package requirements in the `bootstrap` section of the image configuration.

`iso://path/to/isofile`

An absolute path to an `.iso` file accessible via the local file system. KIWI will loop mount the `.iso` file to a KIWI created directory with a generated name. The generated path is provided to the specified package manager as a repository location.

Using multiple `.iso` files from the same SLE product, requires that all `.iso` files are located in the same directory. Only the first `.iso` file is to be specified as a repository in the `config.xml`. The first `.iso` file contains all information necessary for the package manager to locate packages that are contained in other `.iso` files of the same product. Attempting to use multiple `.iso` files in a series as standalone repositories will result in an error.

`obs://$dir1/$dir2`

A special network location used with the http protocol. The values of `$dir1` and `$dir2` represent the project location in the openSUSE build service. The location is evaluated as `this://repos/$dir1/$dir2`.

The “`obs:///`” prefix is also valid as part of the value for the `boot` attribute of the type. If used with the `boot` attribute it is evaluated as `this://images/$dir1/$dir2`.

`opensuse://PROJECTNAME`

A special network location used with the http protocol. The given `PROJECTNAME` specifies a project in the openSUSE builds service. The repository is a repository of type `rpm-md`. For example: `path = "opensuse://openSUSE:10.3/standard"`.

`plain://URI`

A plain resource string. Everything following ‘`plain://`’ will be forwarded to the package manager without further modification. This type of location specification is useful when KIWI does not support a specific URI but the specified package manager does.

`smb://Samba share pathname`

A path to a samba share using the cifs protocol. KIWI creates a mount point and mounts the share including username and password, if specified. Access to the smb share from within the new root tree is provided via a cifs mount. Therefore, the package providing the cifs tools must be included in the package list for the `bootstrap` section of the image configuration. At the time of this writing the package providing the cifs tools is called `cifs-utils`. If any packages provided by the Samba share are used as part of the boot image the cifs tools must also be included in the boot image. This is accomplished with the `bootinclude` attribute of the package element. This is shown in the example below:

```
<packages type="bootstrap">
```

```

...
<package name="cifs-utils" bootinclude="true"/>
</packages>

```

this://PATH

PATH is the relative location to the image description directory for the current image.

5.1.8. packages Element

```

<packages type="type" profiles="name" patternType="type"
  <package name="name" arch="arch"/>
  <package name="name" replaces="name"/>
  <package name="name" bootinclude="true" bootdelete="true"/>
  <archive name="name" bootinclude="true"/>
  <package .../>
  <opensusePattern name="name"/>
  <opensusePattern .../>
  <opensuseProduct name="name"/>
  <opensuseProduct .../>
  <ignore name="name"/>
  <ignore .../>
</packages>

```

The mandatory packages element specifies the list of packages (element package) and patterns (element opensusePattern) to be used with the image. The value of the `type` attribute specifies how the packages and patterns listed are handled, supported values are as follows:

bootstrap

Bootstrap packages, list of packages for the new operating system root tree. The packages list the required components to support a chroot environment in the new system root tree, such as glibc.

delete

Delete packages, list of packages to be deleted from the image being created.

When using the delete type only package elements are considered, all other specifications such as opensusePattern are ignored. The given package names are stored in the `$delete` environment variable of the `/.profile` file created by KIWI. The list of package names is returned by the `baseGetPackagesForDeletion` function. This list can then be used to delete the packages ignoring requirements or dependencies. This can be accomplished in the **config.sh** or **images.sh** script with the following code snippet:

```

rpm -e --nodeps --noscripts \
$(rpm -q 'baseGetPackagesForDeletion' | grep -v "is not installed")

```

Note, that the delete value is indiscriminate of the image type being built.

image

Image packages, list of packages to be installed in the image.

iso

Image packages, a list of additional packages to be installed when building an ISO image.

oem

Image packages, a list of additional packages to be installed when building an OEM image.

pxe

Image packages, a list of additional packages to be installed when building an PXE image.

usb

Image packages, a list of additional packages to be installed when building a USB image.

vmx

Image packages, a list of additional packages to be installed when building a vmx virtual image of any format.

5.1.8.1. Using Patterns

Using a pattern name allows you to considerably shorten the list of specified packages in the `config.xml` file. A named pattern, specified with the `opensusePattern` element is a representation of a predefined list of packages. Specifying a pattern will install all packages listed in the named pattern to be installed in the image. Support for patterns is SUSE-specific, and available with openSUSE 10.1 or later. The optional `patternType` attribute on the `packages` element allows you to control the installation of dependent packages in the image. You may assign one of the following values to the `patternType` attribute:

onlyRequired

Incorporates only patterns and packages that the specified patterns and packages require. This is a "hard dependency" only resolution.

plusRecommended

Incorporates patterns and packages that are required and recommended by the specified patterns and packages in `config.xml`.

By default, only required patterns and packages are installed. KIWI depends on the package manager to resolve the specified list of patterns and packages against the specified repositories and complete the installation. Note that not all supported package managers support the use of named patterns, thus the value of the `packageManager` element determines whether you are able to use named patterns or not. Should the list of specified packages result in a conflict the image creation process will stop and the information provided by the package manager will be captured in the build log and will be displayed in the terminal window where KIWI was started. The `ignore` element may be of use in resolving such conflicts. However, the `ignore` element is limited to effect packages named explicitly. Packages installed in the image through a named pattern are not effected by the `ignore` element setting. Therefore, package conflicts created by packages within named patterns cannot be resolved using the `ignore` mechanism. Further, if a package is specified to be ignored, but is required by another package, then the required package is installed in the image via the automatic dependency resolution by the package manager in use.

5.1.8.2. Architecture Restrictions

To restrict a package to a specific architecture, use the `arch` attribute to specify a comma separated list of allowed architectures. Such a package is only installed if the build systems architecture (`uname -m`) matches one of the specified values of the `arch` attribute.

5.1.8.3. Image Type Specific Packages

If a package is only required for a specific type of image and replaces another package you can use the `replaces` attribute to tell KIWI to install the package by replacing another one. For example you can specify the kernel package in the `type="image"` section as

```
<package name="kernel-default" replaces="kernel-xen"/>
```

and in the `type="xen"` section as

```
<package name="kernel-xen" replaces="kernel-default"/>
```

The result is the xen kernel if you request a xen image and the default kernel in any other case.

5.1.8.4. Packages to Become Included Into the Boot Image

The optional attributes `bootinclude` and `bootdelete` can be used to mark a package inside the system image description to become part of the corresponding boot image (`initrd`). This feature is most often used to specify `boot splash` and/or `graphics boot` related packages inside the system image description but they are required to be part of the boot image as the data is used at boot time of the image.

Packages included into the boot image with the `bootinclude` are still included into the system image as well. If packages should only be included into the boot image, but not the system image, they need to be added to the packages section of `type=delete`.

If the `bootdelete` attribute is specified along with the `bootinclude` attribute this means that the selected package will be marked as a “to become deleted” package and is removed by the contents of the `images.sh` script of the corresponding boot image description.

5.1.8.5. Data not Available as Packages to Become Included

With the optional `archive` element it's possible to include any kind of data into the image. The `archive` element expects the name of a tarball which must exist as part of the system image description. KIWI then picks up the tarball and installs it into the image. If the `bootinclude` attribute is set along with the `archive` element the data will also become installed into the boot image.

6 Creating Appliances with KIWI

Table of Contents

6.1. Overview	41
6.2. The KIWI Model	42
6.3. Cross Platform Appliance Build	43

6.1. Overview

Traditionally, computing functions such as word processing or e-mail handling are delivered as software applications. These applications are targeted to run on a computer with an installed general purpose operating system. Applications often have a specialized installer that must be run by the consumer (whether home computer user or an administrator in an IT department of a company) to install the application on the computer in question. For installation of an application on multiple computers the installation program must often be run on each computer where the application is to be installed. In most cases a given application uses only a small part of the capabilities provided by the general purpose operating system running on a computer. Additionally if an application needs special settings to be applied to the general purpose operating system, these often have to be set by the consumer after the installation is complete. These settings are often documented in an installation guide that consumers may or may not read. Last but not least, running a general purpose operating system to support an application that only requires a small part of the functionality provided by the general purpose OS is a waste of computing resources.

An appliance is the combination of the parts of a general purpose OS needed by a given application and the application itself, bundled and delivered as one unit. This unit may be delivered in a variety of formats, for example a ready to run virtual machine or a self installing system on optical media or a USB stick.

Compared to the traditional model of application delivery the appliance model has a number of advantages. The consumer no longer has to install a general purpose OS and the application in separate steps. The application is part of the appliance and the appliance provider, as the application expert, takes care of the application "installation". Further, the appliance provider takes care of any OS tuning that may benefit the application. Last but not least, the reduced size of the OS does not only consume fewer resources than a full blow "regular" install of a general purpose OS, but it also provides a reduced footprint for potential security exposure. From the application providers point of view there may be an opportunity to drop the implementation and maintenance of a specialized installer as the application installation no longer has to be "consumer friendly".

The traditional software delivery model certainly has it's place. However, for many purposes appliances present a more convenient mechanism for consumers.

6.2. The KIWI Model

With KIWI we started to use a different model. Instead of installing firewall software on top of a general purpose computer/operating system, the designers/engineers built images that are designed specifically for the task. These are so called appliances. When building appliances with KIWI the following proceeding has proven to work reliably. Nevertheless the following is just a recommendation and can be adapted to special needs and environments.

1. Choose an appropriate image description template from the provided KIWI examples. Add or adapt repositories, package names or both, according to the distribution you want to build an image for.

2. Allow the image to create an in-place git repository to allow tracking of non-binary changes. This is done by adding the following line into your **config.sh** script:

```
baseSetupPlainTextGITRepository
```

3. Prepare the preliminary version of your new appliance by calling **kiwi --prepare**
4. Decide for a testing environment. In my opinion a real hardware based test machine which allows to boot from USB is a good and fast approach.

```
<type image="iso" boot="isoboot/suse-..." flags="clib" hybrid="true"/>
```

5. Create the preliminary live stick image of your new appliance by calling **kiwi --create...** After successful creation of the image files find an USB stick which is able to store your appliance and plug it into a free USB port on your image build machine. The deployment can be performed from any OS including Windows as long as a tool to dump data onto a disk device exists and is used.
6. Plug in the stick on your test machine and boot it.
7. After your test system has successfully booted from stick login into your appliance and start to tweak the system according to your needs. This includes all actions required to make the appliance work as you wish. Before you start take care for the following:

- Create an initial package list. This can be done by calling:

```
rpm -qa | sort > /tmp/deployPackages
```

- Check the output of the command **git status** and include everything which is unknown to git and surely will not be changed by you and will not become part of the image description overlay files to the **/.gitignore** files

After the initial package list exists and the git repository is clean you can start to configure the system. You never should install additional software just by installing an unmanaged archive or build and install from source. It's very hard to find out what binary files had been installed and it's also not architecture safe. If there is really no other way for the software to become part of the image you should address this issue directly in your image description and the **config.sh** script but not after the initial deployment has happened.

8. As soon as your system works as expected your new appliance is ready to enter the final stage. At this point you have done several changes to the system but they are all tracked and should now become part of your image description. To include the changes into your image description the following process should be used:

- Check the differences between the currently installed packages and the initial deployment list. This can be done by calling:

```
rpm -qa | sort > /tmp/appliancePackages  
diff -u /tmp/deployPackages /tmp/appliancePackages
```

Add those packages which are labeled with (+) to the `<packages type="image">` section of your `config.xml` file and remove those packages which has been removed (–) appropriately. If there are packages which has been removed against the will of the package manager make sure you address the uninstallation of these packages in your `config.sh` script. If you have installed packages from repositories which are not part of your `config.xml` file you should also add these repositories in order to allow KIWI to install the packages

- Check the differences made in the configuration files. This can be easily done by calling:

```
git diff >/tmp/appliancePatch
```

The created patch should become part of your image description and you should make sure the patch is applied when preparing the image. According to this the command:

```
patch -p0 < appliancePatch
```

needs to be added as part of your `config.sh` script.

- Check for new non binary files added. This can be done by calling:

```
git status
```

All files not under version control so far will be listed by the command above. Check the contents of this list make sure to add all files which are not created automatically to become part of your image description. To do this simply clone (copy) these files with respect to the filesystem structure as overlay files in your image description root/directory.

9. All your valuable work is now stored in one image description and can be re-used in all KIWI supported image types.

Congratulation! To make sure the appliance works as expected prepare a new image tree and create an image from the new tree. If you like you can deactivate the creation of the git repository which will save you some space on the filesystem. If this appliance is a server I recommend to leave the repository because it allows you to keep track of changes during the live time of this appliance.

6.3. Cross Platform Appliance Build


Building appliances for one processor architecture on another processor architecture is in general not possible with KIWI. The exception is that it is possible to build 32 bit (ix86) appliances on a 64 bit system running on the x86-64 architecture. This cross-platform limitation is based on the requirement that KIWI be able to execute installed software inside the unpacked image tree. If the software installed inside the unpacked image tree does not run on the architecture of the build platform then KIWI cannot build the appliance.

While KIWI has the `--target-arch` command line argument to instruct the package manager *zypper* to install packages for the specified architecture, this option is not intended to support cross-platform appliance builds.

7 Maintenance of Operating System Images

Creating an image often results in an appliance solution for a customer and gives you the freedom of a working solution at that time. But software develops and you don't want your solution to become outdated. Because of this together with an image people always should think of *image-maintenance*. The following paragraph just reflects ideas how to maintain images created by KIWI:

Figure 7.1. Image Maintenance Scenarios



- ❶ Image Description changes, requires free space to store /tmp/myOSImage, slower, because Subversion, etc. to track changes
- ❷ Software package source changes
- ❸ Faster, because already prepared, cannot handle image description changes, requires free space to store /tmp/myOSImage
- ❹ Image Description changes
- ❺ Covers all possible changes, does not require storage for prepared trees, slower, because KIWI prepare runs again

The picture in Figure 7.1 shows two possible scenarios which requires an image to become updated. The first reason for updating an image are changes to the software, for example a new kernel should be used. If this change doesn't require additional software or changes in the configuration the update can be done by KIWI itself using its `--upgrade` option. In combination with `--upgrade` KIWI allows to add an additional repository which may be needed if the updated software is not part of the original repository. An important thing to know is that this additional repository is *not* stored into the original `config.xml` file of the image description.

Another reason for updating an image beside software updates are configuration changes or enhancements, for example an image should have replaced its browser with another better browser or a new service like apache should be enabled. In principle it's possible to do all those changes manually within the physical extend but concerning maintenance this would be a nightmare. Why, because it will leave the system in an unversioned condition. Nobody knows what has changed since the very first preparation of this image. So in short:

Don't modify physical extends manually!

Changes to the image configuration should be done within the image description. The image description itself should be part of a versioning system like subversion. All changes can be tracked down then and maybe more important can be assigned to product tags and branches. As a consequence an image must be prepared from scratch and the old physical extend could be removed.

8 System to Image Migration

Table of Contents

8.1. Create a Clean Repository Set First	49
8.2. Watch the Overlay and Unpackaged Files	50
8.3. Checklist	50
8.4. Turn my System Into an Image... ..	50

KIWI provides an experimental module which allows you to turn your running system into an image description. This migration allows you to clone your currently running system into an image. The process has the following limitations at the moment:

- Works for SUSE systems only (with zypper on board)
- The process works semi automatically which means depending on the complexity of the system some manual postprocessing might be necessary

When calling KIWI's migrate mode it will try to find the base version of your operating system and uses the currently active repositories specified in the zypper database to match the software which exists in terms of packages and patterns. The result is a list of packages and patterns which represents your system so far. Of course there are normally some data which doesn't belong to any package. These are for example configurations or user data. KIWI collects all this information and would copy it as overlay files as part of the image description. The process will skip all remote mounted filesystems and concentrate only on local filesystems.

8.1. Create a Clean Repository Set First

When starting with the migration it is useful to let kiwi know about all the repositories from which packages has been installed to the system. In a first step call:

```
kiwi --migrate mySystem
```

This will create an HTML report where you can check which packages and patterns could be assigned to the given base repository. In almost all cases there will be information about packages which couldn't be assigned. You should go to that list and think of the repository which contains that packages (Packman, etc). If something is missing add it either to the zypper list on your system or use the KIWI options `--add-repo ... --add-repo-type`.

Continue calling the following command until your list is clean You should continue the migration if you have a clean list of solved packages without any package skipped except you know that this package can't be provided or is not worth to become part of the migration.

```
kiwi --migrate mySystem --nofiles [--skip package ... ]
```

8.2. Watch the Overlay and Unpackaged Files

Files which has been modified but belong to a package will be automatically copied into the overlay directory below `/tmp/mySystem/root`. You should check that no modified file is a binary because such a binary would be replaced by a new install of the package anyway. Software developers tend to compile software from source and copy/install them into their system. Doing this could cause binary files previously installed by a package to be reported as modified. You should remove such files from your overlay tree.

The migration also copy the entire `/etc` directory into the overlay root directory because it stores all important configuration files. Beside the important files there are most probably a bunch of file which doesn't belong to any package exists only for historical reasons. kiwi creates a list of files and directories to support you best in sorting out what is required and what can be ignored. Nevertheless this is the most time consuming part of your migration review. Simply click on the *all unpackaged files* link to take a look at the complete list. Those files you want to have in your image needs to be copied over to the `/tmp/mySystem/root` directory

8.3. Checklist

After that you should walk through the following check list

- Change author and contact in `config.xml`
- Set appropriate name for your image in `config.xml`.
- Add/modify default type (oem) set in `config.xml` if needed
- Make sure your X11 configuration is appropriate according to the new target. A failsafe version was created in `/tmp/mysys/root/etc/X11/xorg.conf.install -> fbdev` based
- Make sure **yast2** is installed to be able to reconfigure the system. If **yast2** is not installed these tasks needs to be done else. Otherwise yast's second stage is started on first boot of the migrated image
- If you want to access any remote filesystem it's a good idea to let AutoYaST add them on first boot of the system
- Check your network setup in `/etc/sysconfig/network`. Is this setup still possible in the cloned environment? Make sure you check for the MAC address of the card first.

8.4. Turn my System Into an Image...

After the process has finished you should check the size of the image description. The description itself shouldn't be that big. The size of a migrated image description mainly depends on how many overlay files exists in the `root/` directory. You should make sure to maintain only required overlay files. Now let's try to create a clone image from the description. By default an OEM image which is a virtual disk which is able to run on real hardware too is created.

On success you will also find a ISO file which is an installable version of the OEM image. If you burn the ISO on a DVD you can use that DVD to install your cloned image on another computer.

```
kiwi -p /tmp/migrated --root /tmp/mySys  
kiwi --create /tmp/mySys -d /tmp/myResult
```

If everything worked well you can test the created OEM image in any full virtual operating system environment like Qemu or VMware™. Once created the image description can serve for all image types KIWI supports.

9 Installation Source

Table of Contents

9.1. Adapt the Example's config.xml	53
9.2. Create a Local Installation Source	53

Before you start to use any of the examples provided in the following chapters your build system has to have a valid installation source for the distribution you are about to create an image for. By default, all examples will connect to the network to find the installation source. It depends on your network bandwidth how fast an image creation process is and in almost all cases it is better to prepare a local installation source first.

9.1. Adapt the Example's config.xml

If you can make sure you have a local installation source it's important to change the path attribute inside of the repository element of the appropriate example to point to your local source directory. A typically default repository element looks like the following:

```
<repository type="yast2">
  <source path="opensuse://openSUSE:##.##/standard/" />
</repository>
```

9.2. Create a Local Installation Source

The following procedure describes how to create a local SUSE installation source which is stored below the path `/images/CDs`. If you are using the local path as described in this document you only need to flip the given path information inside of the example `config.xml` file.

1. Find your SUSE standard installation CDs or the DVD and make them available to the build system. Most Linux systems auto-mount a previously inserted media automatically. If this is the case you simply can change the directory to the auto mounted path below `/media`. If your system doesn't mount the device automatically you can do this with the following command:

```
mount -o loop /dev/drive-device-name /mnt
```

2. If you do not have a DVD but a CD set, copy the contents of *all* CDs into one directory. It's absolutely important that you first start with the *last* CD and copy the first CD at last. In case of CDs you should have a bundle of 4 CDs. Copy them in the order 4 3 2 1.
3. Copy the contents of the CDs/DVD to your hard drive once you have access to the media. You need at least 4GB free space available. The following is intended to create an openSUSE installation source:

Create a Local Installation Source

```
mkdir -p /image/CDs/full-##.#-i386/  
cp -a /mnt/* /image/CDs/full-##.#-i386/
```

Remember if you have a CD set start with number 4 first and after that, eject the CD and insert the next one to repeat the copy command until all CDs are copied into to /image

10 ISO Image—Live Systems

Table of Contents

10.1. Building the suse-live-iso Example	55
10.2. Using the Image	55
10.3. Flavours	55
10.4. USB stick images	56

A live system image is an operating System on CD or DVD. In principle one can treat the CD/DVD as the hard disk of the system with the restriction that you can't write data on it. So as soon as the media is plugged into the computer, the machine is able to boot from that media. After some time one can login to the system and work with it like on any other system. All write actions takes place in RAM space and therefore all changes will be lost as soon as the computer shuts down.

10.1. Building the suse-live-iso Example

This example is based on openSUSE and includes the KDE desktop.

```
cd /usr/share/doc/packages/kiwi/examples
==> select the example directory for the desired distribution change into it
cd suse-...
kiwi --build ./suse-live-iso -d /tmp/myiso-result --type iso
```

10.2. Using the Image

There are two ways to use the generated ISO image:

- Burn the .iso file on a CD or DVD with your preferred burn program. Plug in the CD or DVD into a test computer and (re)boot the machine. Make sure the computer boot from the CD drive as first boot device.
- Use a virtualization system to test the image directly. Testing an iso can be done with any full virtual system for example:

```
cd /tmp/myiso-result
qemu -cdrom ./suse-*-live-iso.*.iso
```

10.3. Flavours

KIWI supports different filesystems and boot methods along with the ISO image type. The provided example by default uses a clicfs compressed root filesystem. clicfs is a fuse user

space filesystem which reads in data from a compressed image and writes data into a cow file which can exist in RAM or in persistent area on a disk. The result is a full writable live-system. The flags attribute in config.xml exists to be able to have the following alternative solutions:

`flags="compressed"`

Does filesystem compression with squashfs, but don't use an overlay filesystem for write support. A symbolic link list is used instead and thus a split element is required in config.xml. See the split mode section below for details.

`flags="clirc|clirc_udf"`

Creates a FUSE based clicfs image and allows write operations into a cow file. In case of an ISO the write happens into a ramdisk. If clic_udf is specified the iso is created with an udf filesystem and thus this allows to create live systems bigger than 4G

Flags Not Set

If no `flags` attribute is set no compressed filesystem, no overlay filesystem will be used. The root tree will be directly part of the ISO filesystem and the paths: `/bin`, `/boot`, `/lib`, `/lib64`, `/opt`, `/sbin`, and `/usr` will be read-only.

10.3.1. Split mode

If no overlay filesystem is in use but the image filesystem is based on a compressed filesystem KIWI allows to setup which files and directories should be writable in a so called split section. In order to allow to login into the system, at least the `/var` directory should be writable. This is because the PAM authentication requires to be able to report any login attempt to `/var/log/messages` which therefore needs to be writable. The following split section can be used if the flag `compressed` is used:

```
<split>
  <persistent>
    <file name="/var"/>
    <file name="/var/*"/>
    <file name="/boot"/>
    <file name="/boot/*"/>
    <file name="/etc"/>
    <file name="/etc/*"/>
    <file name="/home"/>
    <file name="/home/*"/>
    <file name="/tmp"/>
    <file name="/tmp/*"/>
  </persistent>
</split>
```

10.3.2. Hybrid mode

A hybrid image is a iso image including a partition table and can therefore be attached as a CD/DVD *and* as a normal disk to the system. This has the advantage that a hybrid iso live system can be burned to a CD/DVD as well as uploaded to a USB stick. In order to activate the hybrid feature the hybrid flag must be set to true as indicated below.

```
<type image="iso" ... hybrid="true"/>
```

10.4. USB stick images

kiwi supports two types of USB stick images. The first type which are the hybrid ISO images and basically the same as the live ISO images and the second type which are the OEM virtual disk images. The deployment of both types can be performed from any OS including Windows as long as a tool to dump data onto a disk device exists and is used.

10.4.1. ISO Hybrid stick

As indicated above a hybrid iso image also works as USB stick image. If a hybrid iso is used like a disk image on a writable medium like a USB stick it's possible to write into a persistent area on the stick instead of the RAM. kiwi will create an additional ext2 partition to store that information on the disk if the attribute `hybridpersistent` is set to `true`.

```
<type image="iso" ... hybridpersistent="true"/>
```

10.4.2. OEM USB stick

In contrast to the hybrid iso image it's also possible to create a oem virtual disk image which is dumped on the stick. The big advantage with this approach is, that it's possible to create a stick which contains a live OS but also a data partition for custom data. The data partition is a fat partition also recognized by the Windows operating system. In order to create such a Windows friendly stick one has to pass the option `--fat-storage <size-in-MB>`.

```
kiwi --create ... --fat-storage 500
```

If this option is set kiwi will use the syslinux bootloader for the image as well as the first partition as fat partition of the specified size. The live OS itself will live in a LVM which allows easy manipulation of the logical root volume. For further information about the OEM image type please refer to the OEM chapter Chapter 14, *OEM Image—Preload Systems*

10.4.2.1. OEM compressed / readonly USB stick

If a compressed filesystem type like `clircfs` is used for the image root directory it's also possible to allow persistent writing on the USB stick or alternatively disallow that and let all write actions perform in RAM only. kiwi provides the type attribute `ramonly` for this purpose. So in order to create a read-only oem stick with compressed root filesystem the following type section is required:

```
<type image="oem" filesystem="clircfs" ramonly="true" .../>
```

11 VMX Image—Virtual Disks

Table of Contents

11.1. Building the suse-vm-guest Example	59
11.2. Using the Image	59
11.3. Flavours	59

A VMX image is a virtual disk image for use in full virtualization systems like Qemu or VMware. The image is a file containing the system represented by the configured packages in `config.xml` as well as partition data and bootloader information. The size of this virtual disk can be specified by using the `size` element in the `config.xml` file or by adding the `--bootvm-disksize` command line argument.

11.1. Building the suse-vm-guest Example

The `vm-guest` example provided with KIWI is based on recent openSUSE releases, one example configuration per release. The example uses base pattern and the virtual disk is formatted using the distribution default filesystem.

```
cd /usr/share/doc/packages/kiwi/examples
cd suse-...
kiwi --prepare ./suse-vm-guest --root /tmp/myvm

kiwi --create /tmp/myvm --type vmx -d /tmp/myvm-result
```

11.2. Using the Image

The generated virtual disk image serves as the hard disk of the selected virtualization system (QEMU, VMware, etc.). The virtual hard disk format differs across virtualization environments. Some virtualization environments support multiple virtual disk formats. Using the QEMU virtualization environment test the created image with the following command:

```
cd /tmp/myvm-result
qemu suse-##.##-vm-guest.i686-1.1.2.raw -m 256
```

11.3. Flavours

KIWI always generates a file in the `.raw` format. The `.raw` file is a disk image with a structure equivalent to the structure of a physical hard disk. Individual virtualization systems have

specific formats to facilitate improved I/O performance to the virtual disk, represented by the image file, or additional specified virtual hard disk files. KIWI will generate a specific format when the format attribute of the type element is added.

```
<type image="vmx"... format="name"/>
```

The following table lists the supported virtual disk formats:

Table 11.1. Supported Virtual Disk Formats

Name	Description
vmdk	Disk format for VMware
vhd	Disk format for Microsoft HyperV
ovf	Open Virtual Format requires VMware's ovftool
qcow2	QEMU virtual disk format

11.3.1. VMware support

A VMware image is accompanied by a guest configuration file. This file includes information about the hardware to be represented to the guest image by the VMware virtualization environment as well as specification of resources such as memory.

Within the config.xml file it is possible to specify the VMware configuration settings. In addition it is possible to include selected packages in the created image that are specific to the VM image generation. The following config.xml snippet provides general guidance on the elements in config.xml.

```
<packages type="vmx">
  <!-- packages you need in VMware only -->
</packages>
<type.....>
  <machine memory="512">
    <vm disk controller="ide" id="0"/>
  </machine>
</type>
```

Given the specification above KIWI will create a VMware guest configuration specifying the availability of 512 MB of RAM and an IDE disk controller interface for the VM guest. For additional information about the configuration settings please refer to the *machine* section.

The guest configuration can be loaded through VMware user interface and may be modified through the GUI. The configuration file has the .vmx extension as shown in the example below.

```
/tmp/myvm-result/suse-##.##-vm-guest.i686-1.1.2.vmx
```

Using the `format="vmdk"` attribute of the `<type>` start tag will create the VMware formatted disk image (.vmdk file) and the required VMware guest configuration (.vmx) file.

In addition it is possible to create an image for the Xen virtualization framework. By adding the `bootprofile` and `bootkernel` attributes to the `<type>` start tag with values of `xen` and `xenboot`, respectively. Please refer to the Chapter 15, *Xen Image—Paravirtual Systems* for additional details.

11.3.2. LVM Support

KIWI also provides support for LVM (Logical Volume Management). In this mode the disk partition table will include one lvm partition and one standard ext2 boot partition. KIWI creates the kiwiVG volume group and adds logical volumes as they are needed and configured according to the image type and filesystem. After boot of the system the user has full control over the volume group and is free to change/resize/increase the group and the volumes inside. Support for LVM has been added for all image types which are disk based. This includes vmx, oem and usb. In order to use LVM for the vmx type just add the `--lvm` option as part of the KIWI create step or add the attribute `lvm = "true"` as part of the type section in your `config.xml` file.

```
kiwi --create /tmp/myvm --type vmx -d /tmp/myvm-result --lvm
```

With the optional `systemdisk` section you can set one or more top level directories into a separate volume. See Chapter 5, *KIWI Image Description* for a detailed explanation.

12 Linux Container

Table of Contents

12.1. Building the suse-lxc-guest Example	63
12.2. Using the Image	63
12.3. Image Configuration Details	64

Linux Containers (LXC) [<http://lxc.sourceforge.net/>] provide operating system-level virtualization, utilizing *Control Groups (cgroups)* [<https://www.kernel.org/doc/Documentation/cgroups/cgroups.txt>]. The virtualization is similar to technologies in OpenVZ, Linux-VServer, FreeBSD jails, AIX Workload Partitions, and Solaris Containers. The network and process space of the container is separated from the host resources using namespaces. Kernel space information is mounted into the container root filesystem using the `fstab` file in the configuration directory. The container root filesystem provides the new user space on top of the running kernel of the host. A linux container has two components, the container root file system stored in `/var/lib/lxc/CONTAINER_NAME` and the container configuration stored in `/etc/lxc/CONTAINER_NAME`. The kiwi created container image is packaged in a tarball containing the root filesystem and the configuration. The tarball is expected to be inflated at the root level (/) of the target system that functions as host.

12.1. Building the suse-lxc-guest Example

The lxc-guest example provided with KIWI is based on recent openSUSE releases, one example configuration per release. The example provides a very minimal system.

```
cd /usr/share/doc/packages/kiwi/examples
cd suse-...
kiwi --prepare ./suse-lxc-guest --root /tmp/mylxc

kiwi --create /tmp/mylxc --type lxc -d /tmp/mylxc-result
```

12.2. Using the Image

The created container is packaged in a tarball in the destination directory, `mylxc-result` in the example above. Move this tarball to the root level (/) of the host machine and unpack it. The following commands assume that the image build machine is also the host machine.

```
cp /tmp/mylxc-result/suse-##.#-lxc-guest-lxc.*-1.0.0.tbz /
cd /
tar -xjf suse-##.#-lxc-guest-lxc.*-1.0.0.tbz
```

```
lxc-start -n os### -f /etc/lxc/os###/config
```

12.3. Image Configuration Details

The configuration for a container does not need to contain a kernel package. The container represents the user space that runs on top of the kernel of the container host system.

The container itself must contain the Linux user space container tools and thus the *lxc* package must be included in the container image.

Configure the network configuration for the container using the *vmnic* element. The *mode* attribute indicates the network mode, *veth* by default. While it is possible to configure multiple network interfaces in the *config.xml* file, the written container configuration will only reflect the information configured for the first *vmnic* element found in the *config.xml* file. The configuration for the container expects that the host has a configured network bridge with the name *br0*. For complex network setup implementations it is necessary to edit the *config* file.

The generated configuration file restricts the device access of the container according to a generally accepted best practice security model. The device access permissions may be modified by editing the *config* file for the container.

13 PXE Image—Thin Clients

Table of Contents

13.1. Setting Up the Required Services	65
13.2. Building the suse-pxe-client Example	66
13.3. Using the Image	66
13.4. Flavours	67
13.5. Hardware Grouping	76

PXE is a boot protocol implemented in most BIOS implementations which makes it so interesting. The protocol sends DHCP requests to assign an IP address and after that it uses tftp to download kernel and boot instructions.

A PXE image consists of a boot image and a system image like all other image types too. But with a PXE image the image files are available separately and needs to be copied at specific locations of a network boot server.

13.1. Setting Up the Required Services

Before you start to build PXE images with KIWI, setup the boot server. The boot server requires the services atftp and DHCP to run.

13.1.1. Atftp Server

In order to setup the atftp server the following steps are required

1. Install the packages atftp and kiwi-pxeboot.
2. Edit the file /etc/sysconfig/atftpd. Set or modify the following variables:

- `ATFTPD_OPTIONS="--daemon --no-multicast"`

- `ATFTPD_DIRECTORY="/srv/tftpboot"`

3. Run atftpd by calling the command:

```
rcatftpd start
```

13.1.2. DHCP Server

In contrast to the atftp server setup the following DHCP server setup can only serve as an example. Depending on your network structure, the IP addresses, ranges and domain settings needs to be adapted in order to allow the DHCP server to work within your network. If you

already have a DHCP server running in your network, make sure that the filename and next-server information is provided by your server. The following steps describe how to setup a new DHCP server instance:

1. Install the package `dhcp-server`.
2. Create the file `/etc/dhcpd.conf` and include the following statements:

```
option domain-name "example.org";
option domain-name-servers 192.168.100.2;
option broadcast-address 192.168.100.255;
option routers 192.168.100.2;
option subnet-mask 255.255.255.0;
default-lease-time 600;
max-lease-time 7200;
ddns-update-style none; ddns-updates off;
log-facility local7;

subnet 192.168.100.0 netmask 255.255.255.0 {
    filename "pxelinux.0";
    next-server 192.168.100.2;
    range dynamic-bootp 192.168.100.5 192.168.100.20;
}
```

3. Edit the file `/etc/sysconfig/dhcpd` and setup the network interface the server should listen on:

```
DHCPD_INTERFACE="eth0"
```

4. Run the dhcp server by calling:

```
rcdhcpd start
```

13.2. Building the suse-pxe-client Example

The example provided with KIWI is based on openSUSE and creates an image for a Wyse VX0 terminal with a 128MB flash card and 512MB of RAM. The image makes use of the squashfs compressed filesystem and its root tree is deployed as clicfs based overlay system.

```
cd /usr/share/doc/packages/kiwi/examples
==> select the example directory for the desired distribution change into it
cd suse-...
kiwi --build ./suse-pxe-client -d /tmp/mypxe-result --type pxe
```

13.3. Using the Image

In order to make use of the image all related image parts needs to be copied onto the boot server. According to the example the following steps needs to be performed:

1. Change working directory:

```
cd /tmp/mypxe-result
```

2. Copy of the boot and kernel image:

```
cp initrd-netboot-suse-*.splash.gz \
  /srv/tftpboot/boot/initrd
cp initrd-netboot-suse-*.kernel \
```

```
/srv/tftpboot/boot/linux
```

3. Copy of the system image and md5 sum:

```
cp suse-*-pxe-client.* /srv/tftpboot/image
```

4. Copy of the image boot configuration. Normally the boot configuration applies to one client which means it is required to obtain the MAC address of this client. If the boot configuration should be used globally, copy the KIWI generated file as config.default:

```
cp suse-*-pxe-client.*.config \
/srv/tftpboot/KIWI/config.MAC
```

5. Check the PXE configuration file. The PXE configuration controls which kernel and initrd are loaded and which kernel parameters are set. When installing the kiwi-pxe-boot package, a default configuration is added. To make sure the configuration is valid according to this example, insert the following information into the file /srv/tftpboot/pxelinux.cfg/default:

```
DEFAULT KIWI-Boot

LABEL KIWI-Boot
    kernel boot/linux
    append initrd=boot/initrd vga=0x314
    IPAPPEND 1

LABEL Local-Boot
    localboot 0
```

6. Connect the client to the network and boot.

13.4. Flavours

All the different PXE boot based deployment methods are controlled by the config.MAC (or config.default) file. When a new client boots up and there is no client configuration file the new client is registered by uploading a control file to the TFTP server. The following sections inform about the control and the configuration file.

13.4.1. The PXE Client Control File

This section describes the netboot client control file:

```
hwtype.$<$MAC Address>$
```

The control file is primarily used to set up new netboot clients. In this case, there is no configuration file corresponding to the client MAC address available. Using the MAC address information, the control file is created, which is uploaded to the TFTP servers upload directory /var/lib/tftpboot/upload.

13.4.2. The PXE Client Configuration File

This section describes the netboot client configuration file:

```
config.$<$MAC Address>$
```

The configuration file contains data about image, configuration, synchronization, or partition parameters. The configuration file is loaded from the TFTP server directory /var/lib/tftpboot/KIWI via TFTP for previously installed netboot clients. New netboot clients are immedi-

ately registered and a new configuration file with the corresponding MAC address is created. The standard case for the deployment of a PXE image is one image file based on a read-write filesystem which is stored onto a local storage device of the client. Below, find an example to cover this case.

```
DISK=/dev/sda
PART='5;S;x,x;L;/'
IMAGE='/dev/sda2;suse-##.##-pxe-client.i686;1.2.8;192.168.100.2;4096'
```

The following format is used:

```
IMAGE='device;name;version;srvip;bsize;compressed,...,'
CONF='src;dest;srvip;bsize;[hash],...,src;dest;srvip;bsize;[hash]'
PART='size;id;Mount,...,size;id;Mount'
DISK=device
```

IMAGE

Specifies which image (name) should be loaded with which version (version) and to which storage device (device) it should be linked, e. g., /dev/ram1 or /dev/hda2. The netboot client partition (device) hda2 defines the root file system / and hda1 is used for the swap partition. The numbering of the hard disk device should not be confused with the RAM disk device, where /dev/ram0 is used for the initial RAM disk and can not be used as storage device for the second stage system image. SUSE recommends to use the device /dev/ram1 for the RAM disk. If the hard drive is used, a corresponding partitioning must be performed.

srvip

Specifies the server IP address for the TFTP download. Must always be indicated, except in PART.

bsize

Specifies the block size for the TFTP download. Must always be indicated, except in PART. If the block size is too small according to the maximum number of data packages (32768), linuxrc will automatically calculate a new blocksize for the download.

compressed

Specifies if the image file on the TFTP server is compressed and handles it accordingly. To specify a compressed image download only the keyword "compressed" needs to be added. If compressed is not specified the standard download workflow is used. **Note:** The download will fail if you specify "compressed" and the image isn't compressed. It will also fail if you don't specify "compressed" but the image is compressed. The name of the compressed image has to contain the suffix .gz and needs to be compressed with the **gzip** tool. Using a compressed image will automatically *deactivate* the multicast download option of atftp.

CONF

Specifies a comma-separated list of source:target configuration files. The source (src) corresponds to the path on the TFTP server and is loaded via TFTP. The download is made to the file on the netboot client indicated by the target (dest). Download only happens when configuration files are missing on the client or, if md5sum hash is supplied ([hash]), when different. To achieve this, list of CONF files (and VENDOR_CONF) files is kept on the client in the /etc/KIWI/InstalledConfigFiles backup file, and is compared to the CONF data gathered from the config.MAC and also from other configuration files, e.g. config.group, if supplied. Configuration files selected for comparison are those with same (dest) path. If destination path (dest) is same for more configuration files, only the last one is used (and VENDOR_CONF has always precedence to CONF). By comparing configuration file

lists present in the current CONF, VENDOR_CONF variables and stored in the backup file, following actions can result:

Table 13.1. Configuration files synchronization possibilities

cfg CONF,VENDOR_CONF	file in	cfg file in InstalledConfig- Files backup	action
hash_a		hash_a	nothing, keep
hash_a		hash_b	download from server
none		hash	download from server
hash		none	download from server
none		none	nothing, keep
present		not present	download from server (re- gardless hash)
not present		present	delete on client (regardles hash)

Note that actual configuration files (or their md5sum hashes) on the client machine are not tested, only data from the backup file are used. This means that actual configuration files can be altered or even deleted without triggering any action, or, on the other hand, an action can be triggered without modifying the configuration files, only by modifying or removing of the backup file.

PART

Specifies the partitioning data. The comma-separated list must contain the size (size), the type number (id), and the mount point (Mount). The size is measured in MB by default. The mount specifies the directory the partition is mounted to.

- The first element of the list must define the swap partition.
- The second element of the list must define the root partition.
- The swap partition must not contain a mount point. A lowercase letter x must be set instead.
- If a partition should take all the space left on a disk one can set a lower x letter as size specification.

RAID

In addition to the PART line it's also allowed to add a raid array setup. The first parameter of the RAID line is the raid level. So far only raid1 (mirroring) is supported. The second and third parameter specifies the raid disk devices which makes up the array. If a RAID line is present all partitions in PART will be created as raid partitions. The first raid is named md0 the second one md1 and so on. It's required to specify the correct raid partition in the IMAGE line according to the PART setup. A typical raid image setup could look like this:

```
DISK=/dev/sda  
RAID='1;/dev/sda;/dev/sdb'  
IMAGE='/dev/md1;LimeJeOS-openSUSE-##.##.i686;1.11.3;192.168.100.2;4096'  
PART='5;S;x,2000;83;/'
```

DISK

Specifies the hard disk. Used only with PART and defines the device via which the hard disk can be addressed, e.g., /dev/hda.

REBOOT_IMAGE

If set to a non-empty string, this will reboot the system after the initial deployment process is done. This means before the system init process is activated the system is rebooted. If the machine's default boot setup is to boot via PXE it will again boot from the network.

FORCE_KEXEC

During the initial deployment process kiwi checks if the running kernel is the same as the kernel installed via the system image. If there is a mismatch kiwi activates the installed kernel by calling kexec. This is mostly the same as to perform a reboot but without the need of the BIOS or any bootloader. If FORCE_KEXEC is set to a non-empty string kiwi will also perform kexec if the kernel versions matches.

RELOAD_IMAGE

If set to a non-empty string, this forces the configured image to be loaded from the server even if the image on the disk is up-to-date. The primary purpose of this setting is to aid debugging. The option is sensible only for disk based systems.

RELOAD_CONFIG

If set to a non-empty string, this forces all config files to be loaded from the server. The primary purpose of this setting is to aid debugging. The option is sensible only for disk based systems.

COMBINED_IMAGE

If set to a non-empty string, indicates that the both image specified needs to be combined into one bootable image, whereas the first image defines the read-write part and the second image defines the read-only part.

KIWI_INITRD

Specifies the KIWI initrd to be used for local boot of the system. The variables value must be set to the name of the initrd file which is used via PXE network boot. If the standard tftp setup suggested with the kiwi-pxeboot package is used all initrd files resides in the boot/ directory below the tftp server path /var/lib/tftpboot. Because the tftp server do a chroot into the tftp server path you need to specify the initrd file as the following example shows:

```
KIWI_INITRD=/boot/name-of-initrd-file
```

UNIONFS_CONFIG

For netboot images there is the possibility to use clicfs as container filesystem in combination with a compressed system image. The recommended compressed filesystem type for the system image is **clicfs**.

```
UNIONFS_CONFIG=/dev/sda2,/dev/sda3,clicfs
```

In this example the first device /dev/sda2 represents the read/write filesystem and the second device /dev/sda3 represents the compressed system image filesystem. The container filesystem clicfs is then used to cover the read/write layer with the read-only device to one read/write filesystem. If a file on the read-only device is going to be written the changes inodes are part of the read/write filesystem. Please note the device specifications in UNIONFS_CONFIG must correspond with the IMAGE and PART information. The following example should explain the interconnections:

```
DISK=/dev/sda  
IMAGE='/dev/sda3;image/myImage;1.1.1;192.168.1.1;4096'  
PART='200;S;x,300;L;/,x;L;x'  
UNIONFS_CONFIG=/dev/sda2,/dev/sda3,clicfs
```

As the second element of the PART list must define the root partition it's absolutely important that the first device in UNIONFS_CONFIG references this device as read/write device. The second device of UNIONFS_CONFIG has to reference the given IMAGE device name.

KIWI_KERNEL_OPTIONS

Specifies additional command line options to be passed to the kernel when booting from disk. For instance, to enable a splash screen, you might use `vga=0x317 splash=silent`.

KIWI_BOOT_TIMEOUT

Specifies the number of seconds to wait at the grub boot screen when doing a local boot. The default is 10.

NBDROOT

Mount the system image root filesystem remotely via NBD (Network Block Device). This means there is a server which exports the root directory of the system image via a specified port. The kernel provides the block layer, together with a remote port that uses the `nbd-server` program. For more information on how to set up the server, see the `nbd-server` man pages. The kernel on the remote client can set up a special network block device named `/dev/nb0` using the `nbd-client` command. After this device exists, the mount program is used to mount the root filesystem. To allow the KIWI boot image to use that, the following information must be provided:

```
NBDROOT=NBD.Server.IP.address;\
NBD-Port-Number;/dev/NBD-Device;\
NBD-Swap-Port-Number;/dev/NBD-Swap-Device;\
NBD-Write-Port-Number;/dev/NBD-Write-Device
```

The `NBD-Device`, `NBD-Swap-Port-Number`, `NBD-Swap-Device`, `NBD-Write-Port-Number` and `NBD-Write-Device` variables are optional. If the `nbd` root device is not set, the default values (`/dev/nb0`, port 2000) applies and if the `nbd` swap device is not set the default values (`/dev/nb1`, port 9210) applies. The swap space over the network using a network block device is only established if the client has less than 48 MB of RAM. The optional `NBD-Write-Port-Number` and `NBD-Write-Device` specifies a write COW location for the root filesystem. `clifs` is used as overlay filesystem in this case.

AOEROOT

Mount the system image root filesystem remotely via AoE (ATA over Ethernet). This means there is a server which exports a block device representing the root directory of the system image via the AoE subsystem. The block device could be a partition of a real or a virtual disk. In order to use the AoE subsystem I recommend to install the `aoetools` and `vblade` packages from here first: <http://download.opensuse.org/repositories/server:/lts/>. Once installed the following example shows how to export the local `/dev/sdb1` partition via AoE:

```
vbladed 0 1 eth0 /dev/sdb1
```

Some explanation about this command, each AoE device is identified by a couple Major/Minor, with major between 0-65535 and minor between 0-255. AoE is based just over Ethernet on the OSI models so we need to indicate which ethernet card we'll use. In this example we export `/dev/sdb1` with a major value of 0 and minor of 1 on the `eth0` interface. We are ready to use our partition on the network! To be able to use the device KIWI needs the information which AoE device contains the root filesystem. In our example this is the device `/dev/etherd/e0.1`. According to this the `AOEROOT` variable must be set as follows:

```
AOEROOT=/dev/etherd/e0.1
```

KIWI is now able to mount and use the specified AoE device as the remote root filesystem. In case of a compressed read-only image with clicfs, the AOEROOT variable can also contain a device for the write actions:

```
AOEROOT=/dev/etherd/e0.1,/dev/ram1
```

Writing to RAM is the default but you also can set another device like another aoe location or a local device for writing the data

NFSROOT

Mount the system image root filesystem remotely via NFS (Network File System). This means there is a server which exports the root filesystem of the network client in such a way that the client can mount it read/write. In order to do that, the boot image must know the server IP address and the path name where the root directory exists on this server. The information must be provided as in the following example:

```
NFSROOT=NFS.Server.IP.address;/path/to/root/tree
```

KIWI_INITRD

Specifies the KIWI initrd to be used for a local boot of the system. The value must be set to the name of the initrd file which is used via PXE network boot. If the standard TFTP setup suggested with the kiwi-pxeboot package is used, all initrd files reside in the /srv/tftpboot/boot/ directory. Because the TFTP server does a chroot into the TFTP server path, you must specify the initrd file as follows:

```
KIWI_INITRD=/boot/name-of-initrd-file
```

KIWI_KERNEL

Specifies the kernel to be used for a local boot of the system. The same path rules as described for KIWI_INITRD applies for the kernel setup:

```
KIWI_KERNEL=/boot/name-of-kernel-file
```

ERROR_INTERRUPT

Specifies a message which is displayed during first deployment. Along with the message a shell is provided. This functionality should be used to send the user a message if it's clear the boot process will fail because the boot environment or something else influences the PXE boot process in a bad way.

13.4.3. User another than tftp as Download Protocol

By default all downloads controlled by the KIWI linuxrc code are performed by an atftp call and therefore uses the tftp protocol. With PXE the download protocol is fixed and thus you can't change the way how the kernel and the boot image (initrd) is downloaded. As soon as Linux takes over control the following download protocols http, https and ftp are supported too. KIWI makes use of the **curl** program to support the additional protocols.

In order to select one of the additional download protocols the following kernel parameters needs to be setup:

kiwiserver

Name or IP address of the server who implements the protocol

kiwiservertype

Name of the download protocol which could be one of http, https or ftp

To setup this parameters edit the file `/srv/tftpboot/pxelinux.cfg/default` on your PXE boot server and change the append line accordingly. Please note all downloads except for kernel and initrd are now controlled by the given server and protocol. You need to make sure that this server provides the same directory and file structure as initially provided by the kiwi-pxeboot package.

13.4.4. RAM Only Image

If there is no local storage and no remote root mount setup the image can be stored into the main memory of the client. Please be aware that there should be still enough RAM space available for the operating system after the image has been deployed into RAM. Below, find an example:

- Use a read-write filesystem in `config.xml`, for example `filesystem="ext3"`
- Create `config.MAC`

```
IMAGE='/dev/ram1;suse-##.##-pxe-client.i686;1.2.8;192.168.100.2;4096'
```

13.4.5. Union Image

As used in the `suse-pxe-client` example it is possible to make use of the `clifs` overlay filesystem to combine two filesystems into one. In case of thin clients there is often the need for a compressed filesystem due to space limitations. Unfortunately all common compressed filesystems provides only read-only access. Combining a read-only filesystem with a read-write filesystem is a solution for this problem. In order to use a compressed root filesystem based on `clifs` make sure your `config.xml`'s filesystem attribute contains `clifs`. As an alternative to `clifs` kiwi also supports the fuse based `unionfs` utility. In contrast to `clifs` which writes a block list on the write device, `unionfs` points all write operations into another filesystem which allows to mount and watch this location separately. In order to use a compressed root filesystem based on `unionfs` make sure your `config.xml`'s filesystem attribute contains `squashfs`. Below find examples for the different union modes.

13.4.5.1. Download to local storage, write to local storage

```
DISK=/dev/sda  
PART='5;S;x,400;L;/,x;L;x'  
IMAGE='/dev/sda2;suse-##.##-pxe-client.i386;1.2.8;192.168.100.2;4096'  
UNIONFS_CONFIG=/dev/sda3,/dev/sda2,unionfs  
KIWI_INITRD=/boot/initrd
```

13.4.5.2. Download to local storage, write to ram

```
DISK=/dev/sda  
PART='5;S;x,400;L;/'  
IMAGE='/dev/sda2;suse-##.##-pxe-client.i386;1.2.8;192.168.100.2;4096'  
UNIONFS_CONFIG=tmpfs,/dev/sda2,unionfs
```

13.4.5.3. Mount from remote, write to local storage

For all of the following modes I strongly recommend to check on a separate client machine in the network if it is possible to access the exported read-only and read-write device locations. If accessing devices works the image should also be able to access them on boot. If the boot fails it should be clear that the reason is not the exported device.

- NFSROOT

```
PART='5;S;x,x;L;x'  
NFSROOT="192.168.100.2;/srv/kiwi-read-only-path"  
UNIONFS_CONFIG=/dev/sda2,nfs,unionfs
```

- AOEROOT

```
PART='5;S;x,x;L;x'  
AOEROOT=/dev/etherd/e0.1,/dev/sda2  
UNIONFS_CONFIG=/dev/sda2,aoe,unionfs
```

- NBDROOT

```
PART='5;S;x,x;L;x'  
NBDROOT=192.168.100.7;2000;/dev/nbd0;;;/dev/sda2  
UNIONFS_CONFIG=/dev/sda2,nbd,unionfs
```

13.4.5.4. Mount from remote, write to ram

- NFSROOT

```
NFSROOT="192.168.100.2;/srv/kiwi-read-only-path"  
UNIONFS_CONFIG=tmpfs,nfs,unionfs
```

- AOEROOT

```
AOEROOT=/dev/etherd/e0.1  
UNIONFS_CONFIG=tmpfs,aoe,unionfs
```

- NBDROOT

```
NBDROOT=192.168.100.7;2000;/dev/nbd0  
UNIONFS_CONFIG=tmpfs,nbd,unionfs
```

13.4.5.5. Mount from remote, write to remote

- NFSROOT

```
NFSROOT="192.168.100.2;/srv/kiwi-read-only-path"  
UNIONFS_CONFIG=/srv/kiwi-read-write-path,nfs,unionfs
```

- AOEROOT

```
AOEROOT=/dev/etherd/e0.1,/dev/etherd/e1.1  
UNIONFS_CONFIG=aoe,aoe,unionfs
```

- NBDROOT

```
NBDROOT=192.168.100.7;2000;/dev/nbd0;3000;/dev/nbd1;4000;/dev/nbd2  
UNIONFS_CONFIG=nbd,nbd,unionfs
```

13.4.6. Split Image

As an alternative to the UNIONFS_CONFIG method it is also possible to create a split image and combine the two portions with the COMBINED_IMAGE method. This allows to use different filesystems without the need for an overlay filesystem to combine them together. Below find an example:

- Add a split type in config.xml, for example

```
<type fsreadonly="squashfs"
```

```
image="split" fsreadwrite="ext3" boot="netboot/suse-..."/>
```

- Add a split section inside the type to describe the temporary and persistent parts. For example:

```
<split>
  <temporary>
    <!-- allow RAM read/write access to: -->
    <file name="/mnt"/>
    <file name="/mnt/*"/>
  </temporary>
  <persistent>
    <!-- allow DISK read/write access to: -->
    <file name="/var"/>
    <file name="/var/*"/>
    <file name="/boot"/>
    <file name="/boot/*"/>
    <file name="/etc"/>
    <file name="/etc/*"/>
    <file name="/home"/>
    <file name="/home/*"/>
  </persistent>
</split>
```

- Sample config.MAC:

```
IMAGE='/dev/sda2;suse-##.#-pxe-client.i686;1.2.8;192.168.100.2;4096,\
/dev/sda3;suse-##.#-pxe-client-read-write.i686;1.2.8;192.168.100.2;4096'
PART='200;S;x,500;L;/,x;L'
DISK=/dev/sda
COMBINED_IMAGE=yes
KIWI_INITRD=/boot/initrd
```

13.4.7. Root Tree Over NFS

Instead of installing the image onto a local storage device of the client it is also possible to let the client mount the root tree via an NFS remote mount. Below find an example:

- Export the KIWI prepared tree via NFS.

- Sample config.MAC:

```
NFSROOT=192.168.100.7;/tmp/kiwi.nfsroot
```

13.4.8. Root Tree Over NBD

As an alternative for root over NFS it is also possible to let the client mount the root tree via a special network block device. Below find an example:

- Use nbd-server to export the KIWI prepared tree.

- Sample config.MAC

```
NBDR00T=192.168.100.7;2000;/dev/nbd0
```

13.4.9. Root Tree Over AoE

As an alternative for root over NBD it is also possible to let the client mount the root device via a special ATA over Ethernet network block device. Below find an example:

- Use the **vbladed** command to bind a block device to an ethernet interface. The block device can be a disk partition or a loop device (losetup) but not a directory like with NBD.
- Sample config.MAC:

```
A0ER00T=/dev/etherd/e0.1
```

This would require the following command to be called first:

```
vbladed 0 1 eth0 blockdevice
```

13.5. Hardware Grouping

While the PXE standard takes care of the ability to create hardware groups via hardware or IP address groups, it does not take into account groups for non-contiguous hardware or IP addresses. The PXE standard makes the assumption that each hardware group will be clearly delineated by a range of IP addresses, or the hardware is from the same vendor. While an ideal scenario, this may not be the case in an established, slightly dated installation where the hardware itself has out-lived the vendors that made them.

KIWI has the ability to create groups for non-contiguous configurations where different hardware types may be involved due to newer equipment being rotated into production or older hardware failing and replacements are from different vendors. In addition, an organization might decide to organize their equipment by function, rather than by vendor, and may not be able to use the same hardware from one end to the other.

13.5.1. The Group Configuration File

To make use of the grouping functionality, some new configuration files will be required. These configuration files currently have to be manually managed rather than provided, however future versions of KIWI may provide a means of managing groups more effectively once this feature stabilizes. The number of configuration files required will depend on the number of hardware groups that will be created, rather than one configuration file for each MAC address that will reside on the network.

There will be one configuration file that will always be required if using groups, called:

```
/srv/tftpboot/KIWI/config.group
```

This file has a new static element that must exist, and one or more dynamic elements depending on the number of groups that will be created. For example, the config.group file defined below lists 3 distinct groups:

```
KIWI_GROUP="test1, test2, test3"
test1_KIWI_MAC_LIST="11:11:11:11:11:11, 00:11:00:11:22:CA"
test2_KIWI_MAC_LIST="00:22:00:44:00:4D, 99:3F:21:A2:F4:32"
test3_KIWI_MAC_LIST="00:54:33:FA:44:33, 84:3D:45:2F:5F:33"
```

Note: The above hardware addresses contain random entries, and may not reflect actual hardware.

As we can see in the above example the file contains 1 static element, KIWI_GROUP, and 3 dynamic elements "test1_KIWI_MAC_LIST, test2_KIWI_MAC_LIST and test3_KIWI_MAC_LIST". The definitions of these elements are as follows:

- **KIWI_GROUP**

This element is the only static definition that needs to exist when using groups. While there is no implicit limit to the number of groups that can be configured, it should be kept to a minimum for reasonable management or it could quickly become un-manageable. It will need to contain one or more group names separated by comma's (,) and spacing (for readability). In the above example, our group names were:

- test1
- test2
- test3

Valid group names are made up of upper and lower case letters, and can use numeric, and underscore characters. The same rules used to define bash/sh variable names should apply here, as these names will have to be used as fully defined bash/sh variables when linking hardware addresses to an assigned group. The following is an example that contains valid names:

```
KIWI_GROUP="test1, test_my_name, LIST_HARDWARE, Multiple_Case_Group_1"
```

- **<GROUP_NAME>_KIWI_MAC_LIST**

The name of this element is dynamic and depends entirely on the list of group names that were previously defined. Each group name that was used in the KIWI_GROUP variable, must contain a matching dynamic element, and have KIWI_MAC_LIST appended to the name. To continue with our previous example, to create hardware lists for the groups already defined, we need 3 dynamic elements called:

- test1_KIWI_MAC_LIST
- test2_KIWI_MAC_LIST
- test3_KIWI_MAC_LIST

These variables will contain a comma delimited list of the hardware addresses for all of the machines being assigned to the appropriate group, but there are some caveats that need to be kept in mind. The first caveat is for hardware addresses that contain the HEX characters A-F. The PXE standard uses capital letters for these characters, and as a result KIWI does upper case comparisons, so a MAC address that is defined with lower case letters in this list will never get matched.

The second caveat is that as the list gets longer, it can be harder to maintain and it has the potential to slow down the booting process. However, testing has been completed with 1500+ hosts defined, and there was little delay when transferring the file to a single host. The file size will have a larger impact when trying to download it to 1500+ hosts, so some consideration will have to take that into account. The comparison itself still occurred in under half a second while searching through all 1500+ MAC addresses across 3 defined groups.

13.5.2. The Group Details File

In addition to the config.group file, each defined group will require a config.<GROUP_NAME> file. This file is exactly like a standard KIWI config.<MAC> file,

but is assigned to a group of hosts rather than a single unit. If we continue with the example we used in the previous section, we would need the following files:

```
/srv/tftpboot/KIWI/config.test1  
/srv/tftpboot/KIWI/config.test2  
/srv/tftpboot/KIWI/config.test3
```

The contents of these files is the same that would normally reside in a config.<MAC> file, and all definitions that would be supported for a single host, are supported for a group of hosts. In addition, if a host is matched to a group, yet the config.<GROUP_NAME> file does not exist, KIWI will error out.

For example, the following configuration file, called config.test1 would be used for the group called "test1":

```
DISK=/dev/sda  
PART='5;S;x,x;L;/'  
IMAGE='/dev/sda2;suse-##.#-pxe-client.i686;1.2.8;192.168.100.2;4096'  
CONF='CONFIGURATIONS/xorg.conf.test1;/etc/X11/xorg.conf;192.168.100.2;4096,\  
CONFIGURATIONS/syslog.conf;/etc/sysconfig/syslog.conf;192.168.100.2;4096'
```

As a result of this configuration file, the image would be configured consistently across all the hosts assigned to test1. The following file called config.test2, contains a small change that may be specific to a function:

```
DISK=/dev/sda  
PART='5;S;x,x;L;/'  
IMAGE='/dev/sda2;suse-##.#-pxe-client.i686;1.2.8;192.168.100.2;4096'  
CONF='CONFIGURATIONS/xorg.conf.test2;/etc/X11/xorg.conf;192.168.100.2;4096,\  
CONFIGURATIONS/syslog.conf;/etc/sysconfig/syslog.conf;192.168.100.2;4096'
```

As we can see, while group 1 and 2 share the syslog.conf configuration file, they have different xorg.conf files defined, therefore two distinct groups with one or more hosts assigned to each group can now be configured by managing a smaller number of files.

13.5.3. Using Hardware Mapping to Provide Overrides

The only issue with running mixed hardware configurations pertains primarily to hardware differences. For instance, it may be possible to create a single, xorg.conf file that is able to work with all of the hardware, but there is a chance it might not be possible to do so. With this in mind, KIWI provides a mechanism to provide "default" configurations that works with the most common hardware configuration, while providing hardware specific overrides to allow for any differences and yet have all hardware linked to the same group.

13.5.3.1. The Hardware Mapping Elements

To make use of the hardware linking mechanism, two additional parameters needs to be added to the group details file, the one named config.<group_name>. These two elements "link" hardware specific configurations to the appropriate systems. A general example would look like this:

```
HARDWARE_MAP="vendor_name_model"  
vendor_name_model_HARDWARE_MAP="00:00:00:11:11:11"
```

These parameters are not required, and the same functionality can be applied by using multiple groups to do the same thing, but that might not be desirable to some administrators.

This feature allows for a slightly more complex group to be defined, but the end result is a single group, that can contain multiple sub-groups ensuring flexibility in using a mixed set of hardware.

The definitions for the above parameters are as follows:

- `HARDWARE_MAP`

This element follows the same rules as defined by the `KIWI_GROUP` element. However, this variable will create sub-groups used to ensure multiple types of hardware vendors can be used within the same group. The name of the group(s) should be clearly defined, and a good convention to follow would be to use a combination of the vendor name with the model number or type. This would allow for cases where the same vendor is used, but differences between alternative models requires different maps to be used.

- `<HARDWARE_MAP_NAME>_HARDWARE_MAP`

This element behaves exactly like the `<GROUP_NAME>_KIWI_MAC_LIST` element defined above, in that it lists all MAC addresses that need to be linked to a hardware map. Any host defined within the list will receive configuration files that have been specifically defined in a `hardware_config.<hardware_map>` file, in addition to any files defined within a `CONF` element.

13.5.3.2. The Hardware Mapping Details File

Once the hardware map has been defined, the last step is to ensure configuration specific elements are linked to the host(s) in question. This is done by creating a new `hardware_config.<hardware_map>` file. The contents of the file is quite simple, and contains only one element called `VENDOR_CONF`, as the following example shows:

```
VENDOR_CONF='CONFIGURATIONS/xorg.conf.hardware_name_model;/etc/X11/xorg.conf;192.168.100.2;4096'
```

The format of the `VENDOR_CONF` values is exactly the same as the `CONF` variable used in the standard host and group configurations. In addition, files defined within this list will overwrite any files defined in the group configuration, if and only if, all of the following cases apply:

- The host is assigned to the current hardware map
- The file is defined within the `CONF` and `VENDOR_CONF` elements

NOTE: If a file is not defined in the `CONF` element, but is defined in the `VENDOR_CONF` element, it is simply downloaded to the host as if it was a `CONF` file. In this case, no overwriting will take place as it is considered a new file.

13.5.3.3. A Complete Example

The following is an example of a group that is using hardware from multiple vendors. For the purposes of this example, let's assume the group will have 10 defined hosts, seven are imaginative HP thinstations, while the remaining three are older Maxterm thinstations. We will also assume that the differences we are trying to address are specific to the video card and X.Org drivers used as a result.

With this in mind, we will need the following KIWI specific files:

```
cd /srv/tftpboot/KIWI
ls
  config.example1
  config.group
  hardware_config.maxterm_3500
```

As we can see, there is a KIWI group file, the group configuration or details file, and a new file that we have not seen before called `hardware_config.maxterm_3500`. We will first look at the contents of the `config.group` file:

```
cat config.group

KIWI_GROUP="example1"
example1_KIWI_MAC_LIST=
  "00:00:00:00:00:01 00:00:00:00:00:02 \
   00:00:00:00:00:03 00:00:00:00:00:04 \
   00:00:00:00:00:05 00:00:00:00:00:06 \
   00:00:00:00:00:07 00:00:00:00:00:08
   00:00:00:00:00:09 00:00:00:00:00:0A"
```

Within the file, there is a group called "example1", with ten hosts defined, in this case with imaginary sequential MAC addresses. Next, we look at the `config.example1` group details/configuration file:

```
cat config.example1

KIWI_INITRD=/boot/initrd
KIWI_KERNEL=/boot/linux
DISK=/dev/sda
PART='5;S;x,769;L;/,x;L;x'
IMAGE='/dev/sda2;exmaple-kiosk-opensuse-##.##-pxe-client.i686;0.0.1;192.168.1.2;4096'
UNIONFS_CONFIG=/dev/sda3,/dev/sda2,clifcs
CONF='prefs.js;/home/kioskuser/.mozilla/firefox/07xvllty.default/prefs.js;192.168.1.2;4096,xorg.conf;/'
RELOAD_IMAGE=yes
RELOAD_CONFIG=yes
HARDWARE_MAP='maxterm_3500'
maxterm_3500_HARDWARE_MAP='00:00:00:00:00:02 00:00:00:00:00:03 00:00:00:00:00:04'
```

Here, most of the standard KIWI configuration elements are in place, with a few extras. There are three areas we want to focus our attention on, the `CONF`, `HARDWARE_MAP` and `maxterm_3500_HARDWARE_MAP` variables, as they are the most critical elements to our example.

The first parameter to look at is the `CONF` parameter, which indicates a `prefs.js` (for Mozilla Firefox), and a `xorg.conf` (for X Windows) files will be copied to the host during boot up. These files should be considered defaults for the group, and all hosts defined in this group will use these files. As such, when the systems boot, both of these files will be copied over to their local file systems when the `CONF` element is processed.

Lastly, we have a hardware mapping group called "maxterm_3500", with three of the groups hosts defined as part of of a sub-group, or hardware map. The content of this file is as follows:

```
cat hardware_config.maxterm_3500

VENDOR_CONF='xorg.conf.maxterm_3500;/etc/X11/xorg.conf;192.168.1.2;4096,
             someconfig.cfg;/etc/sysconfig/someconfig.cfg;192.168.1.2;4096'
```

When the `VENDOR_CONF` definition is used, we are telling KIWI that all files defined within this element, are specific to the hardware map they are linked to. As a result, any files listed here will be transferred to a host if, and only if, the host has been linked to the hardware map via the `maxterm_3500_HARDWARE_MAP` element. In our example the only systems that

will receive the `xorg.conf.maxterm_3500` file will be the three maxterms we linked to the hardware map itself.

In our `VENDOR_CONF` element, we are indicating two files that should be transferred, in addition to any file transferred during the processing of the `CONF` element. A "specific" `xorg.conf` file, as well as `someconfig.cfg`. In the case of the `xorg.conf.maxterm_3500` file, when it is transferred to the host, it will overwrite the `xorg.conf` file that was previously transferred via the `CONF` element. However, with the `someconfig.cfg` file, because it was not previously defined in the `CONF` element, it will simply get transferred over, and is a perfect example of how one could enable functionality that is not otherwise configured.

As a result of this example, we have seven terminals that are using a `prefs.js` and generic `xorg.conf` file for their system configuration, and three terminals that are using `prefs.js`, a new version of the `xorg.conf` file as well as a file called `somconfig.cfg`. For the purposes of our example, the contents of the `prefs.js`, `xorg.conf`, `xorg.conf.maxterm_3500` and `someconfig.cfg` are arbitrary, and don't need to be explained here.

14 OEM Image—Preload Systems

Table of Contents

14.1. Building the suse-oem-preload Example	83
14.2. Using the Image	83
14.3. Flavours	84

An OEM image is a virtual disk image representing all partitions and bootloader information in the same fashion it exists on a physical disk. The image format matches the format of the VMX image type. All flavors discussed previously for the VMX image type apply to the OEM image type.

The basic idea behind an oem image is to provide the virtual disk data for OEM vendors to support easy deployment of the system to physical storage media. The deployment can be performed from any OS including Windows as long as a tool to dump data onto a disk device exists and is used. The oem image type may also be used to deploy an image on a USB stick. A USB stick is simply a removable physical storage device.

14.1. Building the suse-oem-preload Example

The OEM example provided with kiwi is based on recent openSUSE releases, one example configuration per release, and includes the default and x11 patterns. The image type is a split type utilizing the distributions default filesystem format for the read-write partition and the squashfs filesystem for the read-only partition. Using the additional `installiso` attribute creates an installable ISO image. When booting from the ISO image the OEM disk image will be deployed to the storage media on the booting machine (after confirmation by the user).

```
cd /usr/share/doc/packages/kiwi/examples
==> select the example directory for the desired distribution change into it
cd suse-...
kiwi --build ./suse-oem-preload -d /tmp/myoem-result --type split
```

14.2. Using the Image

The virtual disk image created by KIWI with the commands shown above can be tested using virtualization software such as QEMU, VMware, or VirtualBox. The virtual disk is represented by the file with the `.raw` extension, whereas the file with the `.iso` extension represents the

installation disk for this oem image. The ISO image is bootable (`filename.iso`) and can be burned to optical media. It is recommended to test the image on a bare test system. The following command shows how to use QEMU to test the OEM disk image (`filename.raw`).

```
cd /tmp/myoem-result
qemu suse-*-oem-preload.*.raw
```

or using the **dd** command you can dump the image onto a test hard disk or USB stick and upon reboot select the appropriate device as the boot device in the BIOS:

```
cd /tmp/myoem-result
dd if=suse-*-oem-preload.*.raw of=/dev/device bs=32k
```

Note, when testing an oem image using the virtual disk image, i.e. the `.raw` file, the geometry of the disk image is not changed and therefore retains the disk geometry of the host system. This implies that the re-partitioning performed for a physical disk install during the oem boot workflow will be skipped.

You can test the installation procedure in a virtual environment using the `.iso` file. In this case the re-partitioning code in the boot image will be executed. The following commands show this procedure using QEMU.

```
cd /tmp/myoem-result
qemu-img create /tmp/mydisk 20G
qemu -hda /tmp/mydisk -cdrom suse-*-oem-preload.*.iso -boot d
```

14.3. Flavours

As indicated above the use of the `installiso` and `installstick` attributes for the oem image supports the creation of an installation image. The installation image can be created in two formats, one suitable for CD/DVD media and a second suitable for a USB stick. The self installing image deploys the oem image onto the selected storage device. The installation process is a simple image dump using the **dd** command. During this process the target system remains in terminal mode. The following configuration snippets show the use of the `installiso` and `installstick` attributes to create the ISO or USB installation image format respectively.

- `<type image="name" ... installiso="true"/>`

Creates a `.iso` file which can be burned onto a CD or a DVD. This represents an installation CD/DVD

- `<type image="name" ... installstick="true"/>`

Creates a `.raw.install` file which can be dumped (**dd**) onto a USB stick. This represents an installation Stick

14.3.1. Specializing the OEM install process

It is possible to specialize the OEM install process by providing shell scripts with the following names. For more information how to pack the scripts and make them work in the boot code, see the chapter Section 3.3, “Boot image Hook-Scripts”.

- `preHWdetect.sh` This script is executed prior to the hardware scan on the target machine.

- `preImageDump.sh` This script is executed immediately prior to the OEM image dump onto the target storage device.
- `postImageDump.sh` This script is executed directly after the OEM image dump onto the target storage device once the image checksum has been successfully verified.

14.3.2. Influencing the OEM Partitioning

By default the oemboot process will create/modify a swap, /home and / partition. It is possible to influence the behavior with the `oem-*` elements. See Chapter 5, *KIWI Image Description* for details.

14.3.3. LVM Support

KIWI also provides support for LVM (Logical Volume Management). In this mode the disk partition table will include one lvm partition and one standard ext2 boot partition. KIWI creates the `kiwiVG` volume group, unless the `lvmgroup` attribute has been set, and adds logical volumes to the group based on the configuration given by the `systemdisk` block for this type. The filesystem for the volume group is determined by the `filesystem` attribute of the type element. After booting the system the user has full control over the volume group and is free to change (resize/increase) the group and the volumes inside. Support for LVM has been added for all disk based image types. This includes the `vmx` and `oem` image types. In order to use LVM the existence of a `systemdisk` section is required. The `systemdisk` specification may be empty. An empty `systemdisk` specification triggers the creation of one LVM root volume with the default `kiwiVG` name.

```
kiwi --create /tmp/myoem --type oem -d /tmp/myoem-result --lvm
```

With the `systemdisk` section you can specify to have one or more top level directories in a separate volume. See Chapter 5, *KIWI Image Description* for a detailed explanation.

14.3.4. Partition Based Installation

The default installation method of an OEM is dumping the entire virtual disk on the selected target disk and repartition the disk to the real geometry. This works but will also wipe everything which was on the disk before. KIWI also supports the installation into already existing partitions. This means the user can setup a disk with free partitions for the KIWI OEM installation process. This way already existing data will not be touched. In order to activate the partition based install mode the following `oem` option has to be set in `config.xml`:

```
<oem-partition-install>true</oem-partition-install>
```

Compared to the disk based install the following differences should be mentioned:

- The bootloader will be setup to boot the installed system. There is no multiboot setup. The user is expected to implement the setup of a multiboot bootloader.
- The `oem` options for `system`, `swap` and `home` doesn't have any effect if the installation happens in predefined partitions.
- There is no support for remote (PXE) OEM installation because `kiwi` has to loop mount the disk image in order to access the partitions which can't be done remotely.
- The raw disk image is stored uncompressed on the install media. This is because KIWI needs to loop mount the disk image which it can't do if the file is only available as compressed

version. This means the install media in this mode will be approximately double the size of a standard install media.

14.3.5. Network Based Installation

Instead of manually dumping the OEM image on the target device or creating a kiwi install CD, USB stick, there is a third method of deploying the OEM image on the target device. It's possible to let the image be downloaded from a PXE boot server over the network. This requires a PXE network boot server to be setup properly in the first place. For details how to do this refer to the chapter: Chapter 13, *PXE Image—Thin Clients*. If your pxe server is running the following steps are required to setup the install process over the network

- Make sure you have created an install PXE tarball along with your oem image:

```
<type image="oem" ... installpxe="true"/>
```

- unpack the created `<image-name>.tgz` file and copy the `initrd` and `kernel` images over to your pxe server

```
tar -xf <image-name>.tgz
scp initrd-oemboot-*.install.* pxe.server.ip:/srv/tftpboot/boot/initrd
scp initrd-oemboot-*.kernel.* pxe.server.ip:/srv/tftpboot/boot/linux
```

- Next copy the system image and md5 sum over to the pxe boot server

```
scp <image-file>.gz pxe.server.ip:/srv/tftpboot/image/
scp <image-file>.md5 pxe.server.ip:/srv/tftpboot/image/
```

- At last set the kernel commandline parameters to the append line in your PXE configuration (for example: `pxelinux.cfg/default`). The required parameters are stored in the file `<image-file>.append` from the kiwi generated install tarball

Optionally the image can be stored on a FTP,HTTP server specified via the **kiwiserver** and **kiwiservertype** append information. In this case make sure you copied the system image and md5 file to the correct location on the ftp,http,etc. server. kiwi searches the image at one place only which is below the `image/` directory on the root path of the specified server. `initrd` and `linux` kernel are loaded by pxe thus they require a tftp server to be present.

15 Xen Image—Paravirtual Systems

Table of Contents

15.1. Building the suse-xen-guest Example	87
15.2. Using the Image	87
15.3. Flavours	88

Xen is a free software virtual machine monitor. It allows several guest operating systems to be executed on the same computer hardware at the same time.

A Xen system is structured with the Xen hypervisor as the lowest and most privileged layer. Above this layer are one or more guest operating systems, which the hypervisor schedules across the physical CPUs. The first guest operating system, called in Xen terminology “domain 0” (dom0), is booted automatically when the hypervisor boots and given special management privileges and direct access to the physical hardware. The system administrator logs into dom0 in order to start any further guest operating systems, called “domain 0” (domU) in Xen terminology.

A Xen image is a virtual disk like a vmx but with the xen kernel installed. In order to run it a Xen dom0 server needs to run. Xen images in KIWI makes use of the PVGrub method supported by current Xen versions. Xen extracts the kernel and initrd from the virtual disk as well as the grub configuration and displays the menu which allows emulation of the Grub console

15.1. Building the suse-xen-guest Example

The latest example provided with KIWI is based on openSUSE and includes the base pattern.

```
cd /usr/share/doc/packages/kiwi/examples cd suse-...  
kiwi --prepare ./suse-xen-guest --root /tmp/myxen
```

```
kiwi --create /tmp/myxen --type vmx -d /tmp/myxen-result
```

15.2. Using the Image

In order to run a domain U the Xen tool **xm** needs to be called in conjunction with the KIWI generated domain U configuration file

```
xm create -c /tmp/myxen-result/  
the-file-with-suffix.xenconfig
```

15.3. Flavours

With KIWI you can provide the information required to create a guest configuration as part of the `config.xml` file. Additionally you can group special packages which you may only need in this para virtual environment with a profile.

```
<packages type="image" profiles="xenFlavour">
  <package name="kernel-xen" replaces="kernel-ec2"/>
</packages>
<type ....>
  <machine memory="512" domain="domU">
    <vmdisk ... device="/dev/xvda"/>
  </machine>
</type>
```

If this information is present KIWI will create a Xen domain U configuration with 512 MB of RAM and expects the disk at `/dev/xvda`. Additional information to setup the Xen guest machine properties are explained in the machine section. The KIWI Xen domain U configuration is stored in the file `/tmp/myxen-result/suse-##.#-xen-guest.####-#.#.#.xenconfig`.

16 EC2 Image—Amazon Elastic Compute Cloud

Table of Contents

16.1. Building the suse-ec2-guest Example	90
16.2. Using EC2 and the created image	91

The Amazon Elastic Compute Cloud™ (Amazon EC2) provides an environment known as *IAAS* [<http://en.wikipedia.org/wiki/IAAS>] (Infrastructure as a Service). In this environment you have the ability to run Virtual Machines (VMs) on hardware managed by Amazon and the virtualization infrastructure provided by Amazon.

The virtualization infrastructure for EC2 is setup to work with Amazon Machine Images (AMIs). There are two storage models for AMIs:

1. S3 [<http://aws.amazon.com/s3/>] (Simple Storage Service) backed AMI
2. EBS [<http://aws.amazon.com/ebs/>] (Elastic Block Store) backed AMI

The image created with KIWI can be used to create an AMI for both storage models. For an S3 backed AMI a bundle with a manifest XML file is required and KIWI uses the Amazon tools to create this bundle for you. With the bundle created all you have to do is upload the bundle using the EC2 tools, register a new AMI and run the registered AMI.

For an EBS backed AMI the procedure to get to a working AMI requires more manual steps when compared to the S3 backed AMI approach. KIWI does not provide any automation for this process at this time. The KIWI created image needs to be uploaded to EC2 and then it needs to be dumped to an EBS volume. This implies that you need to have a running AMI in EC2

The procedures to handle both storage options are outlined below.

You can work with EC2 using the Amazon Web application found at <http://aws.amazon.com> or you can use the Amazon provided command line tools. In this example we will exclusively interact with EC2 using the command line tools. The command line tools are divided into *AMI* and *API* tools. The AMI tools are designed to operate on images, while the API tools are designed to work with the Amazon *REST API* [http://en.wikipedia.org/wiki/Representational_state_transfer]. In order for KIWI to create the bundle for S3 backed AMIs the Amazon AMI tools must be installed. It is recommended that you install both, the AMI and API tools on your build system.

The Amazon tools are not distributed with KIWI and can be installed using packages from the openSUSE Build Service *Cloud:EC2* [<http://download.opensuse.org/repositories/Cloud/>]

EC2/] repository, or can be downloaded from Amazon at <http://aws.amazon.com/developer-tools/368> and <http://aws.amazon.com/developertools/351>.

Documentation for Amazon EC2 can be found at <http://aws.amazon.com/documentation/ec2/>. The documentation for the command line tools may be accessed at <http://docs.amazonwebservices.com/AWSEC2/latest/CommandLineReference>. All commands also support the customary `--help` command line option to display the supported command line arguments for the given command.

When working with the Amazon tools it is useful to set the `EC2_HOME`, `EC2_PRIVATE_KEY`, and `EC2_CERT` environment variables. Setting `EC2_PRIVATE_KEY`, and `EC2_CERT` allows you to forego specification of the `--private-key` and `--cert` command line options with every command. The `EC2_HOME` environment variable is used by the tools to find required libraries.

- `EC2_HOME`

Location of the bin and lib directories installed by the Amazon tools. A good location for the tools on your system is `/usr/local`.

- `EC2_PRIVATE_KEY`

Path to your private key file (including the filename). For example `/home/USERNAME/AWS/keys/pk-....pem`

- `EC2_CERT`

Path to your certificate file (including the filename). For example `/home/USERNAME/AWS/keys/cert-....pem`

Please note that your account will be billed by Amazon at the published rate for any computing resources you consume in EC2. This includes but is not limited to, running instances, storing data (your image) on S3 or EBS, and network traffic.

One final remark before we get started, the default region for any **ec2-** command that communicates with the REST API or sends files to EC2 is the US-East region, i.e. `us-east-1`. Therefore, if you want to upload any data to other EC2 regions you must specify the desired target region. Specifying a region is accomplished by setting the `EC2_URL` environment variable, by using the `--url` command line option, or by using the `--region` argument. The `EC2_URL` environment variable and the `--url` argument expect a value in the form `https://ec2.amazonaws.com(us-east-1)`. The `--region` argument expects the name of a region as returned by the **ec2-describe-regions** command. The region specified for the **ec2-** commands must match the specified in the `config.xml` if this option was specified in your configuration.

16.1. Building the suse-ec2-guest Example

The example provided with KIWI uses openSUSE as the base distribution and includes the base pattern plus the vim editor. Also included is the `suse-ami-tools` package that provides tools needed in the EC2 environment.

Prior to building the EC2 example verify that the EC2 command line tools are installed and work as expected by executing the

```
ec2-describe-images -a
```

command. The result is a list of all available images in EC2 printed to your terminal. Additionally you need to edit the `config.xml` file of the example and enter your EC2 account information. Replace the values for `ec2accountnr`, `ec2privatekeyfile`, and `ec2certfile`. Note that the account number is displayed with dashes ("-") on the Amazon web site, but the entry in `config.xml` is expected to be numeric only. The private key file is the private file you downloaded when you created your certificate. It is recommended to create a copy of the example `/usr/share/doc/packages/kiwi/examples/suse-Version/suse-ec2-guest`.

You are now ready to build the EC2 example. Lets assume you copied the example configuration directory to `/tmp` prior to modifying the `config.xml` file.

```
kiwi --prepare /tmp/suse-ec2-guest --root /tmp/myec2
```

```
kiwi --create /tmp/myec2 -d /tmp/myec2-result -y
```

16.2. Using EC2 and the created image

Upon completion of the image build all files necessary for an S3 backed AMI and for an EBS backed AMI can be found in your destination directory, `/tmp/myec2-result` if you followed the commands given above. Prior to describing the specifics about using the KIWI produced images the following section will address some rudimentary general EC2 concepts and commands that can be used with existing AMIs and the AMIs you can register with the KIWI created images.

16.2.1. Using a registered AMI

This section is not a replacement for the EC2 documentation mentioned earlier. We will only cover the concepts and commands necessary to get you started such that you can launch the KIWI created image in this example.

Prior to launching any instance in EC2 you need to have a keypair. If you do not already have a keypair in EC2 you can create one using the **ec2-add-keypair** command. This creates a public/private keypair that is used to grant you access to your running instance via the ssh tools. Generate the keypair as shown below, the `gsgkey` name is arbitrary and used in this example, you can choose any name you like. The use of the key is quite frequent. Therefore, you probably want to choose a name that is easy to remember and not too terribly long to type.

```
ec2-add-keypair gsgkey
```

Save the private key returned by the command in a local file. Using your favorite text editor, paste everything between (and including) the `-----BEGIN RSA PRIVATE KEY-----` and `-----END RSA PRIVATE KEY-----` lines into your editor and save the key to a file. The file can have any name. However, it makes sense to name the file after the keypair name you have chosen earlier. If the file is named differently from the keypair you will end up launching instances with `-k mykey` and then accessing the instance with `-i yourkey`, which may be a bit weird. As indicated by the heading, this is your private key, thus make sure you safe guard it appropriately. On Linux the ssh tools will complain if the key file does not have the proper permissions. Change the permissions of your private key file to be read-write by you, the owner, only.

```
chown 600 gsgkey
```

The public key of your keypair is stored in the EC2 infrastructure. EC2 allows you to have multiple keypairs, to review your existing keypairs use the **ec2-describe-keypairs** command.

When you launch an instance of an AMI you must specify a keypair name. This selects the public key to be injected into the instance. The key injection occurs through the `amazon init` script provided by the `suse-ami-tools` package. This package, as mentioned previously, is already included in the example's `config.xml` file. Do not forget to include this package when you create your own image descriptions for EC2 or you will not be able to log into your running instances. Additionally you need to activate this service by adding `suseInsertService amazon` in your `config.sh` file.

The key injection mechanism needs to access the network. Therefore, you must configure the network when you build your image. Configuration of the network can be accomplished through the overlay mechanism or via commands in `config.sh`. The network interface of a guest in EC2 is always `eth0` and it needs to be configured to use DHCP. In the example the overlay mechanism is used to setup the network configuration.

Another prerequisite to launching an instance in EC2 is knowing the AMI you want to instantiate. As you have already seen, the **ec2-describe-images** command with the `-a` option will provide information about all publicly available AMIs. Using the command without the `-a` option will provide information about your registered images.

The Amazon EC2 infrastructure uses *PVGrub (Para-Virtual Grub)* [<http://www.linode.com/wiki/index.php/PV-GRUB>] to boot instances of an AMI. This allows instances to run the kernel that is part of the AMI, rather than some kernel provided by the Amazon infrastructure. However, an Amazon provided kernel is still required to kick things off and in the startup process PVGrub eventually picks up the `/boot/grub/menu.lst` file in your image and then boots the kernel specified. Note, that during the boot process you do not have access to a console and thus it makes no sense to have multiple kernel entries in your `menu.lst` file. Without console access you do not have an opportunity to choose a kernel. Each EC2 region has it's own independent copy of this boot mechanism and the boot mechanism is differentiated between 32 bit and 64 bit. The boot kernels are named with an ID that starts with the TLA (Three Letter Acronym) *aki* followed by a dash ("-") and a hex number. The Amazon Kernel Image IDs table below provides guidelines for the selection of the boot kernel ID based on Region and image architecture.

Table 16.1. Amazon Kernel Image IDs

Region	AKI	Arch	Name
AP-Northeast	aki-ec5df7ed	x86	ec2-public-images-ap-northeast-1/pv-grub-hd0-V1.02-i386.gz.manifest.xml
AP-Northeast	aki-ee5df7ef	x86-64	ec2-public-images-ap-northeast-1/pv-grub-hd0-V1.02-x86_64.gz.manifest.xml
AP-Southeast	aki-a4225af6	x86	ec2-public-images-ap-southeast-1/pv-grub-hd0-V1.02-i386.gz.manifest.xml
AP-Southeast	aki-aa225af8	x86-64	ec2-public-images-ap-southeast-1/pv-grub-hd0-V1.02-x86_64.gz.manifest.xml
EU-West	aki-64695810	x86	ec2-public-images-eu/pv-grub-hd0-V1.02-i386.gz.manifest.xml
EU-West	aki-62695810	x86-64	ec2-public-images-eu/pv-grub-hd0-V1.02-x86_64.gz.manifest.xml

Region	AKI	Arch	Name
SA-East	aki-bc3ce3a1	x86	ec2-public-images-sa/pv-grub-hd0-V1.02-i386.gz.manifest.xml
SA-East	aki-cc3ce3d1	x86-64	ec2-public-images-sa/pv-grub-hd0-V1.02-x86_64.gz.manifest.xml
US-East	aki-805ea7e9	x86	ec2-public-images/pv-grub-hd0-V1.02-i386.gz.manifest.xml
US-East	aki-825ea7eb	x86-64	ec2-public-images/pv-grub-hd0-V1.02-x86_64.gz.manifest.xml
US-West	aki-83396bc6	x86	ec2-public-images-us-west-1/pv-grub-hd0-V1.02-x86_64.gz.manifest.xml
US-West	aki-8d396bc8	x86-64	ec2-public-images-us-west-1/pv-grub-hd0-V1.02-x86_64.gz.manifest.xml
US-West2	aki-c2e26ff2	x86	ec2-public-images-us-west-2/pv-grub-hd0-V1.02-x86_64.gz.manifest.xml
US-West2	aki-98e26fa8	x86-64	ec2-public-images-us-west-2/pv-grub-hd0-V1.02-x86_64.gz.manifest.xml

The information in the table above was extracted from the Amazon documentation found at: http://ec2-downloads.s3.amazonaws.com/user_specified_kernels.pdf.

KIWI images created for EC2 do not contain a partition table. Therefore, PVGrub images designated by *h0* in the manifest name need to be used with the images created by Kiwi.

Specifying the optional `ec2region` element as part of the `ec2config` block in your `config.xml` file instructs KIWI to embed the appropriate PVGrub Kernel ID in the EC2 manifest file of the bundle. Therefore, you can create region specific bundles. KIWI will create a bundle for each specified region. Supported values for the `ec2region` element are *AP-Northeast*, *AP-Southeast*, *EU-West*, *US-East*, *US-West*, and *US-West2*.

Public images in EC2 already have the boot kernel information embedded. If you build a generic image with EC2 by not specifying the EC2 region you will need to specify the appropriate boot kernel on the command line when you launch an instance of your AMI as shown below. The example uses the x86 Kernel in the US-East region.

```
ec2-run-instances ami-... --kernel aki-407d9529 -k gsgkey
```

To check the state of your instance(s) use the command **ec2-describe-instances**.

Once the instance state for an instance is shown as *running* you can log into the instance using `ssh` as follows:

```
ssh -i PATH_TO_PRIVATE_KEY root@PUBLIC_IP_OF_YOUR_INSTANCE
```

The *PUBLIC_IP_OF_YOUR_INSTANCE* is displayed as part of the output of the **ec2-describe-instances** command. If you are unable to log in, it is most likely that the security setting for the instance is blocking the `ssh` access, or that you did not enable the `ssh` daemon process on startup. Your `config.sh` script should have a line activating the `ssh` daemon as follows, *suseInsertService sshd*. If you do not have this entry in `config.sh` you will have to rebuild your image. Addressing the problem with port blocking is accomplished using the **ec2-authorize** command to open port 22 as shown below.

```
ec2-authorize default -p 22
```

16.2.2. Using the bundle for an S3 backed AMI

The destination directory, `/tmp/myec2-result` if you followed the commands given above contains a directory ending with `-any`. This directory contains the AMI bundle created by KIWI with the Amazon tools for use with the S3 storage model. The `-any` ending of the directory name indicates that this bundle may be deployed to any EC2 region. As described earlier you can use the `ec2region` element in `config.xml` to build bundles for specific EC2 regions. If the `ec2region` element is used the directory name containing the bundle will end with the appropriate region string.

The generated bundle needs to be transferred to Amazon using the **ec2-upload-bundle** command line tool. This tool is part of the Amazon API tools. Upload the AMI as follows, replacing `AWS_Key_ID` and `AWS_secret_Key_ID` with your Amazon key information. Also you may want to choose a different name for your bucket than `myImages`. If the bucket does not exist in S3 it will be created.

```
ec2-upload-bundle -b myImages -a AWS_Key_ID -s AWS_secret_Key_ID \
-m /tmp/myec2/suse-##.##-ec2-guest.ARCH-Version.ec2-any/ \
  suse-##.##-ec2-guest.ARCH-Version.ami.manifest.xml
```

Where *ARCH* indicates the architecture to be used for this image, 32 bit (i686) or 64 bit (x86_64), and *Version* matches the value of the version element in the `config.xml` file.

After the upload process is complete, register your image with the EC2 infrastructure using the **ec2-register** command as shown below. The result of the registration process is an AMI ID returned on the command line in the form “ami-” followed by a random key sequence. Use this AMI ID to launch your instance as described in the Using a registered AMI section.

```
ec2-register myImages/suse-##.##-ec2-guest.ARCH-Version.ami.manifest.xml
```

This completes the S3 specific setup using the bundle created with KIWI. The next section explains the use of the disk image file created by KIWI to create and EBS backed AMI.

16.2.3. Using the disk image for and EBS backed AMI

Inside the destination directory, `/tmp/myec2-result` if you followed the commands given above, you will find a file with no extension. This file is the filesystem image KIWI created and used to create the AMI bundle suitable for use with S3. This filesystem image is the one we need to install/dump onto an EBS volume to have an EBS backed AMI. In order to create the EBS backed AMI it is necessary to work on a running AMI and the local machine that contains the disk image file. Further, the running instance must have the **dd_rescue** command available. Using the “regular” **dd** command will result in failure as it does not properly handle sparse images. The SUSE Linux Enterprise AMIs available in EC2 have the **dd_rescue** command available. For the openSUSE AMIs available in EC2 it is easy to install the **dd_rescue** command by executing `zypper` as shown below.

```
zypper in dd_rescue
```

The first step in creating an EBS backed AMI is to create a tarball of the filesystem image file. This is necessary to preserve the sparseness properties of the image. Use the `-S` option as shown below. The following commands are executed on your build machine.

```
cd /tmp/myec2-result
```

```
tar -cSf myImage.tar suse-##.##-ec2-guest.ARCH-Version
```

The tarball needs to be uploaded to EC2 and unpacked. This implies that the running instance of your chosen AMI needs to have sufficient space to store the tarball and the unpacked tarball. Therefore, it is recommended to create a storage volume as shown below. The command used will return a volume ID that you want to remember and that will be referred to as *STORE_VOL_ID* in this example. Execute this command on your local machine.

```
ec2-create-volume -s X -z AV_ZONE
```

The *X* is an integer value representing the size of the volume to be created in GB (Giga Bytes). The *AV_ZONE* value is one of the Amazon availability zones. For all commands presented here the *AV_ZONE* value must be the same. Possible values for *AV_ZONE* are obtained with the **ec2-describe-availability-zones**

The next step is to launch an instance of your liking in EC2. This instance will be used to create the EBS volume that will eventually function as the backing store for your AMI, therefore this instance must have the **dd_rescue** as described earlier. Launching the instance is accomplished with the **ec2-run-instances** command executed on your local machine.

```
ec2-run-instances -k gsgkey -n 1 -g SECURITY_GROUP_NAME -t INSTANCE_TYPE -z AV_ZONE ami-...
```

The *SECURITY_GROUP_NAME* is optional, but it is best to have a group setup that allows ssh access by default to avoid having to open the port all the time as outlined earlier. The *INSTANCE_TYPE* specified with the *-t* depends on the architecture of your image, for example a 64 bit AMI cannot run in a small (m1.small) instance type. The **ec2-run-instances** command returns the instance ID, and you want to remember this ID. This ID will be referred to as *INST_ID* in this example.

Wait until the instance is running, check the status with the **ec2-describe-instances** command. Once the instance is indicated as *running* attach the previously created volume to your instance, by executing the **ec2-attach-volume** command on your local machine.

```
ec2-attach-volume STORE_VOL_ID -i INST_ID -d /dev/sdf
```

The chosen device specified with the *-d* is arbitrary, however, you do obviously not want to pick a device node that is already in use.

In a different shell login to the running instance as shown in the Using a registered AMI section.

Check that the storage volume is attached using the **ec2-describe-volumes** command on your local machine. Once the volume is attached the status will be shown as *in use*.

With the EBS storage volume attached to your instance create a filesystem on the volume. There is no need to partition the volume. In the EC2 instance that you previously logged into execute the **mkfs** command.

```
mkfs -t ext3 /dev/sdf
```

Once the filesystem creation is complete mount the volume.

```
mount dev/sdf /mnt
```

With the storage volume attached, formatted and mounted you now have sufficient space to transfer the image tarball you created earlier. From the destination directory on your local machine that contains your tarball initiate an ftp session to your EC2 instance.

```
sftp -i PATH_TO_PRIVATE_KEY root@PUBLIC_IP_OF_YOUR_INSTANCE
```

```
sftp>cd /mnt
```

```
sftp>put myImage.tar
```

```
sftp>bye
```

While the image is transferring you can create a new EC2 volume. This new volume will become the backing store for your new AMI and be referred to as *VOL_ID* in this example. Use the **ec2-create-volume** command on your local machine as shown previously. This time you want to make sure that the specified size matches the size of your filesystem image file.

```
ec2-create-volume -s Y -z AV_ZONE
```

It is good practice to use the size element in your `config.xml` file when creating EC2 images. Using the size element ensures that you have additional space on your root volume and that you can match the volume size exactly to your image size.

As previously, wait for the volume to be created and check the status with the **ec2-describe-volumes** command. When the volume creation process is complete, attach the new volume to your running instance by executing the **ec2-attach-volume** command on your local machine.

```
ec2-attach-volume VOL_ID -i INST_ID -d /dev/sdg
```

Wait until the status for the volume changes to *in use* before proceeding.

By now the upload of your image tarball has probably completed and you can unpack the image on your storage volume in your EC2 instance. In the shell on your EC2 instance unpack the tarball as shown below.

```
cd /mnt
```

```
tar -xSf myImage.tar
```

After the unpack operation completes you can dump the filesystem image to your EBS volume that will be your backing store for your AMI, this volume if attached to `/dev/sdg` in this example. In the shell on your EC2 instance execute the **dd_rescue** command as shown below.

```
dd_rescue -a /mnt/suse-##.##-ec2-guest.ARCH-Version /dev/sdg
```

The `-a` option instructs `dd_rescue` to handle the sparse file properly. You have to wait until the dump process is complete before proceeding to the next step.

With the dump process complete you can now unmount the storage volume from your EC2 instance by executing the following in the shell of your running instance.

```
umount /mnt
```

You can also exit the shell in your running EC2 instance as all remaining commands are executed on your local machine.

You must detach the volume that is intended as your AMI backing store from your running instance. Detaching and deleting the storage volume as well as terminating the running instance are optional.

```
ec2-detach-volume -i INST_ID -d /dev/sdg VOL_ID
```

```
ec2-detach-volume -i INST_ID -d /dev/sdf STORE_VOL_ID
```

Prior to shutting down the instance or deleting the storage volume, wait until the detach operation has completed. This is indicated by the *available* status in the output of the **ec2-describe-volumes** command.

```
ec2-delete-volume STORE_VOL_ID
```

```
ec2-terminate-instances INST_ID
```

Your next step is to create a snapshot of the EBS volume that contains your image. The **ec2-create-snapshot** command returns a snapshot ID that you want to remember and will be referred to as *SNAP_ID* in this example.

```
ec2-create-snapshot -d A_SHORT_DESCRIPTION VOL_ID
```

The process of creating the snapshot will take a while and depends on the size of your volume. Check the status of the snapshot creation process using the **ec2-describe-snapshots** command. When the process is complete the status will change to *100%*.

Once the snapshot is complete you can register it as an AMI with the EC2 infrastructure using the **ec2-register** command.

```
ec2-register -d A_DESCRIPTION -n A_NAME -s SNAP_ID -a ARCH --kernel BOOT_KERNEL_ID
```

The **ec2-register** command will return the AMI ID that you can then use to launch your instance.

The process of creating an EBS backed AMI is a bit tedious. If you create EBS backed AMIs more often it might be well worth your time to script this process using the Amazon REST API. The EC2 tools provided by Amazon that would by name imply to implement this process (**ec2-import-instance** and **ec2-upload-disk-image**) do not handle Linux based disk images at the time of this writing.

A KIWI Man Pages

Table of Contents

kiwi	100
kiwi::config.sh	107
kiwi::images.sh	111
kiwi::kiwirc	114

The following pages will show you the man page of KIWI and the functions which can be used within **config.sh** and **index.sh**

kiwi

kiwi — Creating Operating System Images

Synopsis

```
kiwi { -l | --list }
```

```
kiwi { -o | --clone } image-path { -d } destination
```

```
kiwi { -b | --build } image-path { -d } destination
```

Basics

KIWI is a complete imaging solution that is based on an image description. Such a description is represented by a directory which includes at least one `config.xml` file and may as well include other files like scripts or configuration data. The `kiwi-templates` package provides example descriptions based on a JeOS system. JeOS means *Just enough Operating System*. KIWI provides image templates based on that axiom which means a JeOS is a small, text only based image including a predefined remote source setup to allow installation of missing software components at a later point in time.

Detailed description of the kiwi image system exists in the system design document in file:///usr/share/doc/packages/kiwi/kiwi.pdf. KIWI always operates in two steps. The KIWI `--build` option just combines both steps into one to make it easier to start with KIWI. The first step is the preparation step and if that step was successful, a creation step follows which is able to create different image output types. If you have started with an example and want to add you own changes it might be a good idea to clone of from this example. This can be done by simply copying the entire image description or you can let KIWI do that for you by using the `kiwi --clone` command.

In the preparation step, you prepare a directory including the contents of your new filesystem based on one or more software package source(s) The creation step is based on the result of the preparation step and uses the contents of the new image root tree to create the output image. If the image type ISO was requested, the output image would be a file with the suffix `.iso` representing a live system on CD or DVD. Other than that KIWI is able to create images for virtual and para-virtual (Xen) environments as well as for USB stick, PXE network clients and OEM customized Linux systems.

Image Preparation and Creation

```
kiwi { -p | --prepare } image-path  
[ -r | --root image-root | --cache directory]
```

```
kiwi { -c | --create } image-root  
{ -d | --destdir destination } [--type image-type]
```

Image Upgrade

If the image root tree is stored and not removed, it can be used for upgrading the image according to the changes made in the repositories used for this image. If a distributor provides an update channel for package updates and an image `config.xml` includes this update channel

as repository, it is useful to store the image root tree and upgrade the tree according to changes on the update channel. Given that the root tree exists it's also possible to add or remove software and recreate the image of the desired type.

```
kiwi { -u | --upgrade } image-root [--add-packagename] [--add-patternname]
```

System to Image Migration

The migration module allows you to migrate your currently running system into an image description. The module will check for files not managed by a package manager and also inspects your system for package pattern and file consistency according to the currently active repositories. The system requires the zypper backend in order to work properly.

The migration process creates a cache file so that subsequent calls of the migration runs much faster. Please have in mind that if your system has changed (files created/deleted, etc.) the cache file might not be worth to become reused. In this case you should remove the cache first and start from scratch. The option `--nofiles` will prevent the system from searching for unpackaged and packaged but modified files. The option `--notemplate` will prevent the creation of the image description files which are needed if you want to use KIWI to create a clone image from the result of the migration. With the options `--exclude` and `--skip` you can tell the system to ignore specific directories and/or packages. This makes sense if you know before that some data is not worth to become migrated or can be restored easily later inside the cloned image like software repositories.

The migration process will always place it's result into the `/tmp/$OptionValueOf-m` directory. The reason for this is because `/tmp` is always excluded from the migration operation and therefore we can safely place new files there without influencing the migration itself. You should have at least 50 MB free space for the cache file and the image description all the rest are just hard links.

As one result a HTML based report file is created which contains important information about the system. You are free to ignore that information but with the risk that the migrated image does not represent the same system which is running at the moment. The less issues left in the report the better is the result. In most cases a manual fine tuning is required. This includes the repository selection and the unmanaged files along with the configuration details of your currently running operating system. You should understand the module as a helper to migrate running servers into images. The implementation is still under construction so expect better migration results in future releases :)

```
kiwi { -m | --migrate } name [--exclude directory...] [--skip package...] [--nofiles] [--notemplate]
```

Image Postprocessing Modes

The KIWI post-processing modes are used for special image deployment tasks, like installing the image on a USB stick. So to say they are the third step after preparation and creation. KIWI calls the postprocessing modules automatically according to the specified output image type and attributes but it's also possible to call them manually.

```
kiwi --bootvm initrd --bootvm-system systemImage [--bootvm-disksize size]
```

```
kiwi --bootcd initrd
```

```
kiwi --installed initrd --installed-system vmx-system-image
```

```
kiwi --installstick initrd --installstick-system vmx-system-image
```

Image format conversion

The KIWI format conversion is useful to perform the creation of another image output format like vmdk for VMware or ovf the open virtual machine format. Along with the conversion KIWI also creates the virtual machine configuration according to the format if there is a machine section specified in the XML description

```
kiwi --convert systemImage [--format vmdk|ovf|qcow2|vhd]
```

Helper Tools

The helper tools provide optional functions like creating a crypted password string for the users section of the `config.xml` file as well as signing the image description with an md5sum hash and adding splash data to the boot image used by the bootloader.

```
kiwi --createpassword
```

```
kiwi --createhash image-path
```

```
kiwi { -i | --info } ImagePath [--select repo-patterns|patterns|types|sources|size|profiles|packages|version ]
```

```
kiwi --setup-splash initrd
```

The following list describes the helper tools more detailed

`[--createpassword]`

Create a crypted password hash and prints it on the console. The user can use the string as value for the `pwd` attribute in the XML users section

`[--createhash image-path]`

Sign your image description with a md5sum. The result is written to a file named `.checksum.md` and is checked if KIWI creates an image from this description.

`[-i | --info image-path --select selection]`

List general information about the image description. So far you can get information about the available patterns in the configured repositories with *repo-patterns*, a list of used patterns for this image with *patterns*, a list of supported image types with *types*, a list of source URLs with *sources*, an estimation about the install size and the size of the packages marked as to be deleted with *size*, a list of profiles with *profiles*, a list of solved packages to become installed with *packages*, and the information about the appliance name and version with *version*

`[--setup-splash initrd]`

Create splash screen from the data inside the `initrd` and re-create the `initrd` with the splash screen attached to the `initrd` cpio archive. This enables the kernel to load the splash screen at boot time. If `splashy` is used only a link to the original `initrd` will be created

Global Options

`[--add-profile profile-name]`

Use the specified profile. A profile is a part of the XML image description and therefore can enhance each section with additional information. For example adding packages.

`[--set-repo URL]`

Set/Overwrite the repo URL for the first repo listed in the configuration file that does not have a "fixed" status. The change is temporary and will not be written to the XML file.

`[--set-repo type type]`

Set/Overwrite repo type for the first listed repo. The supported repo types depends on the packagemanager. Commonly supported are rpm-md, rpm-dir and yast2. The change is temporary and will not be written to the XML file.

`[--set-repo alias name]`

Set/Overwrite alias name for the first listed repo. Alias names are optional free form text. If not set the source attribute value is used and builds the alias name by replacing each "/" with a "_". An alias name should be set if the source argument doesn't really explain what this repository contains. The change is temporary and will not be written to the XML file.

`[--set-repo prio number]`

Set/Overwrite priority for the first listed repo. Works with the smart packagemanager only. The Channel priority assigned to all packages available in this channel (0 if not set). If the exact same package is available in more than one channel, the highest priority is used.

`[--add-repo URL, --add-repo type type --add-repo alias name --add-repo prio number]`

Add the given repository and type for this run of an image prepare or upgrade process. Multiple --add-repo/--add-repo type options are possible. The change will not be written to the config.xml file

`[--ignore-repos]`

Ignore all repositories specified so far, in XML or elsewhere. This option should be used in conjunction with subsequent calls to --add-repo to specify repositories at the command-line that override previous specifications.

`[--logfile Filename | terminal]`

Write to the log file *Filename* instead of the terminal.

`[--gzip-cmd cmd]`

Specify an alternate command to run when compressing boot and system images. Command must accept **gzip** options.

`[--log-port PortNumber]`

Set the log server port. By default port 9000 is used. If multiple KIWI processes runs on one system it's recommended to set the logging port per process.

`[--package-manager smart | zypper]`

Set the package manager to use for this image. If set it will temporarily overwrite the value set in the xml description.

`[-A | --target-arch i586 | x86_64 | armv5tel | ppc]`

Set a special target-architecture. This overrides the used architecture for the image-packages in zypp.conf. When used with smart this option doesn't have any effect.

`[--debug]`

Prints a stack trace in case of internal errors

`[--verbose 1 | 2 | 3]`

Controls the verbosity level for the instsource module

Image Preparation Options

`[-r | --root RootPath]`

Set up the physical extend, chroot system below the given root-path path. If no `--root` option is given, KIWI will search for the attribute `defaultroot` in `config.xml`. If no root directory is known, a **mktemp** directory will be created and used as root directory.

`[--force-new-root]`

Force creation of new root directory. If the directory already exists, it is deleted.

Image Upgrade/Preparation Options

`[- -cache directory]`

When specifying a cache directory, KIWI will create a cache each for patterns and packages and re-use them, if possible, for subsequent root tree preparations of this and/or other images

`[- -add-package package]`

Add the given package name to the list of image packages multiple `--add-package` options are possible. The change will not be written to the XML description.

`[- -add-pattern name]`

Add the given pattern name to the list of image packages multiple `--add-pattern` options are possible. The change will not be written to the xml description. Patterns can be handled by SUSE based repositories only.

`[- -del-package package]`

Removes the given package by adding it the list of packages to become removed. The change will not be written to the xml description.

Image Creation Options

`[-d | --destdir DestinationPath]`

Specify destination directory to store the image file(s) If not specified, KIWI will try to find the attribute `defaultdestination` which can be specified in the *preferences* section of the `config.xml` file. If it exists its value is used as destination directory. If no destination information can be found, an error occurs.

`[-t | --type Imagetype]`

Specify the output image type to use for this image. Each type is described in a *type* section of the preferences section. At least one type has to be specified in the `config.xml` description. By default, the types specifying the *primary* attribute will be used. If there is no primary attribute set, the first type section of the preferences section is the primary type. The types are only evaluated when KIWI runs the `--create` step. With the option `--type` one can distinguish between the types stored in `config.xml`

`[-s | --strip]`

Strip shared objects and executables - only makes sense in combination with `--create`

`[- -prebuiltbootimage Directory]`

Search in *Directory* for pre-built boot images.

`[--isocheck]`

in case of an iso image the checkmedia program generates a md5sum into the ISO header. If the `--isocheck` option is specified a new boot menu entry will be generated which allows to check this media

`[--lvm]`

Use the logical volume manager to control the disk. The partition table will include one lvm partition and one standard ext2 boot partition. Use of this option makes sense for the create step only and also only for the image types: vmx, oem, and usb

`[--fs-blocksize number]`

When calling KIWI in creation mode this option will set the block size in bytes. For ISO images with the old style ramdisk setup a blocksize of 4096 bytes is required

`[--fs-journalsize number]`

When calling KIWI in creation mode this option will set the journal size in mega bytes for ext[23] based filesystems and in blocks if the reiser filesystem is used

`[--fs-inodesize number]`

When calling KIWI in creation mode this option will set the inode size in bytes. This option has no effect if the reiser filesystem is used

`[--fs-inoderatio number]`

Set the bytes/inode ratio. This option has no effect if the reiser filesystem is used

`[--fs-max-mount-count number]`

When calling kiwi in creation mode this option will set the number of mounts after which the filesystem will be checked. Set to 0 to disable checks. This option applies only to ext[234] filesystems.

`[--fs-check-interval number]`

When calling kiwi in creation mode this option will set the maximal time between two filesystem checks. Set to 0 to disable time-dependent checks. This option applies only to ext[234] filesystems.

`[--fat-storage size in MB]`

if the syslinux bootlaoder is used this option allows to specify the size of the fat partition. This is useful if the fat space is not only used for booting the system but also for custom data. Therefore this option makes sense when building a USB stick image (image type: usb or oem)

`[--partitioner parted|fdasd]`

Select the tool to create partition tables. Supported are parted and fdasd (s390). By default parted is used

`[--check-kernel]`

Activates check for matching kernels between boot and system image. The kernel check also tries to fix the boot image if no matching kernel was found.

`[--mbrid number]`

Specifies a custom mbrid. The number value is treated as decimal number which is internally translated into a 4byte hex value. The allowed range therefore is from 0x0 to max 0xffffffff. By default kiwi creates a random value

`[--edit-bootconfig script]`

Specifies the location of a custom script which is called right before the bootloader is installed. This allows to modify the bootloader configuration file written by kiwi. The scripts working directory is the one which represents the image structure including the bootloader configuration files. Please have in mind that according to the image type, architecture and bootloader type the files/directory structure and also the name of the bootloader configuration files might be different.

For More Information

More information about KIWI, its files can be found at:

<http://en.opensuse.org/Portal:KIWI>
KIWI wiki

`config.xml`
The configuration XML file that contains every aspect for the image creation.

`file:///usr/share/doc/packages/kiwi/kiwi.pdf`
The system design document which describes some details about the building process.

`file:///usr/share/doc/packages/kiwi/schema/kiwi.xsd.html`
The KIWI RELAX NG XML Schema documentation.

`file:///usr/share/doc/packages/kiwi/schema/test.xsd.html`
The KIWI RELAX NG XML Schema documentation.

kiwi::config.sh

KIWI::config.sh — Configuration File for KIWI image description

Description

The KIWI image description allows to have an optional `config.sh` script in place. This script should be designed to take over control of adding the image operating system configuration. Configuration in that sense means stuff like activating services, creating configuration files, prepare an environment for a firstboot workflow, etc. What you shouldn't do in `config.sh` is breaking your systems integrity by for example removing packages or pieces of software. Something like that can be done in `images.sh`. The `config.sh` script is called *after* the user and groups have been set up. If there are SUSE Linux related YaST XML information, these are validated before `config.sh` is called too. If you exit `config.sh` with an exit code `!= 0` kiwi will exit with an error too.

Example A.1. Template for config.sh

```
#=====
# Functions...
#-----
test -f /.kconfig && . /.kconfig
test -f /.profile && . /.profile

#=====
# Greeting...
#-----
echo "Configure image: [$kiwi_iname]..."

#=====
# Call configuration code/functions
#-----
...

#=====
# Exit safely
#-----
exit
```

Common functions

The `.kconfig` file allows to make use of a common set of functions. Functions specific to SUSE Linux specific begin with the name *suse*. Functions applicable to all linux systems starts with the name *base*. The following list describes the functions available inside the `config.sh` script.

[baseCleanMount]

Umount the system filesystems `/proc`, `/dev/pts`, and `/sys`.

[baseDisableCtrlAltDel]

Disable the **Ctrl-Alt-Del** key sequence setting in `/etc/inittab`

[baseGetPackagesForDeletion]

Return the name(s) of packages which will be deleted

[baseGetProfilesUsed]

Return the name(s) of profiles used to build this image

[baseSetRunlevel {value}]

Set the default run level

[baseSetupBoot]

Set up the linuxrc as init

[baseSetupBusyBox {-f}]

activates busybox if installed for all links from the busybox/busybox.links file—you can choose custom apps to be forced into busybox with the -f option as first parameter, for example:

```
baseSetupBusyBox -f /bin/zcat /bin/vi
```

[baseSetupInPlaceGITRepository]

Create an in place git repository of the root directory. This process may take some time and you may expect problems with binary data handling

[baseSetupInPlaceSVNRepository {path_list}]

Create an in place subversion repository for the specified directories. A standard call could look like this baseSetupInPlaceSVNRepository /etc, /srv, and /var/log

[baseSetupPlainTextGITRepository]

Create an in place git repository of the root directory containing all plain/text files.

[baseSetupUserPermissions]

Search all home directories of all users listed in /etc/passwd and change the ownership of all files to belong to the correct user and group.

[baseStripAndKeep {list of info-files to keep}]

helper function for strip* functions read stdin lines of files to check for removing params: files which should be keep

[baseStripDocs {list of docu names to keep}]

remove all documentation, except one given as parameter

[baseStripInfos {list of info-files to keep}]

remove all info files, except one given as parameter

[baseStripLocales {list of locales}]

remove all locales, except one given as parameter

[baseStripMans {list of manpages to keep}]

remove all manual pages, except one given as parameter example: baseStripMans more less

[baseStripRPM]

remove rpms defined in config.xml under image=delete section

[baseStripTools {list of toolpath} {list of tools}]

helper function for suseStripInitrd function params: toolpath, tools

[baseStripUnusedLibs]

remove libraries which are not directly linked against applications in the bin directories

[baseUpdateSysConfig {filename} {variable} {value}]

update sysconfig variable contents

[Debug {message}]

Helper function to print a message if the variable DEBUG is set to 1

[Echo {echo commandline}]

Helper function to print a message to the controlling terminal

[Rm {list of files}]

Helper function to delete files and announce it to log

[Rpm {rpm commandline}]

Helper function to the RPM function and announce it to log

[suseActivateDefaultServices]

Call all postin scriptlets which among other things activates all required default services using suseInsertService

[suseActivateServices]

Check all services in /etc/init.d/ and activate them by calling suseInsertService

[suseCloneRunlevel {runlevel}]

Clone the given runlevel to work in the same way as the default runlevel 3.

[suseConfig]

Setup keytable language and timezone if specified in config.xml and call SuSEconfig afterwards

[suseInsertService {servicename}]

Recursively insert a service. If there is a service required for this service it will be inserted first. The suse insserv program is used here

[suseRemoveService {servicename}]

Remove a service and its dependent services using the suse insserv program

[suseService {servicename} {on|off}]

Activate/Deactivate a service by using the chkconfig program The function requires the service name and the value on or off as parameters

[suseServiceDefaultOn]

Activates the following services to be on by default using the chkconfig program: boot.rootfsck boot.cleanup boot.localfs boot.localnet boot.clock policykitd dbus consolekit haldaemon network atd syslog cron kbd

[suseSetupProductInformation]

This function will use zypper to search for the installed product and install all product specific packages. This function only makes sense if zypper is used as packagemanager

[suseStripPackager {-a}]

Remove smart or zypper packages and db files Also remove rpm package and db if -a given

Profile environment variables

The .profile environment file contains a specific set of variables which are listed below. Some of the functions above makes use of the variables.

[\$kiwi_compressed]

The value of the compressed attribute set in the type element in config.xml

[\$kiwi_delete]

A list of all packages which are part of the packages section with `type="delete"` in `config.xml`

[\$kiwi_drivers]

A comma separated list of the driver entries as listed in the drivers section of the `config.xml`.

[\$kiwi_iname]

The name of the image as listed in `config.xml`

[\$kiwi_iverison]

The image version string `major.minor.release`

[\$kiwi_keytable]

The contents of the keytable setup as done in `config.xml`

[\$kiwi_language]

The contents of the locale setup as done in `config.xml`

[\$kiwi_profiles]

A list of profiles used to build this image

[\$kiwi_size]

The predefined size value for this image. This is not the computed size but only the optional size value of the preferences section in `config.xml`

[\$kiwi_timezone]

The contents of the timezone setup as done in `config.xml`

[\$kiwi_type]

The basic image type. Can be a simply filesystem image type of `ext2`, `ext3`, `reiserfs`, `squashfs`, `cpio`, or one of the following complex image types: `iso`, `split`, `usb`, `vmx`, `oem`, `xen`, or `pxe`.

kiwi::images.sh

KIWI::images.sh — Configuration File for KIWI image description

Description

The KIWI image description allows to have an optional `images.sh` script in place. This script is called at the beginning of the KIWI create step. It is allowed to remove software there to shrink down the size of the image. Most often `images.sh` is used for boot images because they need to be small. As `images.sh` is called in the create step you should be aware to design the script in a way that it can be called multiple times without shooting itself into its knee. As KIWI allows to create different image types from one previously prepared tree one needs to take into account that `images.sh` can be called more than one time. If you exit `images.sh` with an exit code `!= 0` KIWI will exit with an error too.

Example A.2. Template for images.sh

```
#####
# Functions...
#-----
test -f /.kconfig && . /.kconfig
test -f /.profile && . /.profile

#####
# Greeting...
#-----
echo "Configure image: [$kiwi_iname]..."

#####
# Call configuration code/functions
#-----
...

#####
# Exit safely
#-----
exit
```

Common functions

The `.kconfig` file allows to make use of a common set of functions. Functions specific to SUSE Linux specific begin with the name *suse*. Functions applicable to all linux systems starts with the name *base*. The following list describes the functions available inside the `images.sh` script.

[baseCleanMount]

Unmount the system filesystems `/proc`, `/dev/pts`, and `/sys`.

[baseGetProfilesUsed]

Return the name(s) of profiles used to build this image.

[baseGetPackagesForDeletion]

Return the list of packages setup in the packages `type="delete"` section of the `config.xml` used to build this image.

[baseSetupOEMPartition]

Writes the file `/config.oempartition` depending on the following `config.xml` parameters: `oem-reboot`, `oem-swapspace`, `oem-systemsize`, `oem-swap`, `oem-boot-title`, `oem-re-`

covery, oem-kiwi-initrd. kiwi takes the information from `config.xml` and creates the `config.oempartition` file as part of the automatically created boot image (initrd). The information must be available as part of the boot image because it controls the OEM repartition workflow on first boot of an OEM image. Detailed information about the meaning of each option can be found in the OEM chapter of the KIWI cookbook.

[suseGFXBoot {theme} {loadertype}]

This function requires the `gfxboot` and at least one `bootsplash-theme-*` package to be installed in order to work correctly. The function creates from this package data a graphics boot screen for the `isolinux` and `grub` boot loaders. Additionally it creates the `bootsplash` files for the resolutions 800x600, 1024x768, and 1280x1024

[suseStripKernel]

This function removes all kernel drivers which are not listed in the `*drivers` sections of the `config.xml` file.

[suseStripInitrd]

This function removes a whole bunch of tools binaries and libraries which are not required in order to boot a suse system with KIWI.

[Rm {list of files}]

Helper function to delete files and announce it to log.

[Rpm {rpm commandline}]

Helper function to the `rpm` function and announce it to log.

[Echo {echo commandline}]

Helper function to print a message to the controlling terminal.

[Debug {message}]

Helper function to print a message if the variable `DEBUG` is set to 1.

Profile environment variables

The `.profile` environment file contains a specific set of variables which are listed below. Some of the functions above makes use of the variables.

[\$kiwi_iname]

The name of the image as listed in `config.xml`

[\$kiwi_iverion]

The image version string `major.minor.release`

[\$kiwi_keytable]

The contents of the keytable setup as done in `config.xml`

[\$kiwi_language]

The contents of the locale setup as done in `config.xml`

[\$kiwi_timezone]

The contents of the timezone setup as done in `config.xml`

[\$kiwi_delete]

A list of all packages which are part of the packages section with `type="delete"` in `config.xml`

[\$kiwi_profiles]

A list of profiles used to build this image

[\$kiwi_drivers]

A comma separated list of the driver entries as listed in the drivers section of the config.xml.

[\$kiwi_size]

The predefined size value for this image. This is not the computed size but only the optional size value of the preferences section in config.xml

[\$kiwi_compressed]

The value of the compressed attribute set in the type element in config.xml

[\$kiwi_type]

The basic image type. Can be a simply filesystem image type of ext2, ext3, reiserfs, squashfs, and cpio or one of the following complex image types: iso split usb vmx oem xen pxe

kiwi::kiwirc

KIWI::kiwirc — Resource file for the Kiwi imaging system

Description

The KIWI imaging toolchain supports the use of an optional resource file named `.kiwirc` located in the users home directory.

The file is sourced by a Perl process and thus Perl compatible syntax for the supported variable settings is required.

Example A.3. Template for `.kiwi.rc`

```
$BasePath='/usr/share/kiwi';  
$Gzip='bzip2';  
$LogServerPort='4455';  
$System='/usr/share/kiwi/image';
```

Supported Resource Settings

KIWI recognizes the `BasePath`, `Gzip`, `LogServerPort`, `LuksCipher`, and `System` settings in the `.kiwirc` file.

[BasePath]

Path to the location of the KIWI image system components, such as modules, tests, image descriptions etc.

The default value is `/usr/share/kiwi`

[Gzip]

Specify the compression utility to be used for various compression tasks during image generation.

The default value is **gzip** -9

[LogServerPort]

Specify a port number for log message queuing.

The default value is off

[LuksCipher]

Specify the cipher for the encrypted Luks filesystem.

[System]

Specify the location of the KIWI system image description.

The default value is the value of `BasePath` concatenated with `/image`.

Index

Symbols

** Other system items **

root, 1, 2

A

Amazon Elastic Compute Cloud (see EC2 image)

attributes

alias, 35, 36

arch, 33

blocksize, 31

boot, 10, 10, 25, 25, 26, 26, 26, 26, 37, 37

bootinclude, 11, 11, 13, 13, 37, 40

bootkernel, 25, 60

bootloader-theme, 27

bootprofile, 25, 60

bootsplash-theme, 27

bootstrap, 7

checkprebuilt, 14, 26, 26

compressed, 26

controller, 33, 33, 34, 34

defaultdestination, 27

defaultroot, 27

description, 24

device, 33

displayname, 23

domain, 33

driver, 34, 34

filesystem, 25, 26, 85

flags, 25, 56, 56, 56

format, 25, 60

freespace, 28, 28

fsreadonly, 26

fsreadwrite, 26

group, 34

guestOS, 33

home, 34

HWversion, 33

id, 23, 33, 33, 34, 34, 34, 34

image, 25, 25, 25, 25, 25, 26, 26, 26, 30

imageinclude, 36, 36

inherit, 23

installiso, 83, 84, 84

installstick, 84, 84

interface, 34, 34

kernelcmdline, 27, 31

keytable, 27

kiwirevision, 23

locale, 27

lvm, 26, 26

memory, 33, 33

mode, 34, 64

name, 23, 23, 24, 28, 34

number, 31

onlyRequired, 39

password, 36, 36, 36

path, 36, 36, 36, 37

patternType, 39, 39

plusRecommended, 39

prefer-license, 36

primary, 25

priority, 36

profiles, 24, 24, 24, 24

pwd, 34

realname, 34

rpm-check-signatures, 27

rpm-excludedocs, 27

rpm-force, 27

server, 31

shell, 34

showlicense, 27

size, 28, 28, 31

status, 36, 36

target, 31

timezone, 27

type, 7, 13, 14, 23, 24, 25, 35, 35, 38, 39, 40, 110, 111, 112

unit, 31, 32

username, 36, 36, 36

B

boot parameters, 15

build process, 7

C

checklist, 50

Container image

lxc image, 63

D

devices

/dev/console, 16

/dev/etherd/e0.1, 71

/dev/hda, 69

/dev/hda2, 68

/dev/nb0, 71, 71

/dev/nb1, 71

/dev/ram0, 68

/dev/ram1, 68, 68

/dev/sda2, 70

/dev/sda3, 70

- /dev/sdb1, 71, 71
- /dev/xvda, 88
- directories
 - *boot, 14
 - any, 94, 94
 - /etc, 32, 32, 32
 - /etc/init.d/, 109
 - /etc/lxc/CONTAINER_NAME, 63
 - /home, 29, 29
 - /images/CDs, 53
 - /lib/modules/Version/kernel, 35
 - /media, 53
 - /srv/tftpboot/boot/, 72
 - /tmp, 91, 101
 - /tmp/myec2-result, 91, 94, 94
 - /usr/share/kiwi/image/*boot, 10, 14
 - /usr/share/zoneinfo, 27, 27
 - /var, 56
 - /var/lib/lxc/CONTAINER_NAME, 63
 - /var/lib/tftpboot, 70
 - boot/, 70
 - config/, 22
 - image, 8
 - kiwi-hooks, 11, 11
 - mylxc-result, 63
 - oemboot/suse-SLES11, 14
 - root, 22
 - root/, 43, 50

E

- EC2 images, 89
- environment variables
 - delete, 38
 - RC_LANG, 27

F

- file extensions
 - *.iso, 1, 1
 - *.kiwi, 21
 - .gz, 68
 - .iso, 55, 83, 84, 84, 100
 - .raw, 59, 83, 84
 - .raw.install, 84
 - .vmdk, 60
 - .vmx, 60, 60
- filesystems
 - clifs, 31, 55, 56, 72, 73
 - ext2, 61
 - squashfs, 66, 73, 83
 - tmpfs, 32

H

- hook scripts, 10

I

- images
 - EC2, 89
 - ISO, 55
 - lxc, 63
 - OEM, 83
 - PXE, 65
 - VMX, 59
 - XEN, 87
- initrd customization, 13
- ISO images, 55

K

- KIWI
 - architecture restrictions, 39
 - boot parameters, 15
 - build process, 7
 - Caches, 17
 - checklist, 50
 - common code, 16
 - compressed root, 57
 - config.xml, 23
 - Container image, 63
 - create -- requested image types, 8
 - create -- user defined scripts images.sh, 8
 - cross-platform, 43
 - distribution specific code, 16
 - EC2 image, 89
 - hook scripts, 10
 - hybrid mode, 56
 - Hybrid stick, 57
 - image description, 21
 - image migration, 49
 - initrd customization, 13
 - Installation, 3
 - installation source, 53
 - Introduction, 1
 - ISO image, 55
 - local installation source, 53
 - LVM support, 60, 85
 - maintenance, 45
 - model, 42
 - OEM image, 83
 - OEM stick, 57
 - overlay files, 50
 - overlay filesystem, 56
 - overview, 41
 - patterns, 39
 - physical extends, 47

- pre-built boot images, 14
- prepare -- apply archives, 8
- prepare -- apply overlay tree, 7
- prepare -- create target root directory, 7
- prepare -- install packages, 7
- prepare -- manage target root tree, 8
- prepare -- user defined scripts config.sh, 8
- PXE image, 65
- RAM only image, 73
- release format, 24
- split image, 74
- split mode, 56
- stages, 8
- union image, 73
- USB, 56
- USB sticks, 56
- virtual disk formats, 60
- VMware, 60
- VMX image, 59
- Workflow, 5
- XEN image, 87
- local-ram, 73
- remote-local, 73
- remote-ram, 74
- remote-remote, 74

V

- virtual disk formats, 60
- VMware, 60
- VMX images, 59

X

- XEN image, 87

M

- macros
 - %arch, 36
- manpages
 - kiwi, 100
 - kiwi::config.sh, 107
 - kiwi::images.sh, 111
 - kiwi::kiwirc, 114

O

- OEM images, 83
- overlay files, 50

P

- pre-built boot images, 14
- PXE images, 65

S

- server
 - atftp, 65
 - dhcp, 66
 - TFTP, 67, 67
- services
 - atftpd, 65
 - insserv, 22
 - NFS, 75

U

- union image
 - local-local, 73

