

KIWI Project

openSUSE - KIWI Image System Cookbook

Project, Design and Implementation
by Marcus Schaefer (ms@suse.de)

*



Author: Marcus Schaefer
Date: May 19, 2008
Revision: 2.38

Contents

1	Introduction	5
2	Basic workflow	7
2.1	Boot process	9
2.2	Common and Distribution specific code	10
3	KIWI image description	11
3.1	config.xml	13
4	Maintenance of Operating System Images	25
5	ISO image - Live Systems	27
5.1	Building the suse-live-iso example	27
5.2	Using the image	28
5.3	Flavours	28
6	USB image - Live-Stick System	31
6.1	Building the suse-live-stick example	31
6.2	Using the image	32
6.3	Flavours	33
7	VMX image - Virtual Disks	35
7.1	Building the suse-vm-guest example	35
7.2	Using the image	36
7.3	Flavours	36
8	PXE image - Thin Clients	39
8.1	Setting up the required services	39
8.2	Building the suse-pxe-client example	41
8.3	Using the image	41
8.4	Flavours	42
9	OEM image - Preload Systems	51
9.1	Building the suse-oem-preload example	51
9.2	Using the image	52
9.3	Flavours	52

10 XEN image - Paravirtual Systems	55
10.1 Building the suse-xen-guest example	55
10.2 Using the image	56
10.3 Flavours	56
11 KIWI testsuite	57
11.1 testsuite packages	57
11.2 Creating a test	57
Index	59

1 Introduction

The openSUSE KIWI Image System provides a complete operating system image solution for Linux supported hardware platforms as well as for virtualisation systems like Xen, VMWare, etc. The KIWI architecture was designed as a two level system. The first stage, based on a valid **software package source**, creates a so called **physical extend** according to the provided image description. The second stage creates from a required physical extend an operating system image. The result of the second stage is called a **logical extend** or short an image.

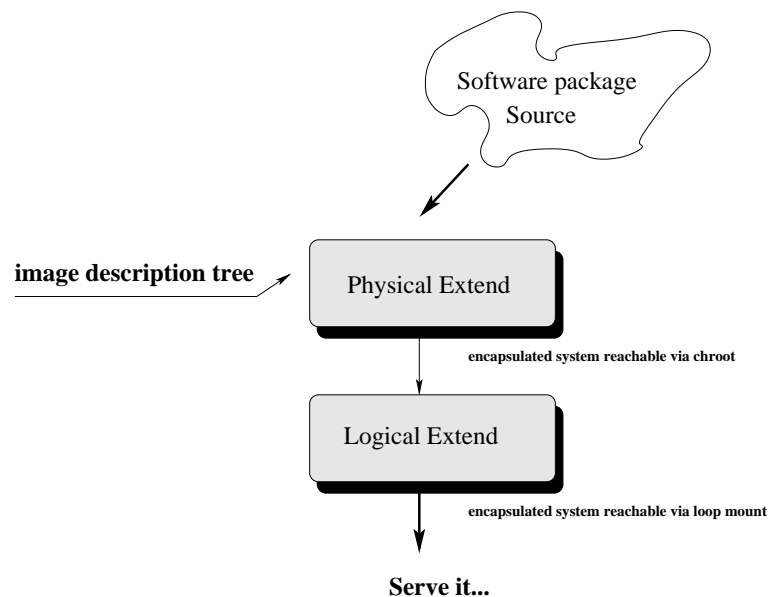


Figure 1.1: Image Serving Architecture

Because this document contains conceptual information about an image system, it is important to understand what an operating system image is all about. A normal installation process is starting from a given installation source and installs single pieces of software until the system is complete. During this process there may be manual user intervention required. However an operating system image represents an already completed *installation* encapsulated as a file and optionally includes the configuration for a specific task. Such an operating system starts working as soon as the image has been brought to a system storage device no matter if this is a volatile or non volatile storage. The process of creating an image takes place without user interaction. This means all requirements of the encapsulated system has to be fulfilled before the image is created. According to this the so called **image description** stores all the information needed to create an image.

2 Basic workflow

Contents

2.1 Boot process	9
2.2 Common and Distribution specific code	10

The creation of an image with KIWI is always divided into two basic steps. These are the **prepare** and the **create** step. the create step requires the prepare step to be exited successfully. Within this first prepare step kiwi builds of a new root tree or, in kiwi-speak, a new physical extend. The building of a new root tree consists of the creation of the directory specified to hold it and the installation of the selected packages on it. The installation of software packages is driven by a packagemanager. KIWI supports the smart and zypper package managers. The prepare step executes the following major stages:

- **Root directory creation**

To prevent accidental deletion of an existing root tree, kiwi will stop with an error message if this folder already exists, unless the option `--force-new-root` is used in which case the existing root will be deleted.

- **Package installation**

First the selected package manager (smart by default) is instructed to use the repositories specified in the image description file. Then the packages specified in the 'bootstrap' section are installed. These packages are installed externally to the target root system (i.e. not chroot'ed) and establish the initial environment so the rest of the process may run chroot'ed. Essential packages in this section are filesystem and glibc-locale. In practice you only need to specify those two, since the rest of the packages will be pulled because of the dependency system. To save space in your image you could schedule a set of packages for deletion after the package installation phase is over by listing them in the 'delete' section.

- **User defined script config.sh**

At the end of the preperation stage the optional script named config.sh is called. This script should be used to configure the system which means for example the activation of services. For a detailed description what functions are already available to configure the system please refer to the KIWI::config.sh manual page

- **Managing the new root tree**

At this point you can make changes on your physical extend so it fits your

purpose better. Bear in mind that changes at this point will be discarded and not repeated automatically if you rerun the 'prepare' phase unless you include them in your original config.xml file and/or config.sh script. Please also note that the image description has been copied into the new root below the directory **<new-root>/image**. Any subsequent create step will read the image description information from the new root tree and not from the original image description location. According to this if you need to change the image description data after the prepare call has finished you need to change it inside the new root tree as well as in your original description directory to prevent losing the change when your root tree will be removed later for some reason.

After the prepare step has finished successfully a subsequent building of an image file or, in kiwi-speak, a new logical extend follows. The building of an image requires a successfully prepared new root tree in the first place. Using this tree multiple image types can be created. So to speak it's possible to create a VMware image and a XEN image from the same prepared root tree. The create step executes the following major stages:

- **User defined script images.sh**

At the beginning of the creation stage the optional script named images.sh is called. This script has no distinctive use case like config.sh but is most often used to remove packages which were pulled in by a dependency but are not really required for the later use of the operating system. For a detailed description what functions are already available to images.sh please refer to the KIWI::images.sh manual page

- **Create the requested image type**

What image type(s) a kiwi image supports depends on what types has been setup in the main image description file config.xml. At least one type must be setup. The following picture shows what image types are currently supported by kiwi:

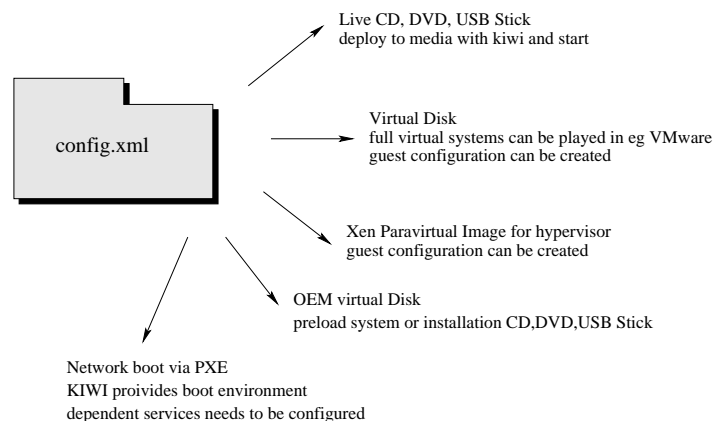


Figure 2.1: Image Types

Detailed information including a step by step guidance how to call kiwi and how to make use of the result image can be found in the image type specific sections later in this document.

2.1 Boot process

Today's linux systems are using a special boot image to control the boot process. This boot image is called **initrd**. The linux kernel loads this initial ramdisk which is a compressed cpio archive into RAM and calls `init` or if present the program named `linuxrc`. The KIWI image system also takes care for the creation of this boot image. Each image type has its own special boot code and shares the common parts in a set of module functions. The image descriptions for the boot images are provided by KIWI and thus the user has in almost all cases no need to take care for the boot image.

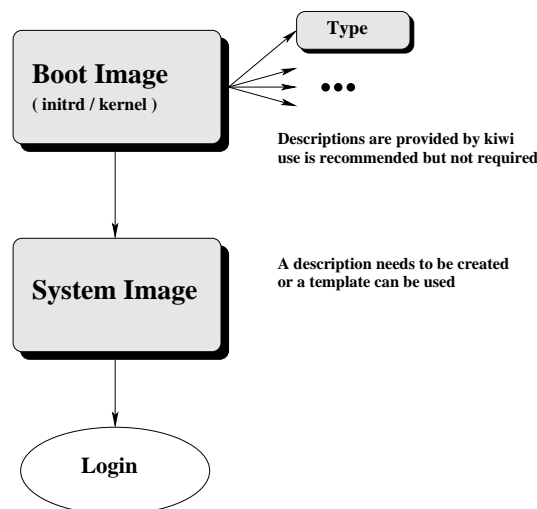


Figure 2.2: Boot process

Furthermore KIWI automatically creates this boot image along with the requested image type. It does that by calling itself in a `prepare` and `create` call. There is no difference in terms of the description of such a boot image compared to the system image description. The system image description is the one the user creates and this image represents the later operating system whereas the boot image only lives temporarily in RAM as long as the system image will be activated. The boot image descriptions are stored in `/usr/share/kiwi/image/*boot` and can be built in the same way as the system image. The boot image without a corresponding system image doesn't make sense though.

2.2 Common and Distribution specific code

KIWI has been developed to be usable for any Linux distribution. By design of a Linux distribution there are differences between each of them. With KIWI we provide on one hand the code which is common to all Linux distributions according to standards and on the other hand there is also code where we have to distinguish between the distribution type.

In case of such specific tasks which are almost all in the area of booting, KIWI provides a set of functions which all have to come with a distribution prefix. As this project uses SUSE Linux as base distribution all required distribution specific tasks has been implemented for SUSE and could be missing for other distributions. The existing implementation for SUSE turns out to be adapted to other distributions very easily though.

A look into the code therefore will show you functions which are prefixed by "suse" as well as scripts whose names starts with "suse-". At any time you see such a script or function you can be assured that this is something distribution specific and needs to be adapted if you plan to use KIWI with another distribution than SUSE. For example the boot workflow is controlled by a program called linuxrc which is in KIWI a script represented by suse-linuxrc. Another example would be the function called suseStripKernel which is able to remove everything but a specified list of kernel drivers from the SUSE kernel.

The prefixed implementation allows us to integrate all the distribution specific tasks into one project but this of course requires the help and knowledge of the people who are familiar with their preferred linux distribution.

3 KIWI image description

Contents

3.1 config.xml	13
--	----

In order to be able to create an image with kiwi a so called image description must be created. The image description is represented by a directory which has to contain at least one file named **config.xml** or alternatively ***.kiwi**. A good start for such a description can be found in the examples provided in `/usr/share/doc/packages/kiwi/examples`.




Figure 3.1: Image description directory

The following additional information is optional for the process of building an image but most often mandatory for the functionality of the later operating system.

- **cdboot/**
For ISO boot images only. The directory contains the configuration files for setting up the isolinux boot loader. Additionally the script `isolinux.sh` must exist which is called to create the iso image by calling `mkisofs`. For SUSE systems `isolinux.sh` is a link to the generic `suse-isolinux` script.
- **images.sh**
Optional configuration script while creating the logical extend. This script is called at the beginning of the image creation process. It is designed to clean-up the image system. Affected are all the programs and files only needed while the physical extend exists.
- **config.sh**
Optional configuration script while creating the physical extend. This

script is called at the end of the installation but **before** the package scripts have run. It is designed to configure the image system, such as the activation or deactivation of certain services (insserv). The call is not made until after the switch to the image has been made with **chroot**.

- **root/**

Subdirectory that contains special files, directories, and scripts for adapting the image environment **after** the installation of all the image packages. The entire directory is copied into the root of the image tree using `cp -a`.

- **config-yast-firstboot.xml**

Configuration file for the control of the yast2 firstboot service. Similar to the autoyast approach yast also provides a boot time service called firstboot. Unfortunately there is no GUI available to setup the firstboot but a good documentation in `/usr/share/doc/packages/yast2-firstboot`. Once you have created such a firstboot file in your image description directory KIWI will process on the file and setup your image as follows:

1. KIWI enables the firstboot service
2. While booting the image YaST is started in firstboot mode
3. The firstboot service handles the instructions listed in the `config-yast-firstboot.xml`
4. If the process finished successfully the environment is cleaned and firstboot won't be called at next reboot.

- **config-yast-autoyast.xml**

Configuration file which has been created by autoyast. To be able to create such an autoyast profile you should first call:

```
yast2 autoyast
```

Once you have saved the information from the autoyast UI as `config-yast-autoyast.xml` file in your image description directory KIWI will process on the file and setup your image as follows:

1. While booting the image YaST is started in autoyast mode automatically
2. The autoyast description is parsed and the instructions are handled by YaST. In other words the **system configuration** is performed
3. If the process finished successfully the environment is cleaned and autoyast won't be called at next reboot.

- **config-cdroot.tgz**

Archive which is used for ISO images only. The data in the archive is

uncompressed and stored in the CD/DVD root directory. This archive can be used, for example, to integrate a license file or readme information directly readable from the CD or DVD.

- **config-cdroot.sh**

Along with the config-cdroot.tgz one can provide a script which allows to manipulate the extracted data.

- **config/**

Optional Subdirectory that contains Bash scripts that are called after the installation of all the image packages, primarily in order to remove the parts of a package that are not needed for the operating system. The name of the Bash script must resemble the package name listed in the config.xml

3.1 config.xml

The mandatory image definition file is divided into different sections which describes information like the image name and type as well as the packages and patterns the image should consist of. The following information explains the basic structure of the XML document. When KIWI is called the XML structure is validated by a RelaxNG based schema. For details on attributes and values please refer to the schema documentation file at `/usr/share/doc/packages/kiwi/kiwi.rng.html`.

```
<image schemeversion="2.0"
      name="iname" inherit="path">
  ...
</image>
```

The image definition starts with an image tag and requires the schema format at version 2.0. The attribute name specifies the name of the image which is also used for the file names created by KIWI. The optional attribute inherit allows to inherit the packages information from another image description. Inside the **image** section the following information applies.

```
<description type="image">
  <author>an author</author>
  <contact>mail</contact>
  <specification>short info</specification>
</description>
```

The mandatory description section contains information about the creator of this image description. The attribute **type** could be either of the value "image" which indicates this is a system image description or at value "boot" for boot image descriptions.

```
<profiles>
  <profile name="name" description="text"/>
  ...
</profiles>
```

The optional profiles section lets you maintain one image description while allowing for variation of the sections packages and drivers that are included. A separate profile element must be specified for each variation. The profile child element, which has name and description attributes, specifies an alias name used to mark sections as belonging to a profile, and a short description explaining what this profile does.

To mark a set of packages/drivers as belonging to a profile, simply annotate them with the **profiles** attribute. It is also possible to mark sections as belonging to multiple profiles by separating the names in the **profiles** attribute with a comma. If a packages/drivers tag does not have a profiles attribute, it is assumed to be present for all profiles.

```
<preferences>
  <version>1.1.2</version>
  <packagemanager>smart</packagemanager>
  <type .../>
</preferences>
```

The mandatory preferences section contains information about the supported image type(s), the used packagemanager, the version of this image and optional attributes. The image version must be a three-part version number of the format: **Major.Minor.Release**. In case of changes to the image description the following rules should apply:

- For smaller image modifications that do not add or remove any new packages, only the release number is incremented. The **config.xml** file remains unchanged.
- For image changes that involve the addition or removal of packages the minor number is incremented and the release number is reset.
- For image changes that change the size of the image file the major number is incremented.

By default kiwi use the **smart** packagemanager but it is also possible to use the SUSE packagemanager called **zypper**.

At least one type must be set to be able to build an image from this description. Multiple type lines are allowed whereas you can specify with the boolean attribute named **primary** which should be the primary image if no type is requested on creation. The following list describes the possible types and their attributes:

- **usb**
Use this type to create a USB stick system along with the attributes **filesystem** and **boot="usbboot/suse-***
- **vmx**
Use this type to create a virtual disk system along with the attributes **filesystem**, **boot="vmxboot/suse-*** and optionally **format**. The format attribute specifies one of the qemu supported virtualisation formats, for example vmdk or qcow2.
- **oem**
Use this type to create a preload virtual disk system along with the attributes **filesystem**, **boot="oemboot/suse-*** and optionally **format**. If the format attribute is set to "iso" or "usb" KIWI will additionally create an installation media suitable for a CD/DVD or an USB stick. This installation media takes over the task of deploying the preload system onto the storage devices which it detects at boot time.
- **pxe**
Use this type to create a network boot image along with the attributes **filesystem** and **boot="netboot/suse-***
- **iso**
Use this type to create a live system on CD or DVD along with the attributes **boot="isoboot/suse-*** and optionally **flags**. If no flags are specified the filesystem will not be compressed and no union filesystem is used. Allowed flags are:
 - **unified**: Compress filesystem with squashfs and mount the system read-write with an aufs based overlay mount
 - **compressed**: Compress filesystem with squashfs and use a link list to mount the system read-write. An additional split section controls the read-write information
 - **unified-cromfs**: Same as unified but uses cromfs
 - **compressed-cromfs**: Same as compressed but uses cromfs

If no flags or the flags "compressed" or "compressed-cromfs" are set an additional **split** section is required.

- **xen**
Use this type to create a Xen enabled para-virtual guest image along with the attributes **filesystem** and **boot="xenboot/suse-"**
- **split**
Use this type if you want to use one of the types **usb,vmx,oem** or **pxe** but as a split image. The split image support allows to create the image as split files whereas one part represents the read-write data and the other part represents the read-only data. Different filesystems can be assigned to each portion. According to this use this type together with the attributes **fsreadwrite, fsreadonly** and **boot="usb|vmx|oem|netboot/suse-"**
- **cpio**
Use this type if your image is a boot image (intrad). Additionally the attribute **bootprofile="default"** can be specified. In a boot image it makes sense to have the different kernels/drivers grouped in a namespace. Within a system image one can select the group by also specifying a bootprofile attribute. If there are profiles defined in a boot image there must be at least one namespace called the **default** namespace.

Within the preferences section there are the following optional attributes:

- **size**
Specifies the size of the image with a numerical value in Megabytes or Gigabytes. Use the "unit" attribute to assign the unit **M** for Megabytes or **G** for Gigabytes. KIWI extends the image size automatically if the specified value is too small. If the actual size is more than 100MB larger than the specified size, KIWI aborts with an error message. KIWI does not automatically reduce the image size if the specified value is too large, because the extra space might be needed to, for example, run custom scripts. If no size is specified, KIWI uses the required size plus approximately 10% free space.
- **compressed**
Specifies whether the image file should be compressed or not. This is not the filesystem compression just the image file compression
- **rpm-check-signatures**
Specifies whether RPM should check the package signature or not
- **rpm-force**
Specifies whether RPM should be called with `-force`
- **keytable**
Specifies the name of the console keymap to use. The value corresponds to a map file in `/usr/share/kbd/keymaps`. The `KEYTABLE` variable in `/etc/sysconfig/keyboard` file is set according to the keyboard mapping.

- **timezone**
Specifies the time zone. Available time zones are located in the /usr/share/zoneinfo directory. Specify the attribute value relative to /usr/share/zoneinfo. For example, specify Europe/Berlin for /usr/share/zoneinfo/Europe/Berlin. KIWI uses this value to configure the timezone in /etc/localtime for the image
- **locale**
Specifies the name of the locale to use, which defines the contents of the RC_LANG system environment variable in /etc/sysconfig/language
- **defaultdestination**
Used if the `-destdir` option is not specified when calling KIWI
- **defaultroot**
Used if the option `-root` is not specified when calling KIWI
- **defaultbaseroot**
Used if the option `-base-root` is not specified when calling KIWI. It's possible to prepare and create an image using a predefined non empty root directory as base information. This could speedup the build process a lot if the base root path already contains most of the image data.

```
<users group="users" id="number">  
  <user pwd="..." home="dir"  
    name="user" id="number"/>  
  ...  
</users> ...
```

The optional `users` element specifies the users to be added to the image. The `group` attribute specifies the group the users belong to. If this group does not exist, it is created. A `user` element must be specified for each group. The `user` child element specifies the users belonging to that group, and the `name`, `pwd` and `home` attributes specifies the username, crypted password, and path to the home directory. The password can be created by the `kiwi --createpassword` tool.

```
<drivers type="type" profiles="name">  
  <file name="filename"/>  
  ...  
</drivers>
```

The optional `drivers` element is only useful for boot images (`initrd`). As a boot image doesn't need to contain the complete kernel one can save a lot of

space if only the required drivers are part of the image. Therefore the drivers section exists. If present only the drivers which matches the file names or glob patterns will be included into the boot image. The type attribute specifies one of the following driver types:

- **drivers**
Each file is specified relative to the `/lib/modules/<Version>/kernel` directory.
- **netdrivers**
Each file is specified relative to the `/lib/modules/<Version>/kernel/drivers/net` directory.
- **scsidrivers**
Each file is specified relative to the `/lib/modules/<Version>/kernel/drivers/scsi`
- **usbdrivers**
Each file is specified relative to the `/lib/modules/<Version>/kernel/drivers/usb` directory.

According to the driver type the specified files are searched in the corresponding directory. The information about the drivernames is provided as environment variable named like the value of the type attribute and is processed by the function **suseStripKernel**. According to this along with a boot image description a script called **images.sh** must exist which calls this function in order to allow the driver information to have any effect.

```
<repository type="type" status="replaceable">
  <source path="URL"/>
</repository>
...
```

The mandatory repository section specifies the source URL and type used by the package manager. The type attribute specifies the repository type which must be supported by the package manager. At the moment KIWI supports the package managers smart and zypper whereas smart has support for more repository types compared to zypper. Therefore the possible values for the type attribute has been copied from smart. The following table shows the possible repo types:

type	smart	zypper
apt-deb	yes	no
apt-rpm	yes	no
deb-dir	yes	no
mirrors	yes	no
red-carpet	yes	yes
rpm-dir	yes	yes
rpm-md	yes	yes
slack-site	yes	no
up2date-mirrors	yes	no
urpmi	yes	no
yast2	yes	yes

The optional attribute **status="replaceable"** makes only sense for boot image descriptions. This attribute indicates that the repository is allowed to become replaced by the repositories defined in the system image descriptions. Because kiwi automatically builds the boot image if required it should create that image from the same repositories which are used to build the system image to make sure both fit together. Therefore it is required to allow the repository to become overwritten which is indicated by the status attribute.

The source child element contains the path attribute, which specifies the location (URL) of the repository. The path specification can be any of the following, and can include the %arch macro which is expanded to the architecture of the image building host.

- **this://<path>**
A relative path name, which is relative to the image description directory being referenced.
- **iso://<path/to/isofile>**
A path to a local .iso file which is then loopback mounted and used as a local path based repository. Alternatively one can do the loop mount himself and point a standard local path to the mounted directory

- **http://<url>**
A http protocol based network location
- **https://<url>**
A https protocol based network location
- **ftp://<url>**
A ftp protocol based network location
- **opensuse://<Project-Name>**
A special http based network location which is created from the given openSUSE buildservice project name. The result is pointing to an rpm-md repository on the openSUSE buildservice. For example: path="opensuse://openSUSE"
- **local path**
A local path which should be an absolute path description

Multiple repository sections are allowed and combined by the used package manager. By default the package manager will always use the latest packages available.

```
<packages type="type" profiles="name"
    patternType="type"
    patternPackageType="type"
    memory="MB"
    disk="/dev/device">
  <package name="name" arch="arch"/>
  <package .../>
  <opensusePattern name="name"/>
  <opensusePattern .../>
  <ignore name="name"/>
  <ignore .../>
</packages>
```

The mandatory packages element specifies the list of packages and pattern names to be used with the image. There are five different types of package sets or patterns, specified with the type attribute:

- **image**
Image packages, list of packages used to finish the image installation. All packages which make up the image are listed here
- **bootstrap**
Bootstrap packages, list of packages used to start creating a new operating system root tree. Basic components which are required to chroot into that system, such as glibc, are listed here.

- **delete**

Delete packages, list of packages stored for later deletion. The package names are available in the `$delete` environment variable of the `/.profile` file created by KIWI. The `baseGetPackagesForDeletion()` function returns the contents of this environment variable, and can be used to delete the packages while ignoring requirements or dependencies. According to this a `config.sh` or `images.sh` script needs to be provided such as the following code snippet shows:

```
rpm -e --nodeps --noscripts \  
$(rpm -q 'baseGetPackagesForDeletion' | grep -v "is not installed")
```

- **xen**

Xen required packages, list of packages used when the image needs support for Xen-based virtualization. The memory and disk attributes define how much memory the virtual system requires and the kind of device the disk should appear as in the virtual instance. This information is used to create the appropriate Xen configuration (`.xenconfig`) file

- **vmware**

VMware required packages, list of packages used when the image needs support for VMware- or generic based full virtualization. The memory and disk attributes define how much memory the virtual system requires and the kind of device the disk should appear as in the virtual instance. This information is used to create the appropriate VMware (`.vmx`) configuration file.

Using a pattern name enhances the package list with a number of additional packages belonging to this pattern. Support for patterns is SUSE-specific, and available with openSUSE 10.1 or later. The optional `patternType` and `patternPackageType` attributes specify which pattern references or packages should be used in a given pattern. The values of these attributes are only evaluated if the KIWI pattern solver is used. If the new (up to SUSE 11.0) `satsolver` pattern solver is used these values are ignored because the `satsolver` can't handle that at the moment. Allowed values for the `pattern*` attributes are:

- **onlyRequired**

Incorporates only patterns and packages that are required by the given pattern

- **plusSuggested**

Incorporates patterns and packages that are required and suggested by the given pattern

- **plusRecommended**

Incorporates patterns and packages that are required and recommended by the given pattern.

By default, only required patterns and packages are used. The result list of packages is solved into a clean conflict free list of packages by the package manager. This for example means that including a suggested package may include required and recommended packages as well according to the dependencies. If a pattern contains unwanted packages, you can use the ignore element to specify an ignore list, with the name attribute containing the package name. Please note that you can't ignore a package if it is required by a package dependency of another package in your list. The packagemanager will automatically pull in the package even if you have ignored it. To restrict a package to a specific architecture, use the arch attribute to specify a comma separated list of allowed architectures.

```
<deploy server="IP" blocksize="4096">
  <timeout>seconds</timeout>
  <commandline>kernel-options</commandline>
  <kernel>kernel-file</kernel>
  <initrd>initrd-file</initrd>
  <partitions device="/dev/sda">
    <partition type="swap" number="1" size="MB"/>
    <partition type="L" number="2" size="MB"
      mountpoint="/" target="true"/>
    <partition type="fd" number="3"/>
  </partitions>
  <union ro="dev" rw="dev" type="aufs|unionfs"/>
  <configuration source="/KIWI/../../file" dest="/../../file"/>
  <configuration .../>
</deploy>
```

The optional deploy section is only useful if you build the **pxe** image type. For this type an additional network boot infrastructure needs to be set up. To ease the process of setting up such a boot server kiwi provides a package called kiwi-pxeboot. This package sets up the basic pxe boot environment like kiwi expects it. The package will setup a directory structure in /srv/tftpboot. The result files of the kiwi image build process needs to be copied to that location. A detailed explanation of what file needs to be copied and where is provided in the PXE Image chapter later in this document. Among the image files itself it is required to provide an information how KIWI should handle the machine it should be install with the created image. Information like which image should be used or how to partition the machine needs to be provided in a file called config.<MAC-Address> below the directory /srv/tftpboot/KIWI. The reason for the deploy section is to allow KIWI to create that file according to the information provided in the image description.

1. The `server` and `blocksize` attributes specify the TFTP server which controls the download of image files. KIWI also supports other protocols than tftp but in order to do that the variables **kiwiserver** and **kiwiservertype** must be set as kernel parameter when the client boots.
2. The optional `timeout` section specifies the grub timeout in seconds which is used when the KIWI initrd configures and installs the grub boot loader on the machine after the first deployment to allow standalone boot.
3. The optional `commandline` section specifies the kernel options which should be passed to the kernel by the grub bootloader. the KIWI initrd includes this kernel options when installing the grub for standalone boot
4. The optional `kernel` and `initrd` sections specifies the KIWI kernel and initrd files on the boot server. In case of a special boot method which is not supported by the distribution standard mkinitrd the KIWI initrd needs to stay on the system and needs to be used for local boot as well. So if your system image makes use of the split type or your deploy section includes any union information the kernel and initrd sections must be provided.
5. The `partitions` section is required if you want to install the system image on a disk or any other permanent storage device. Each partition is specified by one partition subtag which defines the type (see `sfdisk -list-type`), partition number, size, optional mountpoint, and optional information on if this partition is the system image target partition. With the KIWI netboot boot image, the first partition is always the swap partition, while the second partition is used, by default, for the system image. With the optional `target` flag, you can specify a partition other than the second partition to install the system image on. If `size` is set to "image", KIWI calculates the required size for this partition in order to have enough space for the later image.
6. The optional `union` section is used if the system image is based on a read-only filesystem such as squashfs and should be mounted read-write by using an overlay filesystem like aufs or unionfs. In this case, KIWI creates an additional write partition, then combines both partitions with the given overlay filesystem. Currently, there are two such filesystems: unionfs and aufs (aufs is the preferred file system). The partition that holds the read-only system image must be set as the `ro` attribute value, and the partition that serves as the write partition must be set the `rw` attribute value.
7. The optional `configuration` section can be used to integrate a network client's configuration files which are stored remotely on the server. The `source` attribute specifies the path on the server used by a TFTP client program to download the file, and the `dest` attribute specifies the target relative to the root (/) of the network client. Each file is specified by one configuration section.

```
<split>
  <temporary>
    <!-- read/write access to: -->
    <file name="/var"/>
    <file name="/var/*"/>
    <!-- but not on this file: -->
    <except name="/etc/shadow"/>
  </temporary>
  <persistent>
    <!-- persistent read/write access to: -->
    <file name="/etc"/>
    <file name="/etc/*"/>
    <!-- but not on this file: -->
    <except name="/etc/passwd"/>
  </persistent>
</split>
```

The optional split section is used if your image type is split or iso combined with the attributes compressed or compressed-cromfs. The split section controls which files of your splitted image should be writable and whether they are persistantly writable or only temporarily. In case of an iso image all data specified can only be temporarily writable by design of the live system image type.

The split section distinguishes between directory and files. The information of `"/etc"` would make `/etc` a writable directory however none of the files *within* `/etc` are affected. They remain symbolic links to the real files in the read-only area. The main advantage to putting just a directory in the read-write area is that any new files created there are stored on the disk instead of tmpfs. If all the files in `/etc` should also be part of the read-write area and according to this put a complete directory including all its files into the read-write area two lines are required as shown above.

4 Maintenance of Operating System Images

Creating an image often results in an appliance solution for a customer and gives you the freedom of a working solution at that time. But software develops and you don't want your solution to become outdated. Because of this together with an image people always should think of **image-maintenance**. The following paragraph just reflects ideas how to maintain images created by kiwi:

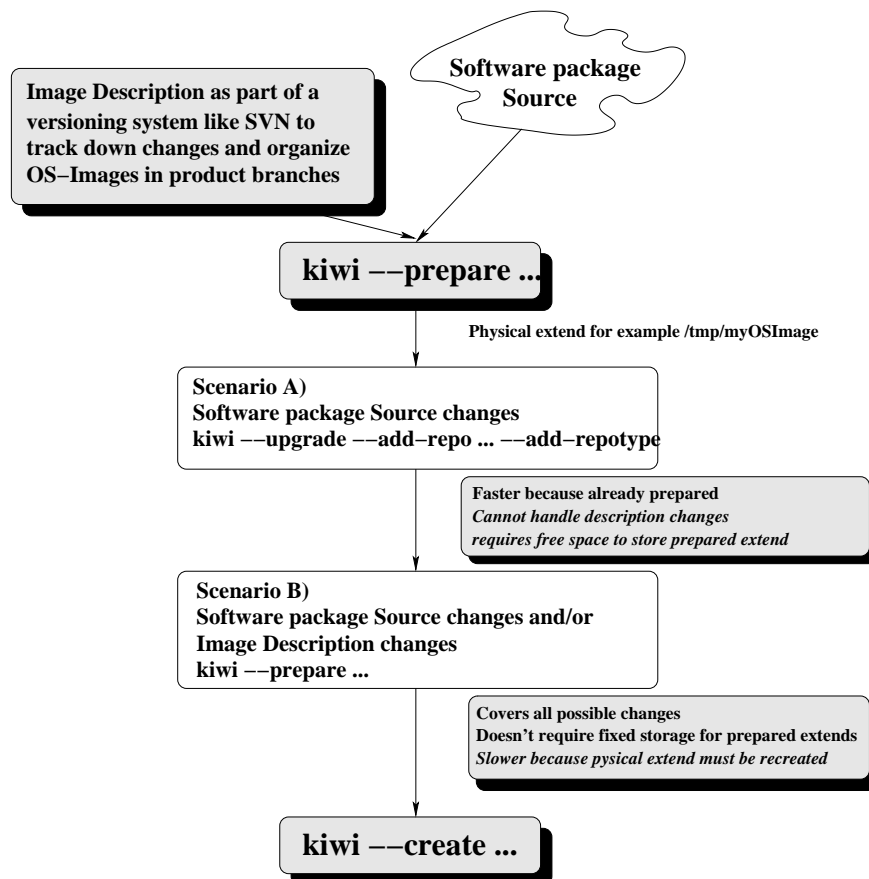


Figure 4.1: Image maintenance scenarios

The picture above shows two possible scenarios which requires an image to become updated. The first reason for updating an image are changes to the software, for example a new kernel should be used. If this change doesn't require additional software or changes in the configuration the update can

be done by kiwi itself using its **upgrade** option. In combination with **upgrade** kiwi allows to add an additional repository which may be needed if the updated software is not part of the original repository. An important thing to know is that this additional repository is **not** stored into the original `config.xml` file of the image description.

Another reason for updating an image beside software updates are configuration changes or enhancements, for example an image should have replaced its browser with another better browser or a new service like apache should be enabled. In principal it's possible to do all those changes manually within the physical extend but concerning maintenance this would be a nightmare. Why, because it will leave the system in an unversioned condition. Nobody knows what has changed since the very first preparation of this image. So in short **don't modify physical extends manually**. Changes to the image configuration should be done within the image description. The image description itself should be part of a versioning system like subversion. All changes can be tracked down then and maybe more important can be assigned to product tags and branches. As a consequence an image must be prepared from scratch and the old physical extend could be removed.

5 ISO image - Live Systems

Contents

5.1 Building the suse-live-iso example	27
5.2 Using the image	28
5.3 Flavours	28
5.3.1 Split mode	29

A live system image is an operating System on CD or DVD. In principal one can treat the CD/DVD as the hard disk of the system with the restriction that you can't write data on it. So as soon as the media is plugged into the computer the machine is able to boot from that media. After some time one can login to the system and work with it like on any other system. All write actions takes place in RAM space and therefore all changes will be lost as soon as the computer shuts down.

5.1 Building the suse-live-iso example

The example provided with kiwi is based on openSUSE 10.3 and includes a KDE and GNOME profile. The source repository is expected to be available below the path /image/CDs/full-10.3-i386. If you don't have a 10.3 repository there edit the config.xml file of the image description and change the path accordingly or use the suggested remote repositories.

```
cd /usr/share/doc/packages/kiwi/examples
kiwi --prepare ./suse-live-iso \
    --root /tmp/myiso --add-profile KDE
```

```
kiwi --create /tmp/myiso \
    --type iso -d /tmp/myiso-result
```

5.2 Using the image

There are two ways to use the generated ISO image:

- Burn the .iso file on a CD or DVD with your preferred burn program. Plug in the CD or DVD into a test computer and (re)boot the machine. Make sure the computer boot from the CD drive as first boot device.
- Use a virtualisation system to test the image directly. Testing an iso can be done with any full virtual system for example:

```
cd /tmp/myiso-result
qemu -cdrom \
    ./suse-10.3-live-iso.i686-2.5.1.iso -m 256
```

5.3 Flavours

KIWI supports different filesystems and boot methods along with the ISO image type. The provided example by default uses a squashfs compressed root filesystem. By design of this filesystem it is not possible to write data on it. To be able to write on the filesystem another filesystem called aufs is used. aufs is an overlay filesystem which allows to combine two different filesystems into one. In case of a live system aufs is used to combine the squashfs compressed read only root tree with a tmpfs RAM filesystem. The result is a full writable root tree whereas all written data lives in RAM and is therefore not persistent. squashfs and/or aufs does not exist on all versions of SUSE and therefore the flags attribute in config.xml exists to be able to have the following alternative solutions:

- **flags="unified"**
Compressed and unified root tree as explained above
- **flags="unified-cromfs"**
Same as unified but uses the cromfs filesystem instead of the squashfs filesystem. cromfs provides better compression but takes longer to become created
- **flags="compressed"**
Does filesystem compression with squashfs but don't use an overlay filesystem for write support. A symbolic link list is used instead and thus a split element is required in config.xml. See the Split mode section below for details.

- **flags="compressed-cromfs"**
Same as compressed but uses the cromfs filesystem instead of the squashfs filesystem.
- **flags not set**
If no flags attribute is set no compressed filesystem and no overlay filesystem will be used. The root tree will be directly part of the ISO filesystem and the paths: /bin, /boot, /lib, /lib64, /opt, /sbin and /usr will be read-only.

5.3.1 Split mode

If no overlay filesystem is in use but the image filesystem is based on a compressed filesystem KIWI allows to setup which files and directories should be writable in a so called split section. In order to allow to login into the system at least the /var directory should be writable because the PAM authentication requires to be able to report any login attempt to /var/log/messages which therefore needs to be writable. The following split section can be used if the flags compressed or compressed-cromfs are used:

```
<split>
  <temporary>
    <!-- allow read/write access to: -->
    <file name="/var"/>
    <file name="/var/*"/>
  </temporary>
</split>
```


6 USB image - Live-Stick System

Contents

6.1 Building the suse-live-stick example	31
6.2 Using the image	32
6.3 Flavours	33
6.3.1 Split stick	33

A live USB stick image is a system on USB stick which allows you to boot and run from this device without using any other storage device of the computer. It is urgently required that the BIOS of the system which you plug the stick in supports booting from USB stick. Almost all new BIOS systems support that. The USB stick serves as OS system disk in this case and you can read and write data onto it.

6.1 Building the suse-live-stick example

The example provided with kiwi is based on openSUSE 10.3 and makes use of the default plus x11 pattern. The operating system is stored on a standard ext3 filesystem. The source repository is expected to be available below the path `/image/CDs/full-10.3-i386`. If you don't have a 10.3 repository there edit the `config.xml` file of the image description and change the path accordingly or use the suggested remote repositories.

```
cd /usr/share/doc/packages/kiwi/examples
kiwi --prepare ./suse-live-stick \
    --root /tmp/mystick
```

There are two possible image types which allows you to drive the stick. Both are added into the `config.xml` of this example image description. If you already have access to the stick you want to run the image on the first approach should be preferred over the second one.

- The first image type named "usb" creates all required images for booting the OS but requires you to plug in the stick and let kiwi deploy the data onto this stick.

```
kiwi --create /tmp/mystick --type usb \  
-d /tmp/mystick-result
```

- The second image type named "oem" allows you to create a virtual disk which represents a virtual disk geometry including all partitions and boot information in one file. You simply can "dd" this file on the stick.

```
kiwi --create /tmp/mystick --type oem \  
-d /tmp/mystick-result
```

6.2 Using the image

To make use of the created images they need to be deployed on the USB stick. For the first image type (usb) you need kiwi itself to be able to deploy the image on the stick. The reason for this is that the usb image type has created the boot and the system image but there is no disk geometry or partition table available. kiwi creates a new partition table on the stick and imports the created images as follows:

```
kiwi --bootstick \  
/tmp/mystick-result/\  
initrd-usbboot-suse-10.3.i686-2.1.1.splash.gz \  
--bootstick-system \  
/tmp/mystick-result/\  
suse-10.3-live-stick.i686-1.1.2
```

In case of the second image type (oem) you only need a tool which allows you to dump data onto a device. On Linux the most popular tool to do this is the **dd** command. The oem image is represented by the file with the .raw extension. As said this is a virtual disk which already includes partition information. But this partition information does not match the real USB stick geometry which means the kiwi boot image (oemboot) has to adapt the disk geometry on first boot. To deploy the image on the stick call:


```
dd if=/tmp/mystick-result/\
   suse-10.3-live-stick.i686-1.1.2.raw \
   of=/dev/<stick-device> bs=32k
```

Testing of the live stick can be done with a test machine which boots from USB or with a virtualisation system. If you test with a virtualisation system for example qemu you should be aware that the USB stick looks like a normal disk to the system. The kiwi boot process searches for the USB stick to be able to mount the correct storage device but in a virtual environment the disk doesn't appear as a USB stick. So if your virtualisation solution doesn't provide a virtual BIOS which allows booting from USB stick you should test the stick on real hardware

6.3 Flavours

USB sticks weren't designed to serve as storage devices for operating systems. By design of these nice little gadgets their storage capacity is limited to only a few G-bytes. According to this KIWI supports compressed filesystems with USB sticks too:

- **filesystem="squashfs"**
This will compress the image using the squashfs filesystem. The boot process will automatically use aufs as overlay filesystem to mount the complete tree read-write. For the write part an additional ext2 partition will be created on the stick. The support for this compression layer requires squashfs and aufs to be present in the distribution KIWI has used to build the image
- **filesystem="cromfs"**
The same as with squashfs but cromfs is used for compression

6.3.1 Split stick

If there is no overlay filesystem available it is also possible to define a split section in config.xml and use the split support to split the image into a compressed read-only and a read-write portion. To create a split stick the types needs to be adapted as follows:

- **type setup for split usb type:**

```
<type fsreadwrite="ext3" fsreadonly="squashfs"
      boot="usbboot/suse-10.3">split</type>
```

- **type setup for split oem type:**

```
<type fsreadwrite="ext3" fsreadonly="squashfs"
      boot="oemboot/suse-10.3">split</type>
```

For both types a split section is required which defines the read-write data. A good starting point is to set /var, /home and /etc as writable data.

```
<split>
  <persistent>
    <!-- allow read/write access to: -->
    <file name="/var"/>
    <file name="/var/*/"/>
    <file name="/etc"/>
    <file name="/etc/*/"/>
    <file name="/home"/>
    <file name="/home/*/"/>
  </persistent>
</split>
```

7 VMX image - Virtual Disks

Contents

7.1 Building the suse-vm-guest example	35
7.2 Using the image	36
7.3 Flavours	36
7.3.1 VMware support	36

A VMX image is a virtual disk image for use in full virtualisation systems like QEMU or VMware. The image represents a file which includes partition data and bootloader information. The size of this virtual disk can be influenced by either the `<size>` element in your `config.xml` file or by the parameter `--bootvm-disksize`

7.1 Building the suse-vm-guest example

The example provided with kiwi is based on openSUSE 10.3 and makes use of the base pattern. The operating system is stored on a standard ext3 filesystem. The source repository is expected to be available below the path `/image/CDs/full-10.3-i386`. If you don't have a 10.3 repository there edit the `config.xml` file of the image description and change the path accordingly or use the suggested remote repositories.

```
cd /usr/share/doc/packages/kiwi/examples
kiwi --prepare ./suse-vm-guest \
    --root /tmp/myvm
```

```
kiwi --create /tmp/myvm \
    --type vmx -d /tmp/myvm-result
```

7.2 Using the image

The generated virtual disk image serves as the harddisk of the selected virtualisation system. The setup of the virtual hard disk differs from the variety of the virtualisation systems. A very simply to use system is the QEMU virtualisation software. To run your image in qemu call:

```
cd /tmp/myvm-result
qemu suse-10.3-vm-guest.i686-1.1.2.raw -m 256
```

7.3 Flavours

Because there are many virtualisation systems available there are also many virtual disk formats. The .raw format KIWI always creates has the same structure as you can find on a real hard disk. For virtualisation software it makes sense to have specific formats to increase the I/O performance when reading or writing data onto the disk from within the virtual system. If you want to tell KIWI to create an additional disk format just extend the type information of the config.xml file by a format attribute.

```
<type ... format='name'>vmx</type>
```

The following table shows a list of supported virtual disk formats

Name	Description
vvfat	Disk format DOS FAT32
vpc	Virtual PC read only disk
bochs	Disk format for Bochs emulator
dmg	Disk format for Mac OS X
cloop	Compressed loop
vmdk	Disk format for VMware
qcow	QEMU virtual disk format
cow	QEMU virtual disk format

7.3.1 VMware support

VMware is a very popular and fast virtualisation platform which is the reason why KIWI has special support for it. VMware requires a so called guest configuration which includes information about what hardware should make up the guest and how much resources should be provided to the guest. With KIWI you can provide the information required to create a guest configuration as part of the config.xml file. Additionally you can group special packages which

you may only need in this virtual environment.

```
<packages type="vmware" memory="512" disk="ide0">  
  <!-- packages you need in VMware only -->  
</packages>
```

If this information is present KIWI will create a VMware guest configuration with 512 MB of RAM and an IDE disk controller interface. Other information like bridged network, CD-ROM drive, power management, etc.. are added as default values and can be changed easily by the native graphics user provided with VMware. The KIWI VMware guest configuration is stored in the file:

```
/tmp/myvm-result/suse-10.3-vm-guest.i686-1.1.2.vmx
```

Together with the **format="vmdk"** attribute KIWI creates a VMware based image (.vmdk file) and the required VMware guest configuration (.vmx)

8 PXE image - Thin Clients

Contents

8.1	Setting up the required services	39
8.1.1	atftp server	39
8.1.2	DHCP server	40
8.2	Building the suse-pxe-client example	41
8.3	Using the image	41
8.4	Flavours	42
8.4.1	The pxe client Control File	42
8.4.2	The pxe client Configuration File	43
8.4.3	User another than tftp as download protocol	47
8.4.4	RAM only image	47
8.4.5	union image	48
8.4.6	split image	48
8.4.7	root tree over NFS	49
8.4.8	root tree over NBD	49

A pxe image consists of a boot image and a system image like all other image types too. But with a pxe image the image files are available separately and needs to be copied at specific locations of a network boot server. PXE is a boot protocol implemented in most BIOS implementations which makes it so interesting. The protocol sends DHCP requests to assign an IP address and after that it uses tftp to download kernel and boot instructions.

8.1 Setting up the required services

Before you start to build pxe images with kiwi you should have setup the boot server. The boot server requires the services **atftp** and **DHCP** to run

8.1.1 atftp server

In order to setup the atftp server the following steps are required

1. install the packages atftp and kiwi-pxeboot

2. edit the file `/etc/sysconfig/atftpd` and set/modify the following variables:
 - `ATFTPD_OPTIONS="--daemon --no-multicast"`
 - `ATFTPD_DIRECTORY="/srv/tftpboot"`
3. run `atftpd` by calling the command: **`rcatftpd start`**

8.1.2 DHCP server

In contrast to the `atftp` server setup the following DHCP server setup can only serve as an example. Please note that according to your network structure the IP addresses, ranges and domain settings needs to be adapted in order to allow the DHCP server to work within your network. If you already have a DHCP server running in your network you should make sure that the filename and next-server information is provided by your server. The following steps describe how to setup a new DHCP server instance:

1. install the package `dhcp-server`
2. create the file `/etc/dhcpd.conf` and include the following statements:

```
option domain-name "example.org";
option domain-name-servers 192.168.100.2;
option broadcast-address 192.168.100.255;
option routers 192.168.100.2;
option subnet-mask 255.255.255.0;
default-lease-time 600;
max-lease-time 7200;
ddns-update-style none; ddns-updates off;
log-facility local7;

subnet 192.168.100.0 netmask 255.255.255.0 {
    filename "pxelinux.0";
    next-server 192.168.100.2;
    range dynamic-bootp 192.168.100.5 192.168.100.20;
}
```

3. edit the file `/etc/sysconfig/dhcpd` and setup the network interface the server should listen on:
 - `DHCPD_INTERFACE="eth0"`
4. run the `dhcp` server by calling: **`redhcpd start`**

8.2 Building the suse-pxe-client example

The example provided with kiwi is based on openSUSE 10.3 and creates an image for a Wyse VX0 terminal with a 128MB flash card and 512MB of RAM. The image makes use of the squashfs compressed filesystem and its root tree is deployed as unified (aufs) based system. The source repository is expected to be available below the path /image/CDs/full-10.3-i386. If you don't have a 10.3 repository there edit the config.xml file of the image description and change the path accordingly or use the suggested remote repository.

```
cd /usr/share/doc/packages/kiwi/examples
kiwi --prepare ./suse-pxe-client \
    --root /tmp/mypxe
```

```
kiwi --create /tmp/mypxe --type pxe \
    -d /tmp/mypxe-result
```

8.3 Using the image

In order to make use of the image all related image parts needs to be copied onto the boot server. According to the example the following steps needs to be performed:

1. Change working directory

```
cd /tmp/mypxe-result
```

2. Copy of the boot and kernel image

```
cp initrd-netboot-suse-10.3.i686-2.1.1.splash.gz \
    /srv/tftpboot/boot/initrd
cp initrd-netboot-suse-10.3.i686-2.1.1.kernel \
    /srv/tftpboot/boot/linux
```

3. Copy of the system image and md5 sum

```
cp suse-10.3-pxe-client.i686-1.2.8 \
    /srv/tftpboot/image
cp suse-10.3-pxe-client.i686-1.2.8.md5 \
    /srv/tftpboot/image
```

4. Copy of the image boot configuration

Normally the boot configuration applies to one client which means it is required to obtain the MAC address of this client. If the boot configuration should be used globally the KIWI generated file can be copied as `config.default`

```
cp suse-10.3-pxe-client.i686-1.2.8.config \  
  /srv/tftpboot/KIWI/config.<MAC>
```

5. Check the PXE configuration file

The PXE configuration controls which kernel and initrd are loaded and which kernel parameters are set. When installing the `kiwi-pxeboot` package a default configuration is added. To make sure the configuration is valid according to this example the file `/srv/tftpboot/pxelinux.cfg/default` should provide the following information:

```
DEFAULT KIWI-Boot  
  
LABEL KIWI-Boot  
    kernel boot/linux  
    append initrd=boot/initrd vga=0x314  
    IPAPPEND 1  
  
LABEL Local-Boot  
    localboot 0
```

6. connect the client to the network and boot

8.4 Flavours

All the different PXE boot based deployment methods are controlled by the `config.<MAC>` (or `config.default`) file. When a new client boots up and there is no client configuration file the new client is registered by uploading a control file to the tftp server. The following sections inform about the control and the configuration file.

8.4.1 The pxe client Control File

This section describes the netboot client control file:

`hwtype.<MAC Address>`

The control file is primarily used to set up new netboot clients. In this case, there is no configuration file corresponding to the client MAC address avail-

able. Using the MAC address information, the control file is created, which is uploaded to the TFTP servers upload directory `/var/lib/tftpboot/upload`.

8.4.2 The pxe client Configuration File

This section describes the netboot client configuration file:

```
config.<MAC Address>
```

The configuration file contains data about image, configuration, synchronization, or partition parameters. The configuration file is loaded from the TFTP server directory `/var/lib/tftpboot/KIWI` via TFTP for previously installed netboot clients. New netboot clients are immediately registered and a new configuration file with the corresponding MAC address is created. The standard case for the deployment of a pxe image is one image file based on a read-write filesystem which is stored onto a local storage device of the client. Below, find an example to cover this case.

```
DISK=/dev/sda
PART=5;S;x,x;L;/
IMAGE=/dev/sda2;suse-10.3-pxe-client.i686;1.2.8;192.168.100.2;4096
```

The following format is used:

```
IMAGE=device;name;version;srvip;bsize;compressed,...,
CONF=src;dest;srvip;bsize,..., src;dest;srvip;bsize
PART=size;id;Mount,...,size;id;Mount
DISK=device
```

- **IMAGE**

Specifies which image (name) should be loaded with which version (version) and to which storage device (device) it should be linked, e.g., `/dev/ram1` or `/dev/hda2`. The netboot client partition (device) `hda2` defines the root file system "/" and `hda1` is used for the swap partition. The numbering of the hard disk device should not be confused with the RAM disk device, where `/dev/ram0` is used for the initial RAM disk and can not be used as storage device for the second stage system image. SUSE recommends to use the device `/dev/ram1` for the RAM disk. If the hard drive is used, a corresponding partitioning must be performed.

- **srvip**

Specifies the server IP address for the TFTP download. Must always be indicated, except in PART.

- **bsize**
Specifies the block size for the TFTP download. Must always be indicated, except in PART. If the block size is too small according to the maximum number of data packages (32768), **linuxrc** will automatically calculate a new blocksize for the download.
- **compressed**
Specifies if the image file on the TFTP server is compressed and handles it accordingly. To specify a compressed image download only the keyword "**compressed**" needs to be added. If compressed is not specified the standard download workflow is used. **Note:** The download will fail if you specify "compressed" and the image isn't compressed. It will also fail if you don't specify "compressed" but the image is compressed. The name of the compressed image has to contain the suffix **.gz** and needs to be compressed with the **gzip** tool. Using a compressed image will automatically **deactivate** the multicast download option of **atftp**.
- **CONF**
Specifies a comma-separated list of source:target configuration files. The source (src) corresponds to the path on the TFTP server and is loaded via TFTP. The download is made to the file on the netboot client indicated by the target (dest).
- **PART**
Specifies the partitioning data. The comma-separated list must contain the size (size), the type number (id), and the mount point (Mount). The size is measured in MB by default. Additionally all size specifications supported by the **sfdisk** program are allowed as well. The type number specifies the ID of the partition. Valid ID's are listed via the *sfdisk – list-types* command. The mount specifies the directory the partition is mounted to.
 - The first element of the list must define the swap partition.
 - The second element of the list must define the **root** partition.
 - The swap partition must not contain a mount point. A lowercase letter **x** must be set instead.
 - If a partition should take all the space left on a disk one can set a lower **x** letter as size specification.
- **DISK**
Specifies the hard disk. Used only with PART and defines the device via which the hard disk can be addressed, e.g., **/dev/hda**.
- **RELOAD_IMAGE**
If set to a non-empty string, forces the configured image to be loaded from the server even if the image on the disk is up-to-date. Used mainly for debugging purposes, this option only makes sense on diskful systems.

- **RELOAD_CONFIG**

If set to a non-empty string, forces all config files to be loaded from the server. Used mainly for debugging purposes, this option only makes sense on diskful systems.

- **COMBINED_IMAGE**

If set to a non-empty string, indicates that the both image specified needs to be combined into one bootable image, whereas the first image defines the read-write part and the second image defines the read-only part.

- **KIWI_INITRD**

Specifies the kiwi initrd to be used for local boot of the system. The variable's value must be set to the name of the initrd file which is used via PXE network boot. If the standard tftp setup suggested with the kiwi-pxeboot package is used all initrd files reside in the **boot/** directory below the tftp server path `/var/lib/tftpboot`. Because the tftpserver does a chroot into the tftp server path you need to specify the initrd file as the following example shows: **KIWI_INITRD=/boot/<name-of-initrd-file>**

- **UNIONFS_CONFIG**

For netboot and usbboot images there is the possibility to use unionfs or aufs as container filesystem in combination with a compressed system image. The recommended compressed filesystem type for the system image is **squashfs**. In case of a usb-stick system the usbboot image will automatically setup the unionfs/aufs filesystem. In case of a PXE network image the netboot image requires a config.<MAC> setup like the following example shows: **UNIONFS_CONFIG=/dev/sda2,/dev/sda3,aufs**. In this example the first device `/dev/sda2` represents the read/write filesystem and the second device `/dev/sda3` represents the compressed system image filesystem. The container filesystem aufs is then used to cover the read/write layer with the read-only device to one read/write filesystem. If a file on the read-only device is going to be written the changes inodes are part of the read/write filesystem. Please note the device specifications in UNIONFS_CONFIG must correspond with the IMAGE and PART information. The following example should explain the interconnections:

```
IMAGE=/dev/sda3;image/myImage;1.1.1;192.168.1.1;4096
PART=200;S;x,300;L;/,x;L;x
UNIONFS_CONFIG=/dev/sda2,/dev/sda3,aufs
DISK=/dev/sda
```

As the second element of the PART list must define the **root** partition it's absolutely important that the first device in UNIONFS_CONFIG references this device as read/write device. The second device of UNIONFS_CONFIG has to reference the given IMAGE device name.

- **KIWI_KERNEL_OPTIONS**

Specifies additional command line options to be passed to the kernel when booting from disk. For instance, to enable a splash screen, you might use 'vga=0x317 splash=silent'.

- **KIWI_BOOT_TIMEOUT**

Specifies the number of seconds to wait at the grub boot screen when doing a local boot. The default is 10.

- **NBDROOT**

Mount the system image root filesystem remotely via NBD (Network Block Device). This means there is a server which exports the root directory of the system image via a specified port. The kernel provides the block layer, together with a remote port that uses the nbd-server program. For more information on how to set up the server, see the nbd-server man pages. The kernel on the remote client can set up a special network block device named /dev/nb0 using the nbd-client command. After this device exists, the mount program is used to mount the root filesystem. To allow the KIWI boot image to use that, the following information must be provided:

```
NBDROOT=NBD.Server.IP.address;\
        NBD-Port-Number;/dev/NBD-Device;\
        NBD-Swap-Port-Number;/dev/NBD-Swap-Device
```

The NBD-Device, NBD-Swap-Port-Number, and NBD-Swap-Device variables are optional. If they are not set, the default values are used (/dev/nb0 for the NBD-Device, port number 9210 for the NBD-Swap-Port-Number, and /dev/nb1 for the NBD-Swap-Device). The swap space over the network using a network block device is only established if the client has less than 48 MB of RAM.

- **NFSROOT**

Mount the system image root filesystem remotely via NFS (Network File System). This means there is a server which exports the root filesystem of the network client in such a way that the client can mount it read/write. In order to do that, the boot image must know the server IP address and the path name where the root directory exists on this server. The information must be provided as in the following example:

```
NFSROOT=NFS.Server.IP.address;/path/to/root/tree
```

- **KIWI_INITRD**

Specifies the KIWI initrd to be used for a local boot of the system. The value must be set to the name of the initrd file which is used via PXE network boot. If the standard TFTP setup suggested with the kiwi-pxeboot package is used, all initrd files reside in the /srv/tftpboot/boot/ directory. Because the TFTP server does a chroot into the TFTP server path, you must specify the initrd file as follows:

```
KIWI_INITRD=/boot/name-of-initrd-file
```

- **KIWI_KERNEL**

Specifies the kernel to be used for a local boot of the system. The same path rules as described for KIWI_INITRD apply for the kernel setup:

```
KIWI_KERNEL=/boot/name-of-kernel-file
```

8.4.3 User another than tftp as download protocol

By default all downloads controlled by the kiwi linuxrc code are performed by an `atftp` call and therefore uses the `tftp` protocol. With PXE the download protocol is fixed and thus you can't change the way how the kernel and the boot image (`initrd`) is downloaded. As soon as linux takes over control the following download protocols `http`, `https` and `ftp` are supported too. KIWI makes use of the **curl** program to support the additional protocols.

In order to select one of the additional download protocols the following kernel parameters need to be setup:

- **kiwiserver**

Name or IP address of the server who implements the protocol

- **kiwiservertype**

Name of the download protocol which could be one of `http`, `https` or `ftp`

To setup these parameters edit the file `/srv/tftpboot/pxelinux.cfg/default` on your PXE boot server and change the **append** line accordingly. Please note all downloads except for kernel and `initrd` are now controlled by the given server and protocol. You need to make sure that this server provides the same directory and file structure as initially provided by the `kiwi-pxeboot` package.

8.4.4 RAM only image

If there is no local storage and no remote root mount setup the image can be stored into the main memory of the client. Please be aware that there should be still enough RAM space available for the operating system after the image has been deployed into RAM. Below, find an example:

- use a read-write filesystem in `config.xml`, for example **filesystem="ext3"**
- sample `config.<MAC>`

```
IMAGE=/dev/ram1;suse-10.3-pxe-client.i686;\
1.2.8;192.168.100.2;4096
```

8.4.5 union image

As used in the suse-pxe-client example it is possible to make use of the aufs or unionfs overlay filesystems to combine two filesystems into one. In case of thin clients there is often the need for a compressed filesystem due to space limitations. Unfortunately all common compressed filesystems provides only read-only access. Combining a read-only filesystem with a read-write filesystem is a solution for this problem. Below, find an example:

```
DISK=/dev/sda
PART=5;S;x,62;L;/,x;L;x,
IMAGE=/dev/sda2;suse-10.3-pxe-client.i386;\
      1.2.8;192.168.100.2;4096
UNIONFS_CONFIG=/dev/sda3,/dev/sda2,aufs
KIWI_INITRD=/boot/initrd
```

8.4.6 split image

As an alternative to the UNIONFS_CONFIG method it is also possible to create a split image and combine the two portions with the COMBINED_IMAGE method. This allows to use different filesystems without the need for an overlay filesystem to combine them together. Below find an example:

- add a split type in config.xml, for example
<type fsreadonly="squashfs" fsreadwrite="ext3" boot="netboot/suse-10.3">split</type>
- add a split section to describe the writable portion, for example:

```
<split>
  <persistent>
    <!-- allow read/write access to: -->
    <file name="/var"/>
    <file name="/var/*"/>
    <file name="/etc"/>
    <file name="/etc/*"/>
    <file name="/home"/>
    <file name="/home/*"/>
  </persistent>
</split>
```

- sample config.<MAC>


```
IMAGE=/dev/sda2;suse-10.3-pxe-client.i686;\
    1.2.8;192.168.100.2;4096,\
    /dev/sda3;suse-10.3-pxe-client-read-write.i686;\
    1.2.8;192.168.100.2;4096
PART=200;S;x,500;L;/,x;L;
DISK=/dev/sda
COMBINED_IMAGE=yes
KIWI_INITRD=/boot/initrd
```

8.4.7 root tree over NFS

Instead of installing the image onto a local storage device of the client it is also possible to let the client mount the root tree via an NFS remote mount. Below find an example:

- Export the kiwi prepared tree via NFS
- sample config.<MAC>

```
NFSROOT=192.168.100.7;/tmp/kiwi.nfsroot
```

8.4.8 root tree over NBD

As an alternative for root over NFS it is also possible to let the client mount the root tree via a special network block device. Below find an example:

- Use nbd-server to export the kiwi prepared tree
- sample config.<MAC>

```
NBDROOT=192.168.100.7;2000;/dev/nbd0
```


9 OEM image - Preload Systems

Contents

9.1 Building the suse-oem-preload example	51
9.2 Using the image	52
9.3 Flavours	52
9.3.1 Influencing the oem partitioning	53

An oem image is a virtual disk image representing all partitions and boot-loader information like it exists on a real disk. The image format is the same compared to the VMX image type. All flavours explained in the VMX chapter also applies to the OEM type.

The original idea of an oem image is to provide this virtual disk data to OEM vendors which now are able to deploy the system independently onto their storage media. The deployment can happen from any OS including Windows if a tool to dump data on a disk device exists. The oem image type is also used to deploy images on USB sticks because in principal it is the same workflow.

9.1 Building the suse-oem-preload example

The example provided with kiwi is based on openSUSE 10.3 and includes the patterns default plus x11. The image type is a split type whereas the read-write filesystem is ext3 and the read-only filesystem is squashfs. The additional format attribute also creates an installable ISO image for deploying the image from CD. The source repository is expected to be available below the path /image/CDs/full-10.3-i386. If you don't have a 10.3 repository there edit the config.xml file of the image description and change the path accordingly or use the suggested remote repository.

```
cd /usr/share/doc/packages/kiwi/examples
kiwi --prepare ./suse-oem-preload \
    --root /tmp/myoem
```

```
kiwi --create /tmp/myoem --type split \  
-d /tmp/myoem-result
```

9.2 Using the image

Testing the oem virtual disk can be done with a virtualisation software like QEMU or VMware. The virtual disk is represented by the .raw extension whereas the .iso extension represents the installation disk for this oem image. The installation disk should be tested on a bare test system For the .raw test just call:

```
cd /tmp/myoem-result  
qemu suse-10.3-oem-preload.i686-1.1.2.raw \  
-m 512
```

or dump the image on a test hard disk and select it as boot device in the BIOS:

```
cd /tmp/myoem-result  
dd if=suse-10.3-oem-preload.i686-1.1.2.raw \  
of=/dev/<device> bs=32k
```

Please note if you test an oem image the virtual disk geometry of the image is the same as the disk geometry inside the host system. According to this the oem boot workflow will skip the re-partitioning which is performed if there would be a real disk

9.3 Flavours

An interesting part of an oem image is that it can be turned into an installation image too. This means it is possible to create an installation CD / DVD or USB stick which deploys the oem based image onto the selected storage device. The installation process is a simply dd of the image onto the selected device so don't expect any user interaction or GUI here to pop up. KIWI supports two types of installation media:

- `<type ... format="iso">...</type>`
Creates a .iso file which can be burned in CD or DVD. This represents an installation CD

- `<type ... format="usb">...</type>`
Creates a `.raw.install` file which can be dumped (dd) on a USB stick. This represents an installation Stick

9.3.1 Influencing the oem partitioning

By default the oemboot process will create/modify a swap, `/home` and `/` partition. It is possible to influence the behavior by the following `oem-*` elements which can be optionally specified within the **preferences** section of your system image XML description. KIWI uses this to create the file `/config.oempartition` as part of the automatically created oemboot boot image. The format of the file is a simple `key=value` format and created by the `KIWIConfig.sh` function named `baseSetupOEMPartition`. Following `oem-*` elements can be specified:

- `<oem-reboot>yes|no</oem-reboot>`
This allows to reboot the oem system after initial deployment. This value is represented by the variable `OEM_REBOOT` in `config.oempartition`
- `<oem-swapsize>number in MB</oem-swapsize>`
Set the size of the swap partition. This value is represented by variable `OEM_SWAPSIZE` in `config.oempartition`
- `<oem-systemsize>number in MB</oem-systemsize>`
Set the size of the `/` partition. This value is represented by the variable `OEM_SYSTEMSIZE` in `config.oempartition`
- `<oem-home>yes|no</oem-home>`
Specify if a home partition should be create. This value is represented by the variable `OEM_WITHOUTHOME` in `config.oempartition`.
- `<oem-swap>yes|no</oem-swap>`
Specify if a swap partition should be create. This value is represented by the variable `OEM_WITHOUTSWAP` in `config.oempartition`.
- `<oem-boot-title>text</oem-boot-title>`
By default the string **OEM** will be appended to the boot manager menu when KIWI creates the grub configuration during first deployment. The `oem-boot-title` value allows to set a custom name which is used instead of **OEM**. This value is represented by the variable `OEM_BOOT_TITLE` in `config.oempartition`.

10 XEN image - Paravirtual Systems

Contents

10.1 Building the suse-xen-guest example	55
10.2 Using the image	56
10.3 Flavours	56

Xen is a free software virtual machine monitor. It allows several guest operating systems to be executed on the same computer hardware at the same time.

A Xen system is structured with the Xen hypervisor as the lowest and most privileged layer.[1] Above this layer are one or more guest operating systems, which the hypervisor schedules across the physical CPUs. The first guest operating system, called in Xen terminology "domain 0" (dom0), is booted automatically when the hypervisor boots and given special management privileges and direct access to the physical hardware. The system administrator logs into dom0 in order to start any further guest operating systems, called "domain U" (domU) in Xen terminology.

A xen image is a filesystem based image file which requires the Xen dom0 running or the project called Xenner which emulates the capabilities of the domain 0. The image created with kiwi can only be used together with the xen tools.

10.1 Building the suse-xen-guest example

The example provided with kiwi is based on openSUSE 10.3 and includes the base pattern. The source repository is expected to be available below the path /image/CDs/full-10.3-i386. If you don't have a 10.3 repository there edit the config.xml file of the image description and change the path accordingly or use the suggested remote repository.

```
cd /usr/share/doc/packages/kiwi/examples
kiwi --prepare ./suse-xen-guest \
    --root /tmp/myxen
```

```
kiwi --create /tmp/myxen \  
    --type xen -d /tmp/myxen-result
```

10.2 Using the image

In order to run a domain U the Xen tool **xm** needs to be called in conjunction with the KIWI generated domain U configuration file

```
xm create -c \  
    /tmp/myxen-result/  
    suse-10.3-xen-guest.i686-1.1.2.xenconfig
```

10.3 Flavours

With KIWI you can provide the information required to create a guest configuration as part of the config.xml file. Additionally you can group special packages which you may only need in this para virtual environment.

```
<packages type="xen" memory="512" disk="/dev/sda">  
    <!-- packages you need in Xen only -->  
    <package name="kernel-xen"/>  
    <package name="xen"/>  
</packages>
```

If this information is present KIWI will create a Xen domain U configuration with 512 MB of RAM and expects the disk at /dev/sda. Other information like bridged network, power management, etc.. needs to be added manually at the moment. The KIWI Xen domain U configuration is stored in the file:

```
/tmp/myxen-result/  
    suse-10.3-xen-guest.i686-1.1.2.xenconfig
```


11 KIWI testsuite

Contents

11.1 testsuite packages	57
11.2 Creating a test	57

The KIWI test suite is useful to perform basic quality checks on the image root directory. The test cases are stored in subdirectories below `/usr/share/kiwi/tests`. To run the testsuite call kiwi as follows:

```
kiwi --testsuite <image-root> \  
[ --test name --test name ... ]
```

If not test names are set the default tests rpm and ldd run. The name of a test corresponds with the name of the directory the test is implemented in.

11.1 testsuite packages

If a test requires special software to be installed but this software is not an essential part of the image itself it can be specified as testsuite packages in the system image config.xml as follows:

```
<packages type="testsuite">  
  <package name="..."/>  
</packages>
```

The testsuite packages are installed when calling kiwi with the testsuite option and are removed after the tests has finished.

11.2 Creating a test

The test itself is defined by a xml description "test-case.xml" and its template definition file `/usr/share/kiwi/modules/KIWISchemeTest.rnc` The following example shows the basic structure of the rpm test:

```
<test_case
  name="rpm"
  summary="check rpm database and verify all rpms"
  description="check if rpm db is present, run rpm's build-in Verify method"

  <requirements>
    <req type="directory">/var/lib/rpm</req>
    <req type="file">/var/lib/rpm/__db.000</req>
    <req type="file">/var/lib/rpm/Packages</req>
  </requirements>

  <test type="binary" place="extern">
    <file>rpm.sh</file>
    <params>CHROOT</params>
  </test>
</test_case>
```

There are basically two sections called "requirements" and "test". In requirements you define what files/directories or packages has to be present in your image to run the test. For example if you need to check the rpm database, the database has to be present within the image. All requirements are checked, and if any of them fail the test won't be executed and an error message is printed. There are three types of requirements:

- **file**
Existence of a file
- **directory**
Existence of a directory
- **rpm-package**
Existence of a package

The test section defines the test script. It could be a binary, shell script or any other kind of executable. Scripts are expected to be in the same directory as where the xml definition for the test resides. There are two types of scripts, extern and intern.

- external scripts are executed outside of the image and are preferred. Their first parameter should be CHROOT. This parameter is changed to the real path of the image chroot directory.
- internal scripts are executed inside image using the "chroot" command. Files are copied into the image and deleted after execution.

A test script always has to return 0 in case of a test to pass, or 1 if any error occur. All messages printed to standard and error output are stored and printed out of the test has failed.

Index

configuration files

- config.<MAC Address>, [43](#)
- hwtype.<MAC Address>, [42](#)

KIWI images

- description, [10](#)
- iso, [26](#)
- maintenance, [24](#)
- oem, [49](#)
- pxe, [37](#)
- testing, [56](#)
- usb, [29](#)
- vmx, [34](#)
- workflow, [5](#)
- xen, [53](#)